
SharpeArena: A Deterministic Point-in-Time Evaluation Sandbox and Reinforcement-Learning Environment for Trading Agents

Tiberiu Toca
General Liquidity
tibi.toca@gmail.com

Abstract

Reinforcement-learning results in trading are produced by the party being evaluated, on an environment that party controls and defends as realistic on the strength of its frictions: transaction costs, slippage, position limits. Frictions constrain the agent, and say nothing about whether the simulated world behaves like a market or about whether the trajectory the environment emitted is the one it ran. We present SharpeArena, a deterministic point-in-time evaluation sandbox and reinforcement-learning environment for trading agents. It excludes three failures at the interface rather than discouraging them: no data-layer operation returns a bar past the engine's cursor, committed canonical fixtures and fingerprints pin the determinism-critical core across tested surfaces, and replay recomputes performance from recorded decisions and frozen engine inputs, so a doctored decision stream produces a different run. Agents are external programs, language models included, that reach the engine through a governed JSON contract. A host-side parser enforces episode semantics, and the checked scoring paths record a wire fault as a typed failed cell rather than as a hold. Around that loop the environment adds a strategy-generation path over a closed non-executable candidate language, a forward paper-trading arm, and prospective forecast contracts that freeze the target, information boundary, deadline, resolution rule and score before submission. Forecast revisions remain append-only, and the producer exports raw evidence for the separate SharpeBench kernel to score without affecting trading rank. Selected deadline, revision, binding and fixed-point Brier rules are modeled in Lean; separate executable tests check the same rules against the implementation, and nothing mechanically links the model to the code. We calibrate the environment's diagnostics with fixed reference policies: nothing is trained and no admitted result is a claim about model quality. The committed F1–F8 artifacts were produced before most current engineering guards, F7 alone having been regenerated for release 0.29.0, and a superseded convenience-sample protocol pilot produced no admitted model-validation result. We find that the canonical generator fails its own realism diagnostic on 23 of 24 seeded panels, that a predeclared sweep of its volatility-clustering knobs finds no Calm setting that survives, that no reference policy is rank-eligible, and that a single-seed reading of the ecology probe does not survive replication. We offer a producer whose current guarantees are properties of committed code rather than of its authors' good faith, and diagnostics that report where its historical tape is not yet market-like.

1 Introduction

Every published number in agent trading evaluation is produced by the party being evaluated, on an environment that party controls. Three failures upstream of scoring void any score, and none of them is visible in the score itself. An agent that observes one future bar has an edge no deflation can remove; look-ahead bugs are a pervasive and usually invisible failure mode of backtests, alongside the selection bias the deflation literature centers [Bailey and de Prado, 2014]. Gençay [2026, p. 7] shows the first point directly: a deliberately leaky oracle posts a deflated Sharpe ratio of 1.00, because deflation has no mechanism against a strategy whose information set is contaminated. Authors of recent systems have found such leaks in their own pipelines only after a result looked implausible, one through a manual audit of an implausible held-out score that a model reviewer had let pass [Guo et al., 2026, pp. 21–22], another after an off-by-one date slice had put validation-year labels into training and inflated test Sharpe by 0.3 [Du, 2026, p. 11]. An environment that is not bit-reproducible cannot support a leaderboard, because the same submission scores differently on different machines and no one can check anyone else’s number. And a trajectory that cannot be recomputed from the environment that allegedly produced it is a self-reported claim, not evidence.

We use *evaluation sandbox* to mean a controlled market interface with governed observations, actions, replay, and evidence. The term does not imply operating-system, process, or container isolation; Section 10 states that separate boundary explicitly.

The trading environments agents are trained and evaluated on close none of the three, and the realism they do claim is claimed on a different axis. FinRL states that it replicates the live stock market and cashes that statement out entirely as frictions: transaction cost, liquidity and risk aversion [Liu et al., 2020]. TradeMaster likewise incorporates practical constraints such as transaction costs, slippage and a non-negative balance in the name of a more realistic environment, defers its transition function to FinRL-Meta by citation, and reports backtest metrics where a validation of the simulated market would go [Sun et al., 2023]. ABIDES claims high fidelity in its title and asserts that its platform produces realistic slippage naturally, without quantitative evidence for either [Byrd et al., 2020]; ABIDES-Gym inherits those background-agent configurations without revalidating them and acknowledges only rare unrealistic behavior in the synthetic market [Amrouni et al., 2021]. FinRL-Meta is the honest exception: it names the sim-to-real gap and enumerates what it leaves unmodelled, including permanent market impact of limit orders and a non-resilient limit order book, though it quantifies none of it and tests no distributional property of returns [Liu et al., 2022]. mbt_gym makes no realism claim at all; it claims correctness and demonstrates it by recovering a known closed-form optimum [Jerome et al., 2023], which is verification rather than validation, and is honest precisely because it never conflates the two. The language-model trading literature has the same gap. A 2025 survey observes that the growth of financial language models has outpaced standardized, time-safe benchmarks that connect text to tradable decisions, and that agent frameworks report simulation results on internal market environments, with little standardization of latency, slippage, queue priority or order mix [Fu, 2025, pp. 11–12].

Read together, the pattern is a substitution: frictions stand in for dynamics. A friction bounds what the agent is allowed to do. It carries no information about whether the process the agent acts on behaves like a market, and none at all about whether the trajectory the environment reports is the one it ran. SharpeArena addresses both halves of what the substitution leaves out: the three producer failures above are excluded structurally, and the question of whether its own tape is market-like is asked as a committed diagnostic that the canonical generator currently fails (Section 7.1). Section 9 makes the comparison on the five property axes.

The companion SharpeBench paper [Toca, 2026] supplies the scorer that consumes these trajectories: it argues that ranking trading agents on raw return over short windows measures luck, and its deterministic kernel deflates for search breadth [Bailey and de Prado, 2014], gates on per-run reliability [Yao et al., 2025], and zeroes entries that bypass risk controls. Every one of its guarantees is conditional on the trajectory being honest, which is the problem this paper takes.

SharpeArena excludes each of the three failures at the environment’s own interface, then adds a detection layer behind the structural guarantee for defense in depth. Look-ahead is prevented by the shape of the data layer, which has no API that returns a future bar, and additionally policed by a deny-list guard at the harness boundary; this is an interface guarantee, and Section 10 states what it does not cover, including in-band prediction of the public deterministic generator. Determinism is achieved by restricting the scenario generator’s arithmetic to operations that are exact and identical on every platform, and additionally pinned by committed golden hashes checked in continuous integration on native, WebAssembly and packaged Python surfaces. Trajectory verification records decisions and recomputes performance from them and the frozen engine inputs. A tamper test changes the decisions and requires the replayed run to diverge; fields outside that replay input are not thereby authenticated.

A fourth failure appears only once the agent is a separate process rather than a Python object, which is the arrangement a language-model agent forces. The transports that carry a decision over a wire have no way to report a wire fault through the interface they implement, so a timeout or a dropped connection arrives as an empty decision, and an empty decision against a stateful engine is a hold that persists the position. A wedged agent therefore completes its window and emits a return series that no scoring gate can distinguish from a deliberately conservative one, because every gate consumes returns. Section 4.5 converts a recorded fault into a typed failed cell instead, and Section 5 applies the same rule to the two other places the agentic layer can silently substitute a plausible default for a fact it does not have.

The interface follows the pattern that made Gym the substrate of a decade of reinforcement-learning research [Brockman et al., 2016, Towers et al., 2025]: a small, stable boundary that outlives any single environment behind it. An agent is a program in any language that reads a JSON *Observation* and writes a JSON *Decision*. The contract is versioned, additive-only, schema-published and governed by a written deprecation policy, so a vendor can build against it once. Conform to the contract and you are scorable everywhere the SharpeBench kernel runs. Whether an ecosystem forms around this contract is an adoption question the paper does not answer; Section 10 records that no third party has yet submitted an agent.

Scope and contribution boundary. This paper claims the environment, its agentic layer and its protocol: the interface-level leak-freedom, the cross-runtime determinism, the replay verifiability, the wire contract and the fault typing on it, the strategy-generation, prospective-forecast and forward-execution paths, the seed-band generalization protocol, and the probe layer, together with the calibration experiments that exercise them. The scoring semantics, the deflation prior and the dispersion-unit bug that SharpeBench 0.3.0 fixed are the companion SharpeBench paper’s contributions [Toca, 2026]; F1 appears here only as validation that the producer-scorer separation catches scorer bugs. Two absences are separate and both belong here rather than in a late caveat. No agent is trained in this paper: every F1–F8 experiment uses fixed reference policies, and a learning baseline through the shipped Gymnasium adapter is named future work in Section 10. No current-model field has been completed or admitted either. A superseded engineering pilot using obsolete convenience-sample checkpoints remains committed for protocol auditability, but it is excluded from the paper’s results and from benchmark rank. The agentic layer of Section 5 is therefore claimed as architecture with verifiable guarantees, not as evidence of model performance, and no third party has submitted an agent (Section 10). The empirical artifacts, archived engineering traces and present guarantees are distinct provenance classes (Section A).

Contributions.

1. A point-in-time (PIT) reinforcement-learning environment for trading agents that is leak-free at its interface boundary, cross-runtime byte-deterministic in its Rust core, and decision-replay-verifiable, with readable canonical pre-hash fixtures, a cross-wrapper tape-semantics fingerprint and effective-configuration readback identifying the enforcing code and the boundary of each property.

2. A governed agent wire contract at version 1.0 with JSON Schemas, conformance fixtures and an additive-only evolution policy. It is the boundary an external program, a language model included, crosses to reach the engine, and it carries the fault typing of the next item.
3. Fault typing on the scored path. A transport fault is converted into a typed failed cell rather than the empty-orders hold the transports emit, because against a stateful engine that hold persists the position and produces a return series indistinguishable from a conservative agent's (Section 4.5). The same rule makes an unknown paper submission a representable state rather than a guessed one (Section 5.7), and records the decisions that were refused alongside the ones that executed (Section 5.6).
4. An evidence-admissibility discipline for the agentic layer: a closed falsifiability manifest binding each generated candidate to a hypothesis, a mechanism, invariants and quantitative kill conditions, recorded with a host-counted trial ordinal before validation (Section 5.5); and a strict silver-to-gold promotion path that refuses malformed production traces rather than skipping them (Section 6.6).
5. A procedural scenario generator with disjoint train, gap and held-out seed bands in the Progen tradition [Cobbe et al., 2020], making the generalization gap a measured quantity for any policy submitted to it, though no learner is measured here (Section 10).
6. A task suite that includes a market-making environment shipping the analytical Avellaneda–Stoikov quoting policy [Avellaneda and Stoikov, 2008, Guéant et al., 2013] as a reference, so agents are scored on regret against a fixed analytical baseline rather than against each other.
7. A diagnostic layer that probes realism, manipulation payoff boundaries, adverse-selection mark-outs, population ecology and generator predictability. Two opt-in paths extend it: a concave-impact path that leaves the canonical linear golden path byte-identical, and a sealed-seed derivation that hides the held-out seeds behind a commit–reveal salt while the seed-band structure keeps disjointness from the train band checkable by construction, without the salt.
8. A prospective forecast-evidence protocol that freezes settlement and scoring semantics before submission, records eligible, rejected and late revisions append-only, binds information exposure, and exports an unscored strict artifact for SharpeBench. An executable cross-product tutorial exercises the boundary without a package cycle. A separate Lean model proves selected deadline, revision, binding and fixed-point Brier properties, and executable tests check the same rules against the implementation independently; no mechanism links the model to the code (Section 5.3).
9. The executed experiments and follow-ups of Sections 7 and 8, in which every number is produced by a committed script listed in Section A.

What the evidence shows. Every finding is an instrument calibration on canonical tape that fails the finite-panel-calibrated realism conjunction in 23 of 24 seeded panels, and a predeclared Calm calibration sweep finds no setting that survives its diagnostic, confirmation and volatility rule (Sections 7.1 and 7.1.1). Under the corrected kernel no reference policy is rank-eligible (Section 8.1), yet a common-random-number oracle witness attains eligibility on every tier and finds the per-run reliability gate, which requires every seed's probabilistic Sharpe ratio (PSR) [Bailey and López de Prado, 2012] to clear its bar, last to open at every attained crossing (Section 8.1.1); Section 8.1 states the Sharpe-ratio convention these statistics use. The manipulation, adverse-selection, predictability and ecology probes each report against the environment as often as for it, and Sections 7 and 8 state the scope of each.

2 Design principles

Every decision in the environment traces to one of eight principles.

Leak-free at the interface. The environment owns the time cursor, and the data layer exposes no operation that returns a bar after it. An agent cannot peek through the interface because there is nothing to call, not because a rule forbids it. Detection layers exist, but they sit behind the structural guarantee, never in place of it. The authors of AQuA draw the same lesson from a leak their reviewer

agent approved, moving their guarantee from “the reviewer should catch leakage” to “leakage cannot be written” [Guo et al., 2026, p. 22]. What an agent might infer in-band from the observations the interface does return, given a public deterministic generator, is a separate question, treated in Section 10.

Deterministic across runtimes. A published number must be byte-identical on any platform and any surface. The determinism-critical core is one Rust engine; its scenario transform uses only multiply, add, divide and max, because `ln` and `exp` differ across libm implementations. Golden hashes committed in the source pin the outputs, and the WebAssembly build asserts equality with the native engine rather than assuming it.

Recompute, do not trust. A trajectory is evidence only if a third party can recompute its claimed performance. Replay consumes the recorded decisions together with the frozen scenario, costs, window and seed; it recomputes the run byte-for-byte, and a tampered decision stream produces something else. This is not a signature over every trace field: metadata not consumed by replay needs its own validator. The counterfactual ledger of Section 5.6 extends the same rule from the decisions that were acted on to the ones that were not, because scoring only the acted decisions is a survivorship filter on the agent’s own policy.

A fault is never a decision. When a failure has no representation, some plausible default absorbs it, and the default is indistinguishable from a real answer. A wire fault must not become a hold (Section 4.5), and a submission whose outcome is unknown must not become an accepted or an absent one (Section 5.7). The remedy in both cases is the same: make the failure state representable and typed, so that a downstream consumer must handle it rather than read past it. Two recent studies of model agents met this failure in practice: in one, a token budget cut a reasoning model’s answer mid-generation and the truncated answer read as an abstention [Barbieri et al., 2026, p. 4]; in the other, 14.9% of one league’s tool calls failed at result level while the transport reported success [Barton et al., 2026, p. 14]. The model-agent field of Section 5.1 now records why each model request that returned a completion stopped, so a truncated completion that still parsed is visible in the evidence. This principle is a property of the evidence rather than of the agent, which is why it sits beside determinism and replay instead of in an operations note.

The interface is stable under governance. The agent interface is a tiny JSON surface that evolves additively under a written governance policy, with schemas and conformance fixtures as the authoritative artifacts. An interface that moves under its implementers cannot be built against; stability under governance is what makes the contract worth conforming to.

Generalization is measured, not presumed. Scenarios are procedural functions of an integer seed, in the Progen model [Cobbe et al., 2020], and the protocol fixes disjoint train, gap and held-out seed bands. The fixed-policy experiments use explicitly named subsets of those bands; a gap-band witness set is not described as the canonical held-out namespace. Overfitting becomes one number, the train-minus-test deflated Sharpe, instead of a suspicion.

The environment produces; SharpeBench owns ranking semantics. SharpeArena exposes and uses published SharpeBench scoring and confidence helpers for diagnostics, calibration and adapter metadata. It does not define an alternative benchmark board: field-level deflation, reliability, process gates and rank eligibility remain the SharpeBench contract [Toca, 2026].

Distrust the simulator too. A synthetic market that pays manipulation without bound, that lets makers profit from informed flow, or that emits Gaussian tape will reward the wrong behavior. The probe layer treats the simulator as a subject: a finite-panel-calibrated realism diagnostic against selected stylized facts [Cont, 2001], payoff-boundary sweeps for manipulation, markout decomposition for adverse selection, and ecology runs that ask which strategies survive contact with each other. The probes report findings against the environment as often as for it: the realism diagnostic fails the canonical tape (Section 7.1) and the adverse-selection levels sit on the wrong side of this principle’s own standard (Section 7.3).

The first three principles are properties of committed code, checkable by the listed tests. The empirical claims are properties of committed evidence artifacts and their producer scripts in the same source tree. Their versions, producers and hashes are stated once, in the provenance paragraph of Section A.

3 The environment

SharpeArena’s Cargo workspace has two members. The crate `sharpearena` holds everything determinism-critical: the stepping engine, the wire contract, the scenario generator, the limit-order-book matching engine, mandates and the execution-noise model, built on the published SharpeBench engine crates rather than a vendored copy. The member `sharpearena-wasm` compiles the same kernel to WebAssembly and tests it against the native engine. A third Rust crate, `sharpearena-py`, is explicitly excluded from that workspace and built separately with `maturin`; it binds the engine into a Python package that adds the ecosystem adapters: Gymnasium [Towers et al., 2025], PettingZoo [Terry et al., 2021], Minari [Farama Foundation, 2024], a `verifiers` training environment, and an MCP server. The adapters are thin by policy; anything that must be byte-identical lives in Rust.

3.1 Point-in-time observation, and the guard behind it

The engine owns the time cursor. On each step it emits a `MarketObservation` whose per-symbol snapshot carries a trailing window of closes up to the decision date, with optional fundamentals and news channels derived from that same trailing window. There is no API on the data layer that returns a bar after the cursor, so a policy cannot observe a future bar. Discipline inside generated code does not give the same guarantee: on the causal task category of AlphaQT-Bench, about half of the twelve evaluated models wrote some executable factor code that uses future data [Luo et al., 2026, p. 7]. The guarantee bounds what can be requested, not what can be inferred from the trailing window a request returns (see Section 10). A property test iterates every scenario tier and asserts that no surfaced window extends past the cleared bar and that each window is a suffix of the true history.

One layer out, at the agent-harness boundary, sits `LookaheadGuard`, a refusal layer for tools and policies that try to read what the environment will not give. It carries a deny-list of named operations, each mapped to the reason it is refused: reading the raw dataset (the answer key), slicing past the current index, peeking the next bar, or feeding the scenario seed to the policy as a feature, which would identify the episode and trivialize generalization. Substring tells catch novel operation names, and an index check bounds any point-in-time slice: reading an index strictly greater than the current bar raises. A research escape hatch exists as an explicit environment variable, so relaxing the guard is a recorded decision rather than an accident. The guard is defense in depth behind the structural guarantee: a harness that bypasses it still cannot make the environment leak.

3.2 Observation and action spaces

An implementer sees two JSON documents and nothing else. Table 1 gives the observation, Table 2 the decision; both schemas are draft 2020-12, both close the object with `additionalProperties: false`, and both are committed beside the conformance kit of Section 4.

The observation is flat. `date` is the clock: an ISO-8601 date naming the decision point, and the bound every other field respects. `cash` is the agent’s own cash. `symbols` carries one snapshot per tradable instrument, and `portfolio` the agent’s own holdings, never a peer’s pending state. A snapshot’s `close_history` is the trailing window: the last L closes up to and including date, oldest first, where L is the scenario’s disclosure setting. The default is 20 closes with no optional channels; the shipped presets are a data-poor tier at 3 closes and a data-rich tier at 50 closes plus fundamentals and news. Both optional channels are derived from the surfaced window itself and therefore inherit its cutoff: `fundamentals` is a named map computed from those closes (trailing return, window high, window low) and `news` is one deterministic headline per symbol whose wording is a function of the window’s trailing return. The v1.0 observation carries no mandate field. The mandate is a scenario-level object (Section 3.8) rendered into the episode prompt by the Python task layer, so

Table 1: The MarketObservation wire object at contract v1.0. Optional fields default to empty and are absent from the schema’s required list by construction, which is what makes an additive field backwards-compatible.

Field	Type	Presence	Meaning
<i>Object MarketObservation</i>			
date	string (ISO-8601)	required	decision point; the point-in-time cutoff
cash	number	required	cash available to this agent
symbols	array of SymbolSnapshot	required	one snapshot per tradable instrument
portfolio	array of PositionState	required	this agent’s holdings; empty at a cold start
<i>Object SymbolSnapshot</i>			
symbol	string	required	instrument identifier
close_history	array of number	required	trailing closes through date, oldest first
fundamentals	map string to number	optional	derived fields; empty by default
news	array of string	optional	derived headlines; empty by default
<i>Object PositionState</i>			
symbol	string	required	instrument identifier
shares	number	required	signed holding; negative is a short
avg_price	number	required	displayed average entry price

the fixed reference policies of Section 8 act mandate-blind, which is what the failure taxonomy of Section 8.4 counts.

The action space is a signed target-weight vector over the observed universe. A Decision is a list of per-instrument orders plus an optional free-text reasoning string that the trajectory records. An empty order list is a hold. Within an order, action is a label drawn from a closed four-variant enum (buy, sell, hold, close) and carries no sizing; all sizing is in target_weight, which is a signed target portfolio weight for that symbol. The sign is the shorting semantics: a negative target_weight is a short position, and the signed-weight-short conformance fixture exercises exactly that case, pairing action: sell with target_weight: -0.3 while a sibling order holds at 0.0 and a third buys at 0.5. Adding a fifth action variant is classified as breaking (Section 4) precisely because a consumer that matches the enum exhaustively would misparse it.

Four bounds are asserted by the conformance kit rather than by the schema’s type system: every action lies in the enum, every target_weight is finite, every order symbol is drawn from the universe the observation actually offered, and every confidence lies in [0, 1]. The schema places no numeric bound on target_weight itself; the gross-exposure ceiling an episode is graded against is a property of its sampled mandate (Section 3.8). A legacy decision that omits confidence, rationale and reasoning still parses and is still well-formed, and the legacy-decision-shape fixture asserts it on every commit. Because both schemas close the object, a field the engine emits but the schema omits is an interoperability break rather than a documentation gap, so the conformance kit compares the serialized keys of fully-populated wire values against the schema’s declared properties in both directions, for the observation and the decision alike.

A stated confidence feeds only a calibration diagnostic, and the Python paths count only confidences an agent actually stated. The local field runner records a confidence–outcome pair only for a decision with at least one order that states a confidence, using the mean over those orders, and pairs it with the sign of the next step’s reward. The next step is the right one because the reward booked at a step is the price move on the holdings the previous decision chose, plus that step’s trading cost; a run’s final decision therefore adds no pair. score_run scores a bare return track, which states no confidence, so it reports no calibration Brier score and zero calibration observations rather than scoring a filled-in 0.5 on every bar. Rust paths that run the pinned SharpeBench 0.27.0 simulator directly, run_backtest_checked among them, keep that release’s rule: an omitted confidence reads as 0.5, as the wire default in Table 2 states, and every step, holds included, is paired with the same step’s return. The two rules coexist until a SharpeBench release carries the new one and the pin here moves to it. No table or figure in this paper reports a calibration statistic.

Table 2: The `Decision` wire object at contract v1.0. Optional fields carry the stated default, so a decision written against an earlier revision of the surface still parses.

Field	Type	Presence	Meaning
<i>Object Decision</i>			
<code>orders</code>	array of <code>Order</code>	required	per-instrument instructions; empty is a hold
<code>reasoning</code>	string	optional ("")	rationale captured into the trajectory
<code>cost</code>	<code>DecisionCost</code>	optional (absent)	self-reported spend for this decision
<i>Object Order</i>			
<code>symbol</code>	string	required	identifier; must be an observed symbol
<code>action</code>	enum	required	buy, sell, hold or close; label only
<code>target_weight</code>	number	required	signed target weight; negative is a short
<code>confidence</code>	number	optional (0.5)	stated conviction in $[0, 1]$
<code>rationale</code>	string	optional ("")	per-order rationale
<i>Object DecisionCost</i>			
<code>cost_usd</code>	number	optional (0)	dollar cost of the compute spent
<code>tokens_in</code>	integer	optional (0)	prompt tokens consumed
<code>tokens_out</code>	integer	optional (0)	completion tokens produced
<code>reasoning_tokens</code>	integer	optional (0)	reasoning tokens, not re-added to the total

3.3 Determinism across runtimes

A scenario is a pure function of one 64-bit seed. The generator’s price transform uses only multiply, add, divide and max; it never calls `ln` or `exp`, which differ across libm implementations, so a generated panel is byte-identical on any platform and any runtime. The source commits FNV-1a fingerprints of canonical generated scenarios and of a scripted order sequence’s fill tape. Each fingerprint is preceded by an exact canonical-JSON fixture comparison, so a failure presents readable bytes before an opaque hash mismatch. Native, Python, host-compiled WebAssembly, `wasm32` and the committed npm binary read the relevant fixtures. Tests recompute the hashes on every commit, and WebAssembly parity tests require native-identical returns, traces and scenarios.

A separate `SPEC_HASH` binds the eight version and source files that determine tape semantics, with CRLF normalized before an FNV-1a fold and a manual epoch for semantics outside that file set. The Rust and WebAssembly engines expose the value; the Python and npm wrappers pin it and refuse to load an absent or mismatched engine with a named diagnosis. This handshake turns a wrapper from one semantics revision paired with an engine from another into a startup error. It is a compatibility fingerprint, not a cryptographic commitment. The generator remains deliberately public and deterministic; that makes the hashes checkable and future bars computable from a recovered seed. Section 7.4 measures the boundary: a causal AR adversary finds no undisclosed stressed-tier edge, but a public bounded band of 2^{16} seeds is exhaustively inverted from one observed bar in about one second. Held-out evaluation therefore requires high-entropy unpublished seeds, not merely disjoint public bands.

3.4 Recompute-from-decisions tamper evidence

A rollout trace is a versioned JSONL artifact whose decision sequence is sufficient to recompute performance against the frozen dataset, costs, window and seed. Because the engine is deterministic, replaying those inputs reproduces the run byte-for-byte. The npm package’s test suite exercises the adversarial case directly: it captures a trajectory, doctors every order’s target weight, replays both, and asserts that the tampered decisions cannot reproduce the honest returns. The low-level replay iterator consumes decisions in order; it does not authenticate every auxiliary trace field, including stored step and observation identifiers, and it treats an exhausted decision iterator as holds. Replay therefore verifies claimed performance from decisions and frozen inputs, not the integrity or completeness of arbitrary metadata. Higher-level field and promotion validators impose their own structural checks.

3.5 Requested versus effective configuration

A reproducible label is insufficient if the requested arm never reached the component that consumed it. The current native binding therefore reports the panel dimensions, window, execution seed and FNV-1a fingerprint of the dataset actually installed in an environment. The Python checker reads that block back, independently regenerates the panel requested by the producer and refuses any disagreement. The current F1, F3, F4, F7, predictability, sealed-seed, throughput and witness producers merge those per-seed readbacks into an `effective_config` block when they are run. Most of the committed evidence predates this field and was not regenerated merely to add it; the check protects future regeneration and does not retroactively attest the configuration of an old artifact. F7 is the one finding regenerated since, so its artifact carries the block and the others do not.

3.6 Procedural scenarios and disjoint seed bands

Scenario families follow Progen’s integer-seed-interval model [Cobbe et al., 2020]: Calm, Hard and Extreme volatility-and-jump tiers, a cointegrated-pairs family with a mean-reverting spread, and a regime-shift family that hands a trend regime over to a whipsaw. The Rust core exposes `train_test_split`, which carves a disjoint test family (two separated integer intervals, with the disjointness asserted in tests) from a bounded train interval and refuses an unbounded one, and `cross_regime_split`, which holds the seed band fixed while swapping the regime, a different and generally harder probe than a within-tier split. The evaluation protocol fixes the canonical bands: train seeds $[0, 256)$, an unsampled gap of 10,000 seeds, and a held-out test band $[10256, 10512)$. The fixed-policy F3 and witness producers instead document their own 16-seed subsets: $[0, 16)$ and $[10016, 10032)$; the latter is a separated gap-band evaluation subset, not the canonical held-out namespace.

The Python layer additionally reserves an absolute evaluation namespace starting at seed 10^6 and commits a named held-out regression set inside it, with a disjointness guard asserting the set never touches the train band. All of these bands are public and therefore enumerable; for an adversarial evaluation the recommended mode is the sealed derivation of Section 7.4.1, which keeps the band structure checkable while hiding the concrete seeds behind a salt.

Overfitting is then a measurement. `generalization_gap` scores the same policy on disjoint named bands and reports the differenced deflated Sharpe, and `cross_regime_transfer` reports the in-distribution minus zero-shot out-of-distribution gap over identical seeds, which is zero by construction when the regimes coincide.

3.7 The task suite and the closed-form reference policy

Beyond the canonical position-trading environment, the suite ships a simplex-allocation portfolio task, a VWAP/TWAP implementation-shortfall execution task, and two market models: an endogenous shared-book market where aggregate agent flow moves the cleared price through a Kyle-style linear permanent-impact term [Kyle, 1985] and an Almgren–Chriss temporary-impact term [Almgren and Chriss, 2001], and a deterministic limit-order-book market with integer ticks, price-time priority, a single-price call-auction uncross and a read-only walk-the-book sweep-cost query.

The two market models clear differently, and the distinction bounds what the environment can be read to say. Neither is a uniform-price mechanism. In the endogenous shared-book market every agent decides against one common reference mid for the bar, the exogenous mid scaled by a permanent-impact multiplier that accumulates only strictly earlier flow, so no order moves the price its own submitter saw; the coefficient of the Kyle-style linear permanent-impact term drives that multiplier forward. The term is a stylized linear permanent impact in the sense of Huberman and Stanzl [2004], not the equilibrium coefficient of Kyle [1985]. Each agent’s execution price then differs from the others’, because the Almgren–Chriss temporary term charges it for both the bar’s net flow and its own size [Almgren and Chriss, 2001]. The limit-order-book market is different again. It matches each submitted order against the resting book under price-time priority, walking as many levels as

the size requires, so one step can execute at several prices and a resting order’s queue position is economically real: reducing an order’s size keeps its place in the level’s queue and increasing it re-queues at the back. That book’s single-price call-auction uncross is a separate read-only query, not the path ordinary orders take. What the two models share is discrete time. Orders arrive once per step in a canonical submission order and nothing trades between steps, which puts both closer to a frequent batch auction [Budish et al., 2015] than to a continuous market and leaves latency races and the sniping of stale quotes with no representation in either. No maker-taker fee or rebate is modeled: every Sharpe, regret and markout in Sections 7 and 8 is gross of fees. The restriction is deliberate, it keeps clearing deterministic and byte-reproducible, but it is a scope restriction and results should be read inside it.

In the book, the canonical submission order is the submitting agent’s index, so when two agents rest at the same price in one step the lower index always queues first, and seat assignment becomes a hidden treatment in any multi-agent comparison on the book. The PettingZoo order-book environment keeps that order by default and offers an opt-in rule, `priority="seeded_shuffle"`, that draws a uniform permutation of the seats for each step from the seed and the step. Every pair of seats then queues each way with probability one half. The rule changes only the codes the Python environment submits, so the native engine, `SPEC_HASH` and the golden fill tape are unchanged. Other simulators handle the same ordering explicitly: the market-making study of Fernández Vicente [2025, PDF p. 52] has investors pick at random among identical quotes to avoid a bias from the order in which quotes were placed, and ABIDES-MARL executes simultaneous actions in an explicitly configured order because independent agent interruptions in ABIDES-Gym cannot guarantee a consistent one [Cheridito et al., 2025, p. 4]. A survey of language models in trading lists queue priority among the microstructure details that agent frameworks leave unstandardized [Fu, 2025, p. 12].

The same environment values each agent’s inventory at a mark that excludes the agent’s own orders. By default (`mark="ex_own_mid"`) the mark is the midpoint of the best bid and best ask among resting orders that other controllers own, so an agent’s own resting quotes never enter its own valuation; the earlier whole-book mid remains available as `mark="book_mid"` to replay earlier runs. Each seat is its own controller unless the `controllers` argument names the seats one entrant runs, and seats that share a controller neither value each other’s quotes nor set each other’s mark by trading together. The principle matches CoffeeBench, which values consumed inventory at a weighted-average cost basis so that an agent cannot inflate its score by self-trading at marked-up prices [Sugiura et al., 2026, p. 5]. The mark has limits. An agent’s quotes still move the reference mid that every seat quotes around, so its rivals’ next quotes, and with them its own mark, can follow its quotes. While the other controllers lack a resting order on either side, the mark is the price of the step’s last fill between two different controllers, and it carries its last value only on a step with no such fill. A lone agent’s only counterparty is the noise trader, so it is marked where it last traded and its reward moves only on the steps where it trades. Over seeds 0 to 31 with no inventory penalty, a lone agent that posted its bid one tick below the reference mid and its ask twenty ticks above it, walking the mid for sixty steps before quoting one tick on each side, earns a mean episode reward of 881.4 against 897.1 for its cash plus inventory at the final reference mid, where the frozen mark it replaced paid 2360.3. Under the default penalty, and with the skewed quote held for the whole episode, the mark ranks that agent below a symmetric three-tick quoter over the same seeds, as `mark="book_mid"` does. The frozen mark ranked it above. The book also has no self-trade prevention; a self-trade leaves one agent’s cash and inventory unchanged.

The endogenous market’s impact coefficients (λ, η) are point estimates, and an opt-in elliptic uncertainty set replaces the pair with an ellipse around it, intersected with the nonnegative quadrant, whose two half-widths and correlation are declared modelling inputs rather than estimates the engine validates. The robust entry points clear each bar at the most expensive coefficients inside that set, found in closed form, and given no set they execute the point-estimate arithmetic byte for byte, which the native equality test named in Section 8 pins. The construction is adapted from Ma and Huang [2025], whose elliptic sets are defined around several foci in the space of transition-kernel perturbations (p. 5) rather than over two impact coefficients, and who report that symmetric uncertainty sets often produce

overly conservative strategies (p. 9). A rank-neutral report, `impact_misspecification_gap`, runs one policy alone on each seed twice on the linear kernel, once at the point estimate and once against a declared set. It refuses with a typed error unless both arms traded on the same exogenous path, and it returns per-seed return and per-bar Sharpe under each arm with t-based intervals on their differences. The difference adapts Ma and Huang’s robustness measure, the relative portfolio gap with and without market impact (p. 9), to a point-estimate-versus-worst-case pair, and it does not reduce to a cost of misspecification. Both arms mark the held position at the exogenous mid, so neither return carries its own permanent impact, and each row reports what the engine’s cleared-mid valuation would add as a separate own-impact mark. With a positive λ radius the worst-case coefficient still moves the cleared mids, so the two arms’ weights, order sizes and later fills can differ, and the gap has no guaranteed sign. Only an η -only set guarantees a nonnegative gap, and only for a policy that sends identical weights in both arms while cleared mids stay positive. The two premises are checked rather than assumed. A cleared mid that is not positive refuses the arm with a typed error, and every seed carries `identical_actions`. The report’s `sign_guaranteed` is true only for an η -only set whose policy sent identical weights on every seed. No result in this paper uses the uncertainty set or the report.

The market-making task is the suite’s calibration instrument. It implements the Avellaneda–Stoikov model [Avellaneda and Stoikov, 2008] and ships the model’s closed-form quoting policy beside the environment: reservation price $r = S - q\gamma\sigma^2\tau$ around the mid price S and half-spread $\delta^* = \gamma\sigma^2\tau/2 + \gamma^{-1} \ln(1 + \gamma/\kappa)$, skewed by inventory. This closed form is the source model’s asymptotic approximation, not an exact optimum: the exact treatment of the same problem is given by Guéant et al. [2013], and no proof is offered that the policy is optimal for the implemented reward, which adds an inventory cap, a running inventory penalty and a terminal liquidation charge the source model lacks. We therefore call it the closed-form reference policy. Both the environment and the policy are pure functions of the same frozen parameter set, so the regret metric `mm_regret` runs the candidate and the reference on identical seeds and reports the mean reward gap; the reference scores zero against itself by construction, which fixes the metric’s zero point without claiming the zero point is unbeatable. Identical seeds pair the two arms only while both consume the environment’s random stream identically. Arrivals, fills and the mid step come from one generator per episode, and numpy’s binomial sampler takes a number of draws that depends on the fill probability, so two quoting policies can desynchronize the stream, after which their arrivals and mid paths differ and the gap mixes policy with luck. `mm_regret` therefore records the mid, the two arrival counts and the generator state after every step, and raises `UnpairedMidPathError` at the first episode where any of the three differs between the arms. Mids alone would not show the desynchronization, because at zero volatility, or at a volatility whose step falls below the float resolution of the mid, every mid equals the opening price while the arrivals differ. The default stream is unchanged, and a committed test reproduces every regret of Section 8.2 bit for bit through the checked function, so each of those comparisons drew the same stream. The check does not couple the two arms’ fill draws, which depend on each arm’s own quotes. ABIDES avoids the problem by giving each agent its own generator seeded from one seed, so that an A/B comparison leaves the other agents’ random streams unchanged (as described by Fernández Vicente, 2025, PDF p. 41).

Each step also splits its reward into four parts that sum to it up to rounding: spread capture on the step’s fills, inventory marked to the mid move, the running inventory penalty, and the terminal liquidation cost. `mm_pnl_split` averages that split per policy over the episodes `mm_regret` uses, which separates a quoter that earns spread from one that earns by holding inventory into a favourable path. The split adapts the four-term reward of Fernández Vicente [2025, Eq. 4.1, PDF p. 67]; this environment has no hedging cost, and its running penalty and liquidation charge take the place of that reward’s inventory penalty. Differencing two policies’ splits attributes their regret component by component only when both arms drew the same stream, which `mm_pnl_split` does not itself check; `mm_regret_split` runs both arms through the check above and returns the per-component regret. Neither output feeds a score. The model also assumes fills that do not select against the maker, which sits in deliberate tension with the adverse-selection probe of Section 7.3. Gym-style market-making environments built on the same model family exist [Jerome et al., 2023]; what this task adds is the

shipped reference policy and the regret-not-PnL scoring convention inside the verifiable-producer stack.

3.8 Mandates and deferred claims

Each scenario can sample a trading mandate, and the objective is conditioned on it. A mandate pairs a structural style with up to three optional constraints layered on top of that style: a realized-drawdown cap, a per-bar gross-exposure cap on $\sum_i |w_i|$, and a benchmark symbol to beat. The wire vocabulary names five styles, long-only, market-neutral, momentum, unconstrained and pairs-convergence, and four of them are drawn. Unconstrained is the permissive control, and the two styles that need a short leg, market-neutral and pairs-convergence, are excluded when the environment disallows shorting, so a sampled mandate is never unsatisfiable by construction: `mandate_breach` detects structural violations from the realized weights and returns, and the training rubric penalizes wrong-objective behavior rather than merely not rewarding it. Momentum is the style that is not drawn. It renders “lean into recent winners, cut losers” in its prompt text, and the breach checker sees per-bar portfolio weights and pooled returns and never per-symbol returns, so nothing in the grader can read that constraint off a book. An objective no grader can read must not be the objective an episode is scored against, so the label stays in the vocabulary, because recorded traces carry it and must still parse, and is absent from the draw. A deterministic failure taxonomy classifies every episode into a single worst-case-first disposition, cascade-wiped before bankrupt before stopped-out before mandate breach before clean, and a suite-level rollup tallies the distribution with a schema that always carries every mode. Two further dispositions describe the evidence rather than the behavior: an episode whose returns, events or mandate are malformed or missing is invalid evidence, and one whose event stream records a protocol, transport, agent or infrastructure error is an execution failure. Neither can count as clean, and neither attributes an infrastructure fault to the agent. The mandate checker, the reward schemes, the reward-misspecification controls and the taxonomy read the per-bar weight vectors through one shared event contract, which refuses a weight payload sent under any event name but the canonical one rather than letting the readers disagree about it. A cell that the checked Rust entry point refuses at the wire (Section 4.5) is different again: it produces no episode to classify, carries a typed transport fault instead of a disposition, and is counted apart from the rollup that Section 8.4 reports.

Questions that resolve after the episode get the same structural treatment as look-ahead. A `DeferredDesk` is the only object the agent touches: its entire state is a list of claims and one integer clock advanced by the harness, it holds no dataset and no environment reference, and it computes no score, so there is no code path from a claim to its resolving datum. Resolution is a free function over claims and outcomes that refuses any outcome available at or before the claim’s commit bar and any outcome predating the stated horizon. The agent cannot see the answer at commit time because the object it holds cannot produce one.

3.9 The probe layer: distrusting the simulator

Realism. A stylized-facts battery computes excess kurtosis, absolute-return autocorrelation, gain/loss skew and aggregational Gaussianity in the sense of the classical catalogue [Cont, 2001], plus a timescale-asymmetry statistic in the sense of Zumbach [2009] and an exploratory Fano exceedance-count ratio (the variance-to-mean ratio of the number of large-return exceedances per equal-length block, which is 1 for an IID Poisson-like count and above 1 when exceedances bunch). `certify_realism` gates kurtosis, an ACF value above the fixed 95th-percentile IID null for its exact short-panel estimator, and movement of absolute excess kurtosis toward zero on aggregation. Fano is reported but ungated: it is normalized by its exact conditional-IID finite-panel expectation, yet 12 count windows remain too noisy for a certificate. Certification is reported beside the results rather than required for the leaderboard, and the canonical configuration fails it on 23 of 24 seeded panels (Section 7.1). The intended endgame is a generator revision that passes on disjoint calibration seeds or leaderboard claims conditioned on certification; until then the gate keeps the failure visible.

Manipulation. A crude push-and-unwind schedule, the classic trade-based manipulation of Allen and Gale [1992], is scored against its own zero-impact reference: the live and reference runs share the seed and follower population, and differ only in impact. A boundary sweep varies one parameter and a size-response sweep asks whether payoff is bounded. Linear permanent impact makes the baseline a model-consistency check under the assumptions of Huberman and Stanzl [2004], Gatheral [2010]. On separate entry points, nonlinear permanent impact is $\lambda \operatorname{sign}(Q/V)|Q/V|^\beta$, with dimensionless flow Q/V and a dimensionless coefficient λ for each declared β ; $\beta = 1$ uses the exact original linear arithmetic. This avoids overloading the market-making model’s κ (order-book-depth parameter). Section 7.2.1 evaluates the symmetric schedule, and Section 7.2.2 reports a 135-cell exploratory asymmetric search with selection-adjusted inference. Each probe result now also carries a per-follower seat-removal externality: a follower’s P&L in the live market with the manipulator trading its schedule, minus the same follower’s P&L on the same seed with the manipulator held flat for the whole episode. Every entry is exactly zero when the push weight or the permanent-impact coefficient is zero. The device adapts the competency measurement of EconEvals, which removes the competitor to turn a two-seller pricing game into a single-agent one [Fish et al., 2026, p. 61]; here the manipulator is held flat to measure what its presence does to the other seats. The externality is reported beside the attribution, feeds no rank, gate or boundary, and postdates the committed F5 artifact.

Adverse selection. Informed meta-order flow is routed against a maker roster, with fills struck at the pre-move mid so the maker’s markout carries the informed trader’s edge. The mechanism the probe operationalizes is the classical one [Glosten and Milgrom, 1985], with one deliberate simplification: the price path is exogenous and quotes do not widen in equilibrium response to toxicity, so the probe measures markout accounting against drift, not equilibrium adverse selection. The design is a paired control: the informed leg sides child orders on the alpha, the uninformed leg sides the identical flow on a coin while holding the price path fixed, and the comparison reports mean markout per filled unit at each horizon. If the informed leg does not mark out worse, the module’s own numbers are declared noise.

Ecology. A replicator dynamic in the market-ecology tradition [Farmer, 2002, Lux and Marchesi, 1999] runs a population of strategies over generations: fitness earned in a shared-book market where a species that gains share becomes a larger fraction of the flow everyone else fills against, seat allocation from population shares, extinction below a threshold, an innovation operator that breeds parameter-perturbed variants, and shock schedules that rotate regimes or scale impact to simulate liquidity droughts. The control schedule holds the environment steady, so any claim about a shock has a baseline to beat.

4 The agent contract and its governance

An agent is an external program in any language that reads a `MarketObservation` as JSON and replies with a `Decision` as JSON. The harness sends an observation, the agent returns a decision, repeat. Everything else in this paper exists to make that loop trustworthy.

4.1 Authoritative artifacts

The contract has four authoritative artifacts, all committed. `CONTRACT_VERSION`, currently 1.0, is the single source of truth for the wire version and lives in the contract module of the core crate. Two JSON Schemas, draft 2020-12, define `MarketObservation` and `Decision` for non-Rust implementers; Section 3.2 enumerates their fields. A conformance kit of fixture documents, covering the normal multi-symbol case, a single symbol, an empty-portfolio start, signed-weight shorting and a legacy decision shape, is exercised by a committed conformance test. And a governance document states the evolution rules in writing. The schema closes the JSON shape but cannot know the symbols in the current observation, so one canonical runtime parser additionally enforces episode semantics. The MCP, training and local-agent surfaces all call it; a semantic fault returns without advancing the environment. Outside the two agent-facing schemas, the engine’s output vocabulary that the npm package restates as TypeScript types (baseline agent names, distribution modes, regime

labels, and the field names of the scenario and observation-richness configurations) is generated into `contract/engine-enums.v1.json` by a test of the core crate, and the npm conformance suite checks the TypeScript declarations against it in both directions. The baseline names and regime labels, which cannot carry their own serde representation, are authored once, in the crate's vocabulary module (`BaselineAgent`, `regime_label` and `REGIMES`), which both the WebAssembly export layer and that generator read. That artifact is a drift check on engine output, not part of the governed agent contract.

4.2 Additive-only evolution

The contract evolves additively, and the compatibility that buys is one-directional. A new field that is optional with a default lets a reader accept a message written against an older version, because the absent field takes its default. It does not let a reader accept a message written against a newer one: the published schemas set `additionalProperties: false`, so a consumer validating against the schema it was written for refuses a newer message carrying the new field. The direction therefore decides whether an additive change is transparent. Decisions travel from agent to harness, where the harness is always the newer reader, so an added optional decision field needs no coordination. Observations travel from harness to agent, where the agent is the older reader, so an added optional observation field is not transparent and needs version negotiation, a parallel namespace, or a declared-unvalidated extension envelope, chosen before the field is added rather than after. The optional fields are absent from the schemas' `required` lists by construction. Removing, renaming or retyping a field, making an optional field required, and removing or renaming an action variant are all defined as breaking and forbidden on the v1 surface. A breaking change never mutates the v1 types in place; it ships a parallel namespace with its own schemas and its own major version, and the two run side by side through a published deprecation window. Adding a new action variant is itself classified as breaking, because consumers that exhaustively match would misparse it, so it also goes through the parallel namespace rather than an additive bump.

4.3 Version semantics

`CONTRACT_VERSION` versions the wire shape only and is deliberately decoupled from the crate and package versions, which move on their own release cadence. A major bump is a breaking change shipped as a parallel namespace. A minor bump is reserved for a batch of additive fields significant enough to advertise. A single additive decision field needs no bump, because the harness is the newer reader and existing agents are unaffected. A single additive observation field does need one, because an agent validating against the older closed schema refuses it (Section 4.2), and it still needs the negotiation, namespace or envelope choice made before the field is added. A package release never, on its own, bumps the contract version.

4.4 Strict inputs and typed native-boundary errors

The JSON Schemas already close the public observation and decision objects. Current Rust and WebAssembly deserializers extend that rule to selected caller-controlled environment, mandate, market, richness and task configurations with `deny_unknown_fields`; a misspelled knob is an error rather than an ignored field that silently selects a default. Result objects remain forward-open so an older consumer can relay output from a newer engine. At the Python native boundary, five stable codes distinguish invalid arguments, invalid JSON, invalid sealed salts, unavailable data and engine failure. They map to typed subclasses of `SharpeArenaError`, itself a `ValueError` for compatibility. An unknown future code falls back to the base class while preserving its `[CODE]` prefix, so a stale wrapper reports what the engine said instead of relabeling it.

4.5 A transport fault is a failed cell, not a hold

An external agent speaks over a wire, and a wire fails. The v1.0 `Decision` carries no error variant, and the external transports cannot signal one through the `Agent` trait they implement, so a timeout, a

dropped connection or an unparseable reply degrades to an empty order list while the fault itself is recorded into a rolling `TransportHealth`. Section 3.2 fixes what an empty order list means: it is a hold, and against a stateful engine a hold is not inaction. The position persists. An agent whose wire wedges on the first bar therefore rides its opening position to the end of the window, the run completes without raising, and the return series it emits is the return series of a deliberately conservative agent. No downstream gate separates the two, because the deflation, reliability and process legs all consume returns. A committed characterization test states the failure rather than describing it: a wedged agent run through the plain engine call yields a scoreable `Run`.

The health record mitigates this failure only when a scored path reads it. The crate therefore re-exports the diagnostics types, and `run_backtest_checked` is the explicit Rust entry point for an external agent: it runs the cell, reads the health afterwards, and returns a `CellOutcome` that is either `Scored(Run)` or `Failed(TransportFault)`. The local-field scheduler and artifact compiler enforce the same distinction in their own typed cell schema, so their faults cannot enter a scored artifact as flat returns. The low-level compatibility function `run_backtest` remains public and degrades an external transport fault to the transport's empty-decision fallback; the type system does not force a wire consumer to choose the checked entry point. The guarantee therefore covers the checked and repository scoring paths, not every possible caller of the compatibility API. A typed failure carries the retryable transport-fault count, the agent-protocol fault count, whether the per-endpoint circuit breaker tripped, and the last decision-level error. Its accessor returns `None` on a failed cell rather than a default-valued `Run`, so a caller that chose the checked path cannot score a failure by accident. Tests preserve both sides of the boundary: a wedged agent through the compatibility call produces a scoreable run, whereas the checked call yields a failed cell; clean and malformed transports exercise the other outcomes.

The fix belongs with the determinism and replay guarantees rather than with operations, because it is the same kind of claim. Leak-freedom bounds what the agent could see and replay bounds what the trajectory could be; neither says anything about a return series the harness synthesized on the agent's behalf. Masking is the one failure mode that produces evidence a verifier cannot reject, since the doctored quantity originates inside the honest producer and replays byte-for-byte. A failure is evidence, and never a hold.

4.6 Why the contract can be trusted without trusting the host

The governance promise is social; the properties of Sections 3.1, 3.3 and 3.4 are technical, and together they remove the host from the trust equation for everything the byte-identity guarantee covers. That coverage has a boundary. The Rust core, its `WebAssembly` build and the scenario bytes returned through the Python bindings are pinned to the byte by committed FNV-1a goldens, and the `WebAssembly` crate asserts byte-identity with the native engine for stepping and scenario generation. The native and Python suites assert all seven scenario fingerprints. Three committed golden files under `crates/sharpearena/contract/attestation/` are each read by three gates: the native suite, a `wasm-pack` run that executes the exported entry points inside a `WebAssembly` module under Node, and the npm package's own suite, which asserts the committed `.wasm` binary it ships rather than a host build of the same source. `scenario-goldens.json` records the canonical and clustered scenarios; `backtest-goldens.json` records two `run_baseline` and two `replay_run` entries on the execution and replay path; `kernel-goldens.json` records three `walk_forward`, two `stress_suite` and four `tag_regime` entries. Each entry carries its exact kernel input and a committed pre-hash fixture. A generator, backtest or regime-tagging change therefore cannot reach a release through one runtime alone. The Python evidence layer behind the remaining probes is deterministic in its seeds but sits outside the golden-hash path (Section 10). Within the covered surface, a leaderboard entry can be recomputed from the frozen seed and submitted decisions on any tested platform. SharpeArena exposes the published scoring helpers used by its diagnostics, while the SharpeBench kernel independently defines board-level deflation, reliability, process gates and rank eligibility [Toca, 2026, Bailey and de Prado, 2014].

5 The agentic layer

The contract of Section 4 defines a loop. This section describes what the environment puts around that loop when the program on the other side of it is a language model: a batched model-agent field, a strategy-generation path that counts its own search, a ledger of the decisions that were not acted on, and a forward paper-execution arm. Everything here is architecture and its guarantees. No agent result is reported from any of it, and Section 10 states that absence as a limitation rather than leaving it to be inferred.

5.1 Model agents and the scorer boundary

The model-agent path keeps environment execution and field judgment separate. SharpeArena owns the point-in-time observation loop, canonical decision parser, execution, process trace and append-only cell journal. The `sharpearena-compile-bench` command validates that a journal contains the complete predeclared Cartesian field, with one consistent completion per cell and matching return hashes, then emits an ordinary SharpeBench submission artifact for each dataset. SharpeBench owns the field-level deflation, bootstrap, process and reliability judgments. The implementation dependency remains directed: SharpeArena uses the small published SharpeBench protocol, simulator and kernel crates, while SharpeBench does not import the full environment package. The products are therefore operationally interdependent without a package cycle.

Model output crosses one canonical boundary. The published `Decision` schema constrains syntax, and a host-side parser separately enforces episode semantics: symbols must have been observed, each symbol may occur once, and weights must be finite and within the configured bound. The signed target is authoritative; the action field remains an audit label because whether a target buys, sells, closes or holds depends on the current position. Every model and training path calls that same parser. The scheduler applies the same rule as the Rust gate of Section 4.5 directly: a malformed response, timeout or transport failure becomes a typed failed cell rather than a hold or a flat return series. The field record binds the model artifact and runtime version, sampling and inference seed, cadence, context and reasoning budgets, prompt/scaffold digest, raw response and its hash, token-count availability, retry count, latency and exact process events. It also records the converter, decoding and parser stack, cache and parallelism settings, and runtime identity; a backend fact that cannot be observed is labeled `unresolved` rather than `reconstructed`. Because a completion the token budget cut off can still parse into a decision, the record also keeps why each successful model request stopped, in one of four classes: `stop` and `length` for exactly those backend spellings, `other` for any other value, and `absent` when the backend reports none, which is never read as a clean stop. A completion the parser rejected keeps its class in the cell’s failure record. The SharpeBench bridge refuses a record whose per-request classes and per-class counts disagree, and publishes per-class totals in the operational profile it writes beside the score submission, with a separate `unrecorded` count for requests in records written before the field existed. The totals are rank-neutral.

The shipped scheduler batches the native vector environment, assigns stable cell ordinals over model, scenario, seed and repetition, supports deterministic sharding, and resumes through append-only journals. A plan-embedded `csv_path` is constrained to a relative path that resolves inside the plan directory; absolute paths, parent escapes and symlinks resolving outside that root are refused before their contents can enter evidence. This is input-file containment, not process isolation. No current-model result is reported in this paper.

5.2 Where the byte-identity guarantee stops

The producer/scorer boundary of Section 6.5 has two parties. An agentic field has three, because the model provider is neither the environment nor the kernel, and this is the exact point at which the paper’s central guarantee ends. Environment replay stays deterministic: the same seed produces the same scenario bytes, and the same decisions replay to the same returns on any tested platform. Model inference does not. Continuous batching, server version, sampling configuration and hardware all move the sampled token sequence, so the paper treats inference as a measured source of agent

variance rather than describing it as byte-identical. What the field record can pin is everything around the sampling step: which checkpoint, which server, which seed, which prompt digest, which raw response hash. What it cannot pin is that the same checkpoint and seed will re-emit that response. Environment bytes and replay remain fixed while agent outputs vary, and the evidence schema keeps the two apart rather than averaging over the distinction.

The configured inference service sits inside the trusted computing base. Containment for an executable third-party entrant is a different problem and is not this artifact's; Section 10 states whose it is and what has been verified about it, which is less than the word sandbox usually implies.

5.3 Prospective forecasts as an unscored producer contract

Historical replay cannot remove pretraining contamination from a model that may already have seen the target period. The FutureX authors make the same argument for forecasting benchmarks: past data carries logical leakage and retrieval contamination, so evaluating prospectively is probably the only methodologically sound way [Zeng et al., 2025, p. 2]. The prospective path therefore commits a question before its outcome exists. A `ForecastContract` freezes the instrument and target, forecast kind, observation source, opening and submission deadline, resolution boundary, unit, exact open and close definitions, boundary ownership, missing-data and fallback policies, and scoring rule. Its canonical JSON digest binds every revision to those bytes. The contract admits six forecast forms (point, binary probability, categorical distribution, Normal predictive distribution, direction outside a frozen neutral band, and prediction interval) and names the scoring rule for each before submission. Directional accuracy, which is not a proper score [Gneiting and Raftery, 2007], is named explicitly rather than grouped with the proper ones. The scores themselves, and the conditions under which inference over them is reported, belong to the consumer and are specified in the companion paper [Toca, 2026].

The `ForecastLedger` is append-only. An idempotency key may replay only byte-equivalent content, a claim may never move to a different contract, and a revision must name the prior ordinal and give a reason. A submission before the opening time is recorded as rejected; one after the closed deadline is recorded as late; neither replaces the latest eligible revision. Each attempt records its information exposure, and a consensus-visible attempt must bind the consensus snapshot it observed. Resolution covers every claim exactly once as pending, resolved, cancelled or rejected. A resolved outcome must have become available after the effective revision and no earlier than the frozen resolution boundary. The producer validates the outcome's shape but writes no score.

Since SharpeArena 0.25.0 the producer exports the ledger as `sharpe.forecast-evidence.v2`: the v1 envelope plus, on every revision, a declaration of the canonical encoding its contract digest was computed under. The archived pilot field under `paper/evidence/` remains v1. The document has a closed shape and strict finite JSON, and is written through an exclusive temporary file, synchronized and atomically replaced, with a parent-directory synchronization on POSIX. SharpeBench 0.19.0 and later reads both envelopes without importing SharpeArena, verifies a v2 digest under its declared encoding only, validates the contract and revision graph again, and scores the forecasts on a report-only surface that cannot alter the trading-rank projection [Toca, 2026]. The executable tutorial under `examples/forecast-quality` exercises this producer-to-consumer boundary end to end (Section A).

Lean supplies a small independent model of the load-bearing rules. It proves that eligibility is exactly the closed interval from opening through deadline, that late and rejected revisions cannot replace the effective ordinal, that every accepted repeated binding preserves the original contract digest, and that truthful reporting uniquely minimizes fixed-point expected binary Brier loss at positive scale. These are proofs about the abstract integer model, not an extraction of Python semantics. Separate executable tests check the same rules against the Python implementation, including a test that restates the fixed-point Brier model and compares it with the floating-point rule. Nothing mechanically links the two: no test reads a Lean declaration and no proof reads the code, so agreement between model and implementation rests on the two being written to the same rules. No current model

result is admitted: the repository retains one superseded convenience-sample lifecycle pilot for audit transparency, but its obsolete checkpoints and scores are excluded from this paper’s empirical conclusions.

5.4 Observed strategy search

The strategy-generation path measures the search performed inside the harness. A model emits a closed JSON language over bounded price, moving-average, momentum, RSI and volatility conditions; it cannot emit Python, shell commands or arbitrary tool calls. The host records the raw generation response and counts every candidate before validation or deduplication. Accepted candidates are compared on an explicit validation interval under a deterministic median-score rule with a fixed tie-break; only the selected candidate touches the disjoint test interval. That observed raw count supplies the deflation trial count for this search, which is the count Section 6.5 asks an entrant to declare honestly and here does not have to take on trust. It does not reveal candidates generated before entry, so it narrows rather than closes the declared-trials problem. It also counts within one search: a journal can hold several searches against the same test interval, and each record reports only its own trials.

The test interval stays out of the loop because a model loop can overfit whatever it is scored on. In one 2026 alpha-mining study, an iterative single-model baseline raised its training information coefficient from 0.038 to 0.053 over ten rounds while its median holdout Sharpe per alpha was -0.676 , which the authors describe as learning to game the training metric [Li et al., 2026, p. 19]. Gençay [2026] builds a close relative of this path for strategy discovery: a validated tool surface whose trial ledger records every strategy evaluation the search performs, with an invalid construction rejected by validation rather than evaluated (p. 4). SharpeArena assigns its ordinal before validation, so its count also includes the candidates validation rejects.

5.5 Binding a candidate to a falsifiable claim

A candidate that arrives with a rule and a backtest number invites the prose that follows selection. The hypothesis, the mechanism and the conditions under which the edge would be abandoned all get written after the winner is known, and are shaped by knowing it. The *EdgeManifest* forecloses that ordering by making the falsifiability terms part of the candidate, under a closed schema, recorded before validation and before deduplication.

A manifest binds a candidate to a hypothesis and a mechanism, each required to be at least eight characters so that an empty string is not a manifest, together with the regimes and instruments the claim is made for, at least one invariant, at least one kill condition, and a verification plan. Every condition applies a comparator from a closed set to a threshold that carries an explicit unit from a ten-member enumeration, so a bare number is a parse error and a string such as "2%" is rejected with the instruction to write the number and put the unit in the unit field. A missing required field invalidates the candidate; nothing is defaulted, and an unknown field is refused alongside a missing one. An empty invariant list is refused on the stated ground that it asserts nothing and can therefore never fail, and one condition identifier may not appear in both the invariant list and the kill list. The verification plan refuses a confirmation split equal to the selection split, and requires at least two observations before it will conclude anything.

The ledger assigns a candidate its trial ordinal before validation and before deduplication, so the count it publishes is the count of candidates the model emitted rather than the count that survived. An invalid candidate is recorded with the reason it was invalid, a repeat is recorded with the ordinal it duplicates, and neither is dropped. The summary labels the count with the stage it was taken at, so a downstream deflation prior consumes the pre-filter number and not a filtered one. A ledger whose trial ordinals do not match its own row count fails to load rather than loading partially.

Kill conditions are evaluated only outside the selection sample: the monitor refuses any sample not named as test, forward or holdout. It returns one of four verdicts under a precedence fixed in code. A fired kill condition retires the edge; failing that, an unresolved check makes the whole report

indeterminate; failing that, a violated invariant reports a violation; otherwise the edge is healthy. A metric that is absent, that arrives in a unit other than its threshold's, or that is not a finite number resolves to unresolved rather than to a pass, so a measurement that was never taken cannot be read as an edge that survived.

5.6 The counterfactual ledger

A journal of executed orders records what the risk guard let through and says nothing about what it stopped, so the distance between what the agent proposed and what the venue received is not a measurable quantity. Scoring only the acted decisions is a survivorship filter, and the thing it filters is the agent's own policy. The counterfactual ledger records every decision, acted on or not, with one ghost order per intended order carrying the intended quantity, the quantity that actually executed, a reference price, and one disposition from a closed five-member set: executed, risk refused, below minimum notional, not submitted, or resized.

The record refuses to contradict itself. A disposition of executed whose intended and executed quantities differ raises; a disposition other than executed whose quantities agree raises with the complaint that it claims a difference that is not there; an executed quantity above the intended one raises; and the flag saying the decision was acted on must agree with whether any quantity reached the broker. Over a window the ledger reports a selection gap: how many decisions were acted on, intended against executed notional, intended against executed profit and loss under supplied settlement prices, and the difference between them, with a disposition histogram that always carries all five keys so a zero is printed rather than omitted. A positive foregone figure says the execution path cost money and a negative one says it avoided a loss. Neither is a claim about the market; both are claims about the gap the harness introduced.

The forward paper arm instantiates the ledger in memory by default and may persist it separately. The historical local-field scheduler does not yet instantiate it: that path submits target weights to the native engine, whose step result exposes process events and the resulting portfolio but not a per-order intended-versus-filled quantity record. Reconstructing those quantities from successive portfolios would conflate slippage, participation caps and carried fills, so this paper does not describe that reconstruction as a ledger. The deterministic scenario makes such a ledger recomputable once the native step surface exposes the required fill facts; that integration remains work rather than a claimed default.

5.7 Forward paper execution and the unknown submission

Forward paper trading is a different evidence class. Read-only market data enters a trusted supervisor; the model sees observations but never credentials; a fixed risk guard preflights the order batch against a shadow account before an in-memory broker or an adapter pinned to Alpaca's paper endpoint can receive it, and that adapter refuses any other origin and follows no redirect. The preflight stops at the first refusal, because a remote paper broker does not update the caller's account between requests and the verdicts after that point would be computed against a state the broker never reached; the orders it never assessed are recorded as not submitted, which Section 5.6 keeps distinct from refused. The journal records the market snapshot, decision, proposed orders, risk outcome and paper-broker response, and a separate private preimage supports the forward commitment. There is no real-capital endpoint and no override. Provider time, provider data and simulated fills make this evidence non-replayable by construction, and every record the forward journal writes is stamped as non-deterministic with its replay guarantee named as none, so a forward result may extend the forward record but may not inherit the byte-identity claim of the historical environment.

The state that matters in a forward arm is the one a backtest never has. A submission whose response never arrives is neither accepted nor absent, and the harness cannot know which without asking the broker. The order lifecycle therefore makes `submission_unknown` a representable state rather than a comment. It is reachable only from `submitted`, and it leaves only to `reconciled_accepted` or `reconciled_absent`. A server-side error or a transport failure on submit puts the lifecycle

there, while a broker that answers with a client-side rejection is a rejection and not an unknown. Acknowledgment latency stays unset until a real acknowledgment arrives, so an order that reaches its terminal state through reconciliation does not acquire a fabricated one.

Reconciliation asks the broker by a client order identifier derived from the agent identity, the model digest, the parsed decision, the market snapshot and the order index, so the same decision addresses the same order on every attempt and a retry cannot create a second order by accident. A replacement is authorized only after the broker confirms absence, and only once: the lifecycle refuses a second. A query that cannot be answered raises rather than resolving, so a network failure is never recorded as a confirmed absence, and the lifecycle stays in the unknown state where no replacement is authorized. State is persisted before the submit call, again on the transition to unknown, and again after every reconciliation, each write going to a temporary file that is then atomically renamed, so an order left unresolved by a crash is still unresolved after the restart rather than quietly absent. Account state is reconciled against the broker rather than read from the plan file and left stale. Every transition and every raw broker acknowledgment hash reaches the forward evidence journal, which stays distinct from the deterministic backtest evidence.

Two consequences of this design are limitations rather than features, and Section 10 states both: partial fills accumulate by last write rather than by summing deltas, and a broker that cannot answer a reconciliation query halts the session instead of retrying.

6 Benchmarking protocol

This section is written for someone who wants to evaluate an agent in SharpeArena rather than for someone reading our results. It fixes the seed bands, the seed counts, the seed-custody mode, the scoring path and the verification step. The rules below are the ones our own producers follow, stated prescriptively; where a producer departs from them, we say so and say why.

6.1 Seed bands: what to train on and what to evaluate on

Train on the bounded band $[0, 256)$. Evaluate on the held-out band $[10256, 10512)$. The unsampled gap of 10,000 seeds between them is part of the protocol, not a formatting artifact: adjacent integer seeds are cheap to enumerate from a published train interval, and the gap makes an accidental overlap visible as an out-of-band seed rather than as a plausible one. `train_test_split` carves that structure and refuses an unbounded train interval; the disjointness is asserted in the crate's tests rather than left to the caller.

Two bands that appear in this paper are *not* the canonical held-out namespace, and results reported on them should not be described as held-out results. The fixed-policy F3 producer uses a 16-seed train subset $[0, 16)$ and a 16-seed gap-band evaluation subset $[10016, 10032)$; the rank-eligibility witness of Section 8.1.1 uses the same gap-band subset and, separately, the F1 table band. Both subsets are separated from the train band, and neither is $[10256, 10512)$. We use them because the fixed reference policies have nothing to overfit, so the smaller subsets cost nothing in interpretation and buy tractable runtimes for a probe repeated under five noise paths per cell. A submitted agent has parameters and therefore does not get that exemption: report it on the full canonical band.

Report the cross-regime transfer matrix alongside the within-tier gap. A within-tier gap is structurally blind to the quantity Section 8.3 measures, namely how much score a regime switch moves for a policy selected on one tier, and both metrics run from the same committed scripts. The Python surface additionally reserves an absolute evaluation namespace starting at seed 10^6 with a committed held-out regression set inside it and a disjointness guard against the train band; use that namespace for regression testing a harness, not for reporting agent scores.

6.2 Seed counts

Run all 256 seeds of the held-out band on each tier you report. The reason is resolution, and this paper measures it directly. At 16 seeds, a per-seed pass rate near one half carries a nominal 95% half-width of about 0.24, and a difference between two such rates about 0.35 (Section 8.1); a bounded, heavily deflated score bootstrapped over 16 seeds produces intervals that span most of $[0, 1]$ (Section 8.3); and on the market-making task, 16 paired episodes detect a mean regret of about 12 reward units at 80% power and a two-sided 5% level, roughly a third of the regret a badly miscalibrated quoter incurs (Section 8.2). Sixteen seeds resolve the shape of a curve and the sign of a collapse. They do not resolve a ranking between two nearby agents, and a leaderboard built on them would be reporting seed luck.

State the dispersion convention beside every number. We report three: seed-resampled percentile bootstrap intervals for deflated-Sharpe quantities, t-based intervals over committed per-episode vectors for paired quantities, and plain counts where the quantity is a count. Whichever convention you use, serialize the per-seed or per-episode vector behind the interval, as the producers under `paper/src/` do, so a reader can recompute it. Any claim narrower than the reported interval is outside the design’s reach and should not be made.

6.3 Seed custody, and what is forbidden

Every band named above is public and therefore enumerable. Section 7.4 measures what that costs: a scan of a public 2^{16} band recovers the true seed from a single observed bar in about a second, and a recovered seed makes every future bar of the episode exactly computable. For any evaluation whose result is meant to survive an adversarial reading, use the sealed derivation of Section 7.4.1: publish the salt hash before entries arrive, withhold the salt during scoring, and reveal it afterwards so anyone can recompute the named seeds and replay the run. The current Rust API accepts only a `SealedSalt`, whose sole constructor refuses fewer than 16 bytes, whose Debug output is redacted, and which has no serialization, display or dereference implementation. The Python binding uses the same constructor and maps a refusal to the typed `InvalidSalt` error; the WebAssembly package exposes no sealed-seed entry point. The band structure stays checkable without the salt, so disjointness from the train band remains verifiable by construction.

Any public bounded-band evaluation should be treated as compromised until it documents sealed seed custody, a pre-entry commitment, a non-reuse policy and a reveal deadline. This is a statement about the protocol, not an accusation about any particular run: the scan is cheap enough that assuming it was not performed is unreasonable. Other evaluation environments guard held-out answers against the same kind of probing; OpenFinGym, for example, rate-limits the submission API of its verifier to prevent reverse engineering of the held-out ground truth [Zhang et al., 2026, p. 7].

Three prohibitions bind a submitted agent, and each has an enforcing layer rather than a promise. Reaching past the cursor is forbidden: reading the raw dataset, slicing past the current index and peeking the next bar are refused by `LookaheadGuard` at the harness boundary and are structurally unavailable at the environment’s own surface (Section 3.1). Feeding the scenario seed to the policy as a feature is forbidden, because it identifies the episode and makes the generalization protocol meaningless. Enumerating or reconstructing the evaluation seeds, by band scan or otherwise, is forbidden, and sealed derivation is what makes that prohibition enforceable rather than honour-based. Relaxing the guard is possible for research use through an explicit environment variable, which records the decision; a run that sets it is not a protocol-conformant evaluation.

6.4 Leakage channels in one place

The controls above and in Sections 3 and 5 close different channels, and each leaves a different residue. Table 3 lists them together, an inventory of the kind Li et al. [2026, Appendix A.4, p. 30] give for their own system, where each audited path by which test values could reach the model is

Table 3: Leakage channels, the control SharpeArena applies to each, and the exposure that remains. Entries restate properties described in the cited sections.

Channel	Control	Remaining exposure
A bar past the cursor, requested through the environment	The data layer has no operation that returns one (Section 3.1)	None at the interface; inference from the bars it does return is the fourth row
Optional observation channels	Fundamentals and news are computed from the surfaced trailing window and inherit its cutoff (Section 3.2)	As the fourth row
Raw dataset, index slicing or next-bar peeks by harness tools	LookaheadGuard deny-list, substring tells and index check (Section 3.1)	Side channels the environment does not own, such as the clock, filesystem or network (Section 10); the explicit research override (Section 6.3)
Inverting the public generator from observed bars	Opt-in sealed seed derivation with a salt hash published before entry and the salt revealed after scoring (Sections 6.3 and 7.4.1)	A public bounded band falls to a one-bar scan (Section 7.4); tape-level inversion is untested; the sealed evidence is a simulated workflow
The scenario seed fed to the policy as a feature	Forbidden by the protocol and refused by the guard (Section 6.3)	A harness that bypasses the guard
Repeated looks at a test interval during strategy search	Selection on a validation interval; only the selected candidate touches test; candidates counted before validation (Section 5.4)	Candidates generated before entry; searches that reuse one test interval are counted per search, not together
A model that has already seen the outcome it forecasts	Forecast contracts frozen before the outcome exists, with information exposure recorded per attempt (Section 5.3)	Historical replay cannot remove pretraining exposure; no current-model field is admitted (Section 10)

listed with how it was sealed. Every entry restates a property established in the section it cites, so the table adds no result.

6.5 Scoring and verification

SharpeArena exposes published SharpeBench scoring and confidence helpers, and its calibration scripts call them; an entrant does not supply a scorer. SharpeBench owns the authoritative board semantics: deflated Sharpe with the declared trial count [Bailey and de Prado, 2014], a per-run reliability gate, and the process gates that zero entries bypassing risk controls [Toca, 2026]. Rank eligibility is the conjunction of those legs, not any one of them, and Section 8.1 shows why: deflation alone crowns drift on the calm tier, and the per-run reliability gate alone misses the search breadth deflation discounts. Declare the trial count honestly. Section 8.3 contains a worked instance of what happens otherwise: the same policy on the same sixteen seeds scores 0.1231 or 0.1114 depending only on how many trials the deflation prior is told to consume. The count covers every configuration whose reward was inspected before the entry was chosen: each training run and random seed, each hyperparameter setting and reward shaping, and each checkpoint evaluated to pick the one submitted. The repository’s evaluation guide asks entrants to declare it that way, and the SharpeBench bridge carries the number into scoring as `in_sample_trials`. Both kinds of search appear in reinforcement-learning market making. Three identically configured deep Q-learning market makers trained in one market ended with mean results of 129, −45 and −5 thousand dollars, and the network transferred to a later experiment was the best of ten evaluated snapshots (Fernández Vicente, 2025, PDF pp. 52, 55–56, 58). A run chosen that way is one draw from a search, and its deflation should count the search.

Submit the decision trajectory, not self-reported returns. The replay verifier regenerates the scenario from the declared seed and recomputes the run from the ordered decisions and frozen execution inputs (`replay_submission` in the Rust core, `replayRun` in the npm package). A mismatch is evidence that the decisions or declared engine inputs do not reproduce the reported run, and the npm package’s test suite exercises the decision-tampering case directly. This guarantee is narrower than full artifact authentication: the low-level replay path does not validate every auxiliary metadata field or prove that the decision list has the length another field claims. Field compilation and trace-promotion validation enforce those separate structural requirements.

6.6 Growing the case set: strict silver-to-gold trace promotion

An evaluation suite written by the people who wrote the environment tests what those people already thought of. The alternative is to grow it from traces the system actually produced, and the hazard there is subtler than a bad case: a pipeline that skips whatever it cannot parse silently selects for the traces that happened to be easy to read, and the resulting suite is biased by the parser rather than by the operator. The promotion path refuses instead of skipping.

A malformed or incomplete trace raises rather than being passed over. A blank line, a line that is not JSON, a record that is not an object, a step record appearing after the metadata record, a missing observation, a missing decision, a missing reward (a missing reward is not zero), a non-finite reward, a step ordinal out of sequence, a declared step count that disagrees with the steps present, or a trace shorter than two steps all abort the load. Eight metadata fields are required and none of them is defaulted.

A trace is identified by a fingerprint over six components: environment, model, scaffold, contract, data and process. The data component digests the dataset hash together with the sorted scenario seeds. The process component digests the sequence of process events, recording per step the ordinal, the order count, the sorted action labels, the sorted environment events and whether the episode terminated. It is derived from actions rather than rewards, so two runs that arrived at the same outcome by different routes do not share a fingerprint.

Deterministic checks run before anything else, and four of the five block: rewards finite, no look-ahead key present in the observation, a structured decision, and seeds declared in the metadata. The fifth, that the reward series varies at all, only warns, because a flat series is a suspicion rather than a defect. Only a failing check produces a silver candidate, and what is frozen is not the transcript. The scenario is minimized by a greedy two-sided shrink to the smallest window on which the check still fails, floored at two steps, with the per-step auxiliary payload stripped to four keys; what is stored beside it is the check that triggered the promotion and the hash of the source trace. Silver rows are immutable: re-appending an identical row is a no-op, and appending different content under the same identifier raises.

Promotion from silver to gold requires a recorded operator decision that names the operator, names the candidate it promotes and carries a rationale of at least twenty characters. A gold case is a minimal scenario plus the invariant it is expected to satisfy, and it is evaluated with sockets disabled, so a case that reaches for the network fails rather than succeeding against a live service. This is the same admissibility discipline as seed custody, applied to the case set instead of to the seeds, and it is what keeps the forward and generated-strategy evidence of Section 5 from being pooled with the deterministic tables of Sections 7 and 8 on the strength of looking similar.

6.7 Compute cost

The protocol above is cheap, and we measure that rather than assert it. Table 4 reports throughput for the three shipped surfaces on the canonical 4×120 panel, measured on one AMD Ryzen Threadripper PRO 5975WX (Zen 3) workstation under Windows 11, as in the companion paper’s compute description, by the single command in Section A (the throughput artifact records the host’s processor identifier, AMD64 family 25 model 8 stepping 2, and its Windows 11 build rather than a model name), which times the Python surface in process and shells out to the other two. The three surfaces do not run identical workloads, and the table says which is which. The scenario-generation column is the same function emitting the same canonical JSON bytes on all three, so those three numbers are directly comparable. The stepping columns are not. The native row steps a scalar `TradingEnv` through the Rust API with no serialization per step. The Python row steps the `Gymnasium` adapter, which crosses the binding and encodes a `JSON Decision` on every step, and pays a factor of about 52 for it. The `WebAssembly` row runs a whole baseline episode inside the kernel with one boundary crossing per episode, so its per-step figure measures the kernel rather than the interface.

Table 4: Measured throughput per surface on a 4×120 panel, median of three repeats of 500 episodes each. Scenario generation produces the canonical JSON serialization on every surface and is the one directly comparable column. Steps per second and episode wall time reflect each surface’s own stepping path, described in the text: the native and Python rows step 120 bars per episode, the WebAssembly row runs a 100-bar baseline episode inside the kernel.

Surface	Stepping path	Scenarios/s	Steps/s	Episode (ms)
Native Rust	scalar TradingEnv through the Rust API, no per-step serialization	50,031	747,468	0.161
WebAssembly	baseline episode inside the kernel, one boundary crossing per episode	29,211	508,127	0.197
Python	Gymnasium adapter, JSON Decision encoded on every step	44,026	14,240	8.43

The number that matters for a submission is the cost of one protocol-conformant evaluation. Running the full six-policy reference field over the canonical held-out band on all three tiers, 4,608 episodes and 552,960 steps, takes 302 s of wall clock on the Python surface, the slowest of the three and the one an entrant is most likely to use. That is roughly 1,830 steps per second rather than the 14,240 of the stepping row, because the evaluation figure also carries the reference policies’ own computation and the SharpeBench kernel’s scoring. A single agent on a single tier is 256 of those 4,608 episodes, one eighteenth of the field; a single agent across all three tiers is one sixth. Both are episode-count ratios, and because policies differ in their own per-step cost they do not promise a proportional share of the wall clock. A learning baseline is not blocked by the environment either: at the Python surface’s measured stepping rate, ten million environment steps cost about 12 min of environment time, and at the native rate about 13 s, so the compute constraint Section 10 names belongs to the learner rather than to the simulator. Throughput is hardware-dependent, and it is the one class of number in this paper a replicator should expect to differ.

7 Instrument validation

A simulator paper that reports a stylized-facts battery it passes has not, in this literature, shown that its simulator is market-like. Cont [2001] says so in the founding paper: the stylized facts are “usually formulated in terms of qualitative properties of asset returns and may not be precise enough to distinguish among different parametric models.” The objection has been restated at every subsequent methodological turn. The four-level empirical-validity scheme of Axtell and Epstein [1994] (see also Barde and van der Hoog, 2017), as surveyed by Fagiolo et al. [2019], runs from a Level 0 caricature through qualitative and then quantitative agreement with macro structures to Level 3 quantitative agreement with micro structures; Fagiolo et al. [2019] observe that publishing standards now require at least Level 1, and survey the techniques built to “better discriminate among different ABMs reproducing the same set of stylized facts.” Lamperti [2018] is blunter, that “moments are poor representations of time series’ behaviour”: two series can agree on every stylized fact and every moment while their dynamics diverge point for point. Twenty-four years after Cont, LOB-Bench restates the objection for machine learning, that prior work “relied on qualitative comparisons of stylized facts, making rigorous model comparisons infeasible and hindering research progress” [Nagy et al., 2025]. Recent simulator papers still grade stylized facts this way. One marks conformity with a Hill tail exponent of “approximately three” by check marks, accepting 2.82 and rejecting 3.23, without a stated tolerance [Hashimoto et al., 2026, pp. 15, 22]; another compares simulators on excess kurtosis and volatility of volatility with arrows that mark larger values as better [Zou et al., 2026, p. 10]; a third reads volatility clustering from a lag-1 absolute-return autocorrelation of 0.18 in one standard configuration [Jayakody and Jayakody, 2026, p. 11]. None of the three states a null distribution for the statistic it grades. The standard failure mode of this genre is therefore not a simulator that fails its realism check. It is a simulator that passes one and is degenerate anyway, because the check cannot tell models apart.

Read against that, the headline result of this section is evidence about the instrument as well as a shortfall in the generator. A public, committed, deterministic generator fails a calibrated three-check conjunction on 23 of 24 seeded panels, against an IID null fixed in source and tuned to no market panel, and the failure is reproducible from a command in Section A by anyone who disbelieves it. No pass rate that would have counted as a realism claim was declared before the gate was run, so the result cannot be read as a test the generator could have passed on a stated criterion. What the failure does show is narrower: this gate can refuse this generator. It does not establish the gate’s discriminating power in general, which would need generators known to be market-like and known not to be, graded against a threshold fixed in advance. We therefore report the result the way Vyetenko et al. [2020] and Nagy et al. [2025] report theirs: we name the fact that fails (absolute-return autocorrelation, on all three tiers), name the configurations that fail it (the canonical generator everywhere, and the Calm tier on excess kurtosis besides), report the opt-in driver’s relative improvement without promoting it to a realism claim, and locate the contribution in the measuring instrument rather than in the thing measured. Vyetenko et al. [2020] set the register in one sentence: “there are still significant discrepancies between simulated markets and real ones. However, this work serves as a benchmark against which we can measure future improvement.” Platt and Gebbie [2018] reach the same conclusion from the other direction, that “traditional stylized fact-centric validation seems unable to detect these potential problems, suggesting that such methods of validation are simply not sufficient for the validation of intraday ABMs employing realistic matching procedures.” Their conclusion is scoped to intraday agent-based models with realistic matching, not to a daily-bar procedural generator like this one, so we cite it for the general point rather than as a finding about this tape. On the four-level scheme above, the canonical tape sits at Level 0 on the axis this section gates and makes no claim above it.

Validation runs on four axes, and the section is organized by them, because the environment fails the first and passes structural checks on the other three. A single verdict about “realism” would misreport all four. Section 7.1 is *unconditional* realism, whether the marginal and autocorrelation structure of the tape resembles the classical catalogue. Section 7.2 is *response* realism, whether the price responds to flow the way a no-manipulation model requires. Section 7.3 is *mechanism* realism, whether informed flow costs a maker what the microstructure literature says it should. Section 7.4 is *adversarial* realism, whether a public deterministic generator withstands an adversary who attacks the generator rather than the interface. Within each axis, every follow-up unit sits directly after the finding it extends.

F4 grades the canonical position-trading generator; F5 and F6 run on the market-making and market-clearing models, which are separate simulators with their own structural and positive-control checks and are not submitted to F4, as Section 10 states. Dispersion follows three conventions: t-based 95% intervals over committed per-seed or per-episode vectors for F5 and F6, count-based pass rates for F4, and across-seed standard deviations for the predictability probe, each stated where it is used. Every entry point, seed set and quantity is named where it is reported, the numbers come from JSON and figures committed with their producer, and the commands are in Section A.

7.1 Unconditional realism: stylized-facts certification (F4)

For each tier we grade price panels from eight seeded episodes with `certify_realism` against three directional bounds from the classical catalogue: positive excess kurtosis, absolute-return autocorrelation above a fixed 95th-percentile IID null for this exact 120×4 , ten-lag estimator, and aggregation moving *absolute* excess kurtosis toward zero [Cont, 2001]. The IID null is a deterministic 20,000-panel simulation fixed in source and tuned to no market panel. Gain/loss skew [Cont, 2001], timescale asymmetry [Zumbach, 2009], and a finite-panel-calibrated Fano exceedance-count ratio are informational. The Fano ratio conditions on the observed exceedance count and divides the Bessel-corrected sample Fano by its exact equal-block IID expectation; it remains ungated because 12 count windows make it too noisy for certification.

Table 5 and Figure 1 give a mixed verdict. Calm is platykurtic (mean excess kurtosis -0.97), although its returns move closer to Gaussian in absolute-kurtosis distance on aggregation (mean $+0.57$), and

Table 5: Stylized-facts results over 8 seeds. Fano is a conditional-IID-calibrated exploratory ratio and has no pass column; the final rate is the conjunction of the three calibrated directional checks, including aggregation toward zero absolute excess kurtosis. Aggregation passes 8/8 seeds in Calm, 8/8 in Hard, and 8/8 in Extreme.

Tier	Excess kurtosis		$ r $ autocorr		Fano proxy	Gated facts
	mean	pass	mean	pass	mean	pass rate
Calm	-0.97	0/8	+0.0002	1/8	1.32	0.000
Hard	+16.30	8/8	-0.0107	0/8	0.93	0.000
Extreme	+11.38	8/8	-0.0118	1/8	0.91	0.125

fails the conjunction on all eight seeds. Hard and Extreme pass excess kurtosis and aggregation on all eight seeds, but neither clears the calibrated ACF threshold consistently: Hard passes 0 of 8 and Extreme 1 of 8. The canonical generator therefore passes 1 of 24 seeded panels, and the single binding fact across all three tiers is volatility clustering as the ACF measures it. The same split appears in an endogenous multi-agent order-book simulator built on ABIDES, whose return distributions all reject normality while Engle’s ARCH-LM test finds no significant volatility clustering [Cheridito et al., 2025, p. 20]. The Fano ratios are observations and do not decide certification.

The opt-in volatility-clustering driver (`vol_clustering`) adds deterministic EMA persistence to the same seeded panel while preserving the zero-knob golden path. Under the calibrated ACF rule at strength 0.5, the three-check pass rates rise to 0/8, 8/8 and 4/8 for Calm, Hard and Extreme, so the driver moves Hard from failing to passing and Extreme halfway. That is a relative improvement in a named opt-in configuration, and we state it as one: the certified path is the zero-knob canonical generator, and it still fails. The opt-in jump-burst knobs and the exploratory Fano ratio support no remediation claim.

7.1.1 Calibrating Calm with the existing knobs: a negative result

The obvious next question is whether the shipped knobs can be tuned into a Calm tier that passes, so we ask it under a rule fixed before the sweep. We sweep the existing opt-in families, `vol_clustering` in $\{0, 0.3, 0.5\}$ and the jump-burst probability, persistence and size grid, for 99 configurations. The protocol is fixed in the producer: a setting must pass the three-check conjunction on at least 7 of 8 diagnostic seeds, have mean realized-volatility ratio at most 1.25 versus default Calm, and pass at least 7 of 8 disjoint confirmation seeds (100–107); among survivors it would minimize normalized RMS return perturbation. The ACF check uses a one-sided 5% IID-null test for the exact 120×4 , ten-lag estimator rather than the raw condition “ $ACF > 0$ ”. No cell satisfies all three predeclared steps, so there is no certified Calm preset and no report-band estimate to promote.

The closest low-volatility candidate is `vol_clustering` = 0.5, burst probability 0.02, persistence 0 and size 0.015. It passes 7/8 diagnostic panels but only 5/8 confirmation panels at a mean volatility ratio 1.150. Several settings reach 7/8 on both bands only by exceeding the 1.25 volatility constraint; the strongest near candidate (clustering 0.5, probability 0.02, persistence 0.5, size 0.015) has 8/8 and 7/8 passes but ratio 1.289. We retain the former as a reproducible `calm_calibration_candidate` in the Rust and Python surfaces, pinned by a golden fingerprint, precisely so a future generator revision can be compared against it; it is not a certified configuration (Figure 2).

7.2 Response realism: the manipulation boundary and size response (F5)

Everything in this subsection is a search over a finite grid of sampled schedules and coefficients. Sampled means are estimates at sampled cells; a negative cell is not a proof over the schedules between the cells, and a positive cell is not an optimum. We state that once here and report cell values plainly below.

We score a push-and-unwind manipulator against its zero-impact paired reference (identical seed and schedule, $\lambda = \eta = 0$) over eight seeds, so `impact_pnl` attributes profit specifically to having moved




Figure 1: Stylized-factor statistics for the canonical zero-knob tiers only, one panel per statistic, each with its own vertical axis. Points are per-seed values over the eight seeded panels of each tier (Calm circles, Hard squares, Extreme triangles), and black bars are tier means. The three gated statistics carry their inclusive rule in the panel title: excess kurtosis ≥ 0 , absolute-return autocorrelation ≥ 0.0146 (the 95th percentile of the fixed IID null for this estimator, 0.014556 before rounding) and aggregational Gaussianity ≥ 0 . In those panels a dashed line marks the bound, the shaded region marks the passing side, and the count under each tier is the number of its 8 seeds that pass. The kurtosis and aggregation panels use a symmetric-logarithmic vertical axis that is linear within ± 1 . The Fano ratio, gain/loss skew and timescale asymmetry are reported but not gated, and their panel titles say so; thin grey lines mark their neutral values, 1 for the Fano ratio and 0 for skew and asymmetry.

the price. `impact_boundary_sweep` sweeps permanent impact (λ , the coefficient of a Kyle-style linear permanent-impact term), temporary impact (η) and the momentum-follower gain over six grid points per axis; `size_response` sweeps the push weight over five points. The evidence commits per-seed impact-PnL vectors at every grid point; intervals below are t-based 95% CIs over the eight seeds.

Under the precise linear, time-independent permanent-impact assumptions studied by Huberman and Stanzl [2004] and the related no-dynamic-arbitrage conditions of Gatheral [2010], round-trip price manipulation is ruled out, so the linear sweep is a model-consistency check on the environment rather than a discovery procedure. The three units that follow test what a linear symmetric sweep structurally cannot reach.

Table 6 and Figures 3a and 3b are consistent with the linear model’s no-manipulation conditions on the sampled grid. On every axis the pump loses money at every sampled coefficient, with every per-point 95% CI excluding zero, including the realistic center of the grid ($\lambda = 0.1$, $\eta = 0.05$: impact




Figure 2: Calm calibration sweep. Each point is one of 99 configurations, placed by its normalized tape perturbation and its number of passing diagnostic seeds out of 8; marker shape encodes the `vol_clustering` setting, and the legend sits above the axes. Filled points satisfy the 1.25 volatility constraint and hollow grey points exceed it. The dashed line labelled “rule: ≥ 7 of 8” is the diagnostic threshold. No configuration meets the full diagnostic–confirmation–volatility rule under the finite-panel-calibrated ACF gate.

PnL -3.4×10^{-4} , CI $[-4.3, -2.6] \times 10^{-4}$). The size response is bounded and decreasing: the least-bad size is the smallest sampled push weight and the loss grows to -5.1×10^{-4} at full size.

7.2.1 Normalized-flow concavity: a falsifiable symmetric probe

Empirical impact estimates are concave in traded quantity rather than linear [Almgren et al., 2005], so a probe that can only express linear impact cannot return a positive even in principle. To make it capable of one, the engine supplies an opt-in permanent-impact exponent on dimensionless crowd flow $x = Q/V$:

$$I_\beta(Q; V) = \lambda \operatorname{sign}(x) |x|^\beta, \quad x = Q/V.$$


The temporary own-size term remains $\eta q_i/V$. This normalization is load-bearing: applying a power to raw shares and dividing by V afterwards would change units with β and make results depend on the arbitrary share denomination. A regression test rescales capital, flow and V together and requires the nonlinear impact fraction to remain unchanged. At $\beta = 1$ the engine uses the original expression $(\lambda Q + \eta q_i)/V$ exactly, preserving every linear golden bit; only $\beta \neq 1$ enters the `powf` path outside the cross-runtime golden guarantee. At fixed λ and sub-unit $|x|$, lowering β also raises the impact magnitude relative to the linear term, so this is a falsifiability stress test rather than a scale-matched comparison of impact shapes. The exponent acts on each bar’s flow, and the resulting increments compound in the permanent multiplier and never decay, so it makes the per-bar term concave without reproducing the concave, decaying impact of a whole meta-order (Section 10).

We rerun the symmetric push-and-unwind sweep at $\beta \in \{0.5, 0.7\}$, eight seeds per point, with paired zero-impact references and pointwise t intervals. No sampled mean is positive. At the canonical calibration, impact-attributed P&L is -8.9×10^{-4} (95% CI $[-19.6, 1.8] \times 10^{-4}$) for $\beta = 0.5$,

Table 6: Manipulation probe, impact-attributable PnL as fraction of capital per episode, in units of 10^{-4} , mean over 8 seeds with t-based 95% CIs. Top: endpoints of each swept axis (6 grid points per axis); every axis is unprofitable everywhere, so no profitability boundary exists to report. Bottom: the full size-response grid in the push weight.

Axis	Low end		High end	Boundary
Permanent λ (0 to 0.8)	-1.3	-31.7	$[-44.3, -19.2]$	none
η (0 to 0.8)	-2.2	-22.4	$[-23.5, -21.4]$	none
Follower gain (0 to 120)	-2.6	-6.7	$[-10.1, -3.3]$	none

Push wt.	PnL	95% CI
0.10	-0.14	$[-0.25, -0.03]$
0.25	-0.50	$[-0.77, -0.24]$
0.50	-1.53	$[-2.05, -1.00]$
0.75	-3.07	$[-3.84, -2.30]$
1.00	-5.14	$[-6.16, -4.12]$



(a) Boundary sweeps, one panel per axis on a shared vertical scale.

(b) Size response in the push weight.

Figure 3: The manipulation probe, mean impact-attributed PnL over 8 seeds (intervals are in Table 6). (a) Impact-attributed PnL is negative at every sampled coefficient on every axis, so no profitability boundary is drawn. (b) The loss is bounded and decreasing in size; the dashed line marks the least-bad push, which is the smallest one sampled.

-13.5×10^{-4} (CI $[-21.2, -5.9] \times 10^{-4}$) for $\beta = 0.7$, and -3.4×10^{-4} (CI $[-4.3, -2.6] \times 10^{-4}$) for the linear arm. For $\beta = 0.5$, 7 of 23 stored sweep slots have intervals crossing zero; repeated canonical slots reduce these to five distinct configurations. At $\beta = 0.7$, 0 of 23 cross zero (Figure 4).

The null identifies the schedule’s limitation rather than closing the economic question: profitable constructions under nonlinear permanent impact can be asymmetric, and the next unit searches such shapes.




Figure 4: Impact-attributed P&L for the sampled symmetric schedule under linear and normalized-flow concave impact, against the permanent-impact coefficient (left) and the push weight (right). Lines are means and shading is the pointwise 95% interval; exponents differ by colour, marker and line style. The two panels have independent vertical scales, so their magnitudes are not comparable by eye.

7.2.2 Exploratory asymmetric-schedule positive control

Concavity makes slicing matter. For positive dimensionless flow $x = Q/V$, the sum of the n per-slice impact increments before the engine’s multiplicative price compounding is

$$nI_\beta(Q/n; V) = n^{1-\beta}I_\beta(Q; V),$$

so a slowly accumulated leg and a block liquidation need not cancel when $\beta < 1$. The identity is about the impact function, not the compounded terminal price, so it motivates a positive control without guaranteeing that any finite schedule profits after compounding, temporary costs and follower flow.

We cross five up/down duration ratios, three block fractions, three exponents and three market arms: 135 cells, each on eight seeds. All displayed cell intervals are pointwise. Because the reported winner is selected after seeing the grid, its headline inference uses a two-sided Bonferroni familywise 95% interval over all 135 cells (Student t , 7 degrees of freedom); the evidence stores the family size, critical value and adjusted interval.

The best observed cell is the no-follower, zero-temporary-impact arm at $\beta = 0.5$: nine bars of uniform accumulation followed by one-bar liquidation yield $+21.8 \times 10^{-4}$. Its pointwise 95% interval is $[17.4, 26.2] \times 10^{-4}$ and its familywise interval remains positive at $[10.0, 33.6] \times 10^{-4}$. The same sampled schedule with temporary impact yields $+18.4 \times 10^{-4}$, with familywise interval $[6.4, 30.4] \times 10^{-4}$. The canonical follower arm has no pointwise-positive cell. Three granularities of the same count are worth separating: of the 135 cells, four have familywise-positive intervals (the 9:1 and 8:2 uniform schedules on each of the two no-follower arms, $+21.8$, $+10.9$, $+18.4$ and $+8.9 \times 10^{-4}$); they occupy two arms; and the two reported winners are the 9:1 cell on each arm. The fourth of those cells, temporary-impact 8:2, clears zero only at the boundary: its familywise lower bound is $+0.04 \times 10^{-4}$, so a reader recomputing it at a slightly different critical value would count three rather than four. At $\beta = 0.7$ no sampled mean is positive. Among the 45 sampled linear cells, every mean is negative; the 1:9 and 9:1 theory-arm values are numerically near-symmetric rather than bit-identical (Figure 5).

We confirm the selected theory-arm cell independently on 32 fresh seeds fixed before that confirmation run: mean $+26.3 \times 10^{-4}$, two-sided 95% t interval $[23.7, 28.9] \times 10^{-4}$ (df 31). This is one fixed-cell interval and carries no multiplicity correction for the exploratory search, so the 135-cell Bonferroni interval answers the selected-grid question while the fresh-band interval answers the one-cell replication question. Detection is itself a finite-sample problem. In an Almgren–Chriss




Figure 5: Exploratory asymmetric-schedule search. Shading shows pointwise 95% intervals; selection-adjusted intervals for the two selected $\beta = 0.5$ winners are reported in the text. Four of the 135 cells retain positive familywise intervals, spread over two sampled arms.

model whose permanent impact has the same signed power form (p. 5), Tsaknaki et al. [2026, p. 18] find that at the volatility they calibrate to an average meta-order, none of their methods, the full-information optimizer included, produces a statistically significant positive terminal cash, and the permanent-impact exponent they estimate from 500 simulated meta-orders is 0.99 with standard error 1.88 against a true 0.1 (p. 12). What the positive control establishes is that this instrument can detect a profitable sampled round trip under normalized concave impact. What it leaves open is the shape of the profitable region: the search fixes push size, λ , V , total leg length and most follower gains, and omits short-side, overshooting and longer-hold trips. Conversely, the absence of a positive among 45 sampled linear cells is consistent with the cited linear-impact model under its assumptions.

7.2.3 Extended sweeps through the positive-control winner

The positive control leaves six axes at their anchor values, so we sweep each of them one at a time through its no-follower, zero-temporary-impact, uniform 9:1 winner: push weight, λ , total leg length, hold length, short-side trips and follower gain, at $\beta \in \{1, 0.5\}$, five values per axis and eight common seeds, for 60 cells. Because the anchor was selected in the earlier 135-cell grid on these same seeds, the familywise intervals are two-sided Bonferroni 95% t intervals over the global 195-cell family, the antecedent grid plus this extension, rather than over 60 cells alone. That correction counts every evaluated cell, but it does not give simultaneous coverage here: the 60 extension cells were chosen around a winner selected from the first grid on the same eight seeds, so the hypotheses they test are data-adaptive, and a Bonferroni bound over the cells evaluated does not account for how they were chosen. We therefore read the extension's familywise intervals as descriptive. The independent




Figure 6: Extended manipulation sweeps after per-leg uniform normalization. Shading is the pointwise 95% interval; error bars are the global 195-cell Bonferroni familywise 95% interval, which is descriptive for these data-adaptively chosen cells rather than simultaneous coverage. The displayed maxima are sampled cells, not optimized values.

evidence remains the fixed-cell confirmation on 32 fresh seeds (Section 7.2.2), which covers the original 9:1 cell and none of the extension cells. Unequal “uniform” legs use separate $1/n_{\text{up}}$ and $1/n_{\text{down}}$ increments, and tests verify that each leg executes exactly one signed peak notional, so changing duration no longer silently changes the amount liquidated.

The results are sharply asymmetric by exponent. Across the 60 extension cells, 23 global-familywise intervals are positive, 31 are negative and 6 cross zero. The longest tested 45:5 theory-arm trip is the largest observed cell, $+135.0 \times 10^{-4}$ with global-familywise interval $[122.1, 148.0] \times 10^{-4}$. Interactions between axes, other exponents, other flow scales, transient impact and overshooting round trips remain untested (Figure 6).

7.3 Mechanism realism: adverse-selection markout (F6)

`compare_informed_vs_uninformed` runs twenty-four paired episodes in which the identical meta-order flow is sided on the alpha in one leg and on a coin flip in the other, holding schedule, sizes and price path fixed, and reports mean maker markout per filled unit at horizons of 1, 5 and 20 bars. Markout is measured in the tape’s absolute price units (the mid starts at 100) per unit of filled quantity; the evidence commits per-episode vectors, and the gap CIs below are t-based 95% over the twenty-four paired episodes.

Two quantities must be separated here: the sign of the *paired estimated gap* and the sign of the *levels*, and they come out on opposite sides.

The paired estimated gap is positive at every horizon (Table 7, Figure 7): makers filled by informed flow mark out worse than the paired-uninformed control at $h = 1, 5$ and 20, the per-unit gap widens from 0.072 to 0.436 as the information realizes, and every horizon’s paired CI excludes zero, including the smallest ($h = 1$: $[0.052, 0.094]$ over per-episode paired gaps). The loss is in the drift after the fill rather than in the fill itself, which is the signature shape of adverse selection. The claim is scoped to the paired estimated gap, the differential cost of informed over uninformed flow, and not to maker profitability levels.

Table 7: Maker markout per filled unit (price units per unit quantity) against informed versus paired-uninformed flow, 24 paired episodes. The point estimates pool markout and filled quantity over all episodes, so they are quantity-weighted; the gap CIs are t-based 95% intervals over the 24 per-episode paired gaps and are centred on the per-episode mean gap (0.073, 0.305 and 0.441; Table 8), not on the pooled gap printed beside them. The gap is the price of trading against information and grows with horizon; the levels stay positive at every horizon, a separate finding discussed in the text.

Horizon (bars)	Informed	Uninformed	Gap	Gap 95% CI
1	0.689	0.761	0.072	[0.052, 0.094]
5	0.489	0.791	0.302	[0.228, 0.383]
20	0.386	0.823	0.436	[0.303, 0.579]

The level result goes against the environment, and we report it the way we report F4. Makers earn a *positive* mean markout against informed flow at every horizon (0.689, 0.489, 0.386 per unit), which is on the wrong side of the Section 2 standard that a realistic market should not let makers profit from informed flow. Three calibration choices produce it: fills are struck at the pre-move mid, so every fill books the full quoted depth as spread capture; the default quote depths are wide relative to the drift the $\alpha = 2.0$ meta-orders realize over 20 bars; and the price path is exogenous, so quotes never widen in equilibrium response to toxicity as they would in Glosten and Milgrom [1985]. At this calibration the probe therefore measures markout accounting against a fixed drift, in which informed flow costs makers relative to uninformed flow but not in absolute terms. Recalibrating the informed alpha or striking fills at post-impact prices so the level turns negative, and adding a variant where quotes widen with realized toxicity, are the named follow-ups; the first of them is the subsection below.

The single-episode maker decomposition illustrates the exact identity on one seeded episode (seed 0); it is an illustration, not an aggregate finding. The wider quoter (maker_0, 135 fills) keeps a positive markout at $h = 20$ (0.526 per unit, toxic-fill rate 0.296): the flow is not selecting against it. The tighter quoter (maker_1, 92 fills) shows the subsidized pattern: aggregate markout is still positive at $h = 20$ (0.197 per unit) but its meta-order fills alone lose 0.197 per unit on 62% of its filled quantity, with a toxic-fill rate of 0.533 at $h = 20$, so noise flow is paying for the informed flow. Spread capture and adverse drift decompose exactly ($\text{spread_capture} + \text{adverse_drift} = \text{markout}$, in price units times quantity): maker_1 captures 82.8 of spread and gives back 64.7 to adverse drift by $h = 20$.

7.3.1 The endogenous arm: maker markout when the flow moves the price

The design above fixes the efficient-price path, so it measures the informed-versus-uninformed *gap* but cannot speak to maker profitability *levels*: a maker whose fills never move the price against it books the full quoted depth on every print, and the positive levels reported above are partly that artifact. This arm makes the level question answerable. We rerun the same twenty-four paired episodes with the cleared price endogenous: the taker flow actually filled against the maker ladder moves the mid through the permanent-impact law of the shared-book clearing engine of Section 3. Let S_t be the efficient price (the meta-orders' alpha increments plus ordinary volatility, identical across both legs and both arms), Q_t the signed taker quantity filled in bar t (meta-order children plus noise flow, buys positive), λ the coefficient of the engine's Kyle-style linear permanent-impact term [Kyle, 1985] and V the engine's ADV-like volume normalizer. Makers then quote around

$$\text{mid}_t = S_t M_t, \quad M_{t+1} = M_t (1 + \lambda Q_t / V), \quad M_0 = 1, \quad (1)$$

the engine's running multiplier, updated after the bar clears so the reference price a maker quotes against is moved only by flow strictly before it. Temporary impact is not a separate term: the maker ladder itself is the temporary component (a taker pays the quoted depth, which is what the engine's Almgren–Chriss η term models in the maker-less market [Almgren and Chriss, 2001]), so the arm adds exactly one mechanism to the exogenous scenario, that a filled unit moves the next reference mid by $\text{mid}_t \lambda / V$ in the taker's direction. At $\lambda = 0$ the arm reproduces the exogenous scenario bit for bit, and a cross-check in the test suite feeds the identical flow sequence to the native clearing




Figure 7: Markout per filled unit by horizon: informed flow versus the paired-uninformed control.

Table 8: Maker markout per filled unit (price units per unit quantity) under the exogenous and endogenous price paths, 24 paired episodes, horizons in bars. Every cell is the mean over per-episode values (episode markout over episode filled quantity, all makers; toxic rate is the share of all maker fills marking out negative), with t-based 95% CIs, $df = 23$; this is the convention of the committed per-episode vectors and of the gap CIs in Table 7, so the exogenous rows differ from the pooled quantity-weighted aggregates of Table 7 in the third decimal. The evidence also serializes paired endogenous-minus-exogenous changes and their CIs. Endogenous arm at $\lambda = 0.1$, $V = 240$.

Path	h	Informed	Uninformed	Gap	Toxic-fill rate (inf / unin)
exogenous	1	0.689 [0.683, 0.695]	0.762 [0.740, 0.783]	0.073 [0.052, 0.094]	0.016
exogenous	5	0.488 [0.470, 0.507]	0.794 [0.720, 0.867]	0.305 [0.228, 0.383]	0.221
exogenous	20	0.385 [0.341, 0.430]	0.827 [0.711, 0.942]	0.441 [0.303, 0.579]	0.381
endogenous	1	0.622 [0.614, 0.629]	0.693 [0.672, 0.715]	0.072 [0.051, 0.093]	0.034
endogenous	5	0.352 [0.327, 0.376]	0.652 [0.579, 0.724]	0.300 [0.225, 0.375]	0.323
endogenous	20	0.240 [0.192, 0.289]	0.670 [0.558, 0.782]	0.430 [0.294, 0.565]	0.437

engine and matches its cleared mids to relative 10^{-12} , so Equation (1) is the engine's law and not a paraphrase of it. The calibration sets $V = 240$, one episode's realized taker volume (mean 239 filled units over the 48 exogenous-arm episodes), and the engine's default $\lambda = 0.1$, so a filled unit displaces the mid by 0.042 price units at the initial mid of 100. A parent order fills a mean 22.4 of its 60 offered units, and the stated law applied to that fill alone gives $22.4 \lambda / V = 0.933\%$. The measured quantity is different: `informed_displacement`, the parent-signed change in the log impact multiplier over the parent's own window, averages 0.81% (t-based 95% CI [0.77, 0.85]% over episodes, sided with the parent in both legs, and zero by construction in the exogenous arm). That window also carries the noise flow filled in the same bars, and its multiplier compounds bar by bar, so the two numbers describe different quantities and neither is derived from the other.

Both effects the design predicts appear (Table 8, Figure 8). First, the gap moves by at most 0.011 per unit: 0.072, 0.300 and 0.430 at $h = 1, 5$ and 20 against 0.073, 0.305 and 0.441 in the exogenous arm, every CI excluding zero. Permanent impact is direction-agnostic from the maker's seat, since

Table 9: Endogenous-arm markout per filled unit at $h = 20$ as the permanent-impact coefficient sweeps, $V = 240$ throughout, the same 24 paired episodes, and exploratory pointwise t-based 95% CIs (df = 23); no sweep-wide multiplicity claim is made. “Impact/unit” is the mid displacement per filled unit at the initial mid, $100 \lambda/V$. The $\lambda = 0$ row is the exogenous arm.

λ	Impact/unit	Informed	Uninformed	Gap	Toxic (inf)
0.00	0.000	+0.385 [+0.341, +0.430]	+0.827 [+0.711, +0.942]	0.441 [0.303, 0.579]	0.381
0.05	0.021	+0.313 [+0.267, +0.359]	+0.748 [+0.635, +0.862]	0.435 [0.299, 0.572]	0.415
0.10	0.042	+0.240 [+0.192, +0.289]	+0.670 [+0.558, +0.782]	0.430 [0.294, 0.565]	0.437
0.20	0.083	+0.095 [+0.039, +0.152]	+0.513 [+0.403, +0.623]	0.418 [0.282, 0.554]	0.479
0.30	0.125	−0.050 [−0.117, +0.017]	+0.356 [+0.245, +0.467]	0.406 [0.268, 0.544]	0.507
0.40	0.167	−0.196 [−0.275, −0.117]	+0.198 [+0.084, +0.312]	0.393 [0.251, 0.536]	0.525

the maker is on the wrong side of the move its own fill causes whichever way the taker trades, so it cancels out of the informed-minus-uninformed difference; the gap is the information, and the arm confirms it is not an artifact of the fixed path. Second, the levels fall in both legs at every horizon: informed markout drops from 0.385 to 0.240 per unit at $h = 20$, uninformed from 0.827 to 0.670, and the informed toxic-fill rate rises from 0.381 to 0.437. That drop is the impact cost the exogenous design could not book.

The economically meaningful question is whether makers still profit against informed flow once their own fills move the price, and at this calibration the answer is yes: the informed level at $h = 20$ is 0.240 with pointwise CI [0.192, 0.289], positive but 0.145 per unit lower than the exogenous arm’s. The sweep in Table 9 is exploratory. Its sampled informed level declines from +0.095 at $\lambda = 0.2$ to −0.050 at $\lambda = 0.3$ (whose pointwise interval crosses zero) and −0.196 at $\lambda = 0.4$, while the uninformed level remains positive over these six values and the paired gap changes from 0.441 to 0.393. The data therefore bracket an observed sign change between sampled values; they do not identify a crossover at $\lambda \approx 0.27$ or attach uncertainty to an interpolated value. The default calibration ($\lambda = 0.1$) is one explicitly sampled point below that bracket.

The impact model is the engine’s stylized linear permanent impact with a single volume normalizer rather than an estimated impact curve, and linearity is itself a modeling commitment [Huberman and Stanzl, 2004]; the temporary component is folded into the maker ladder rather than carried as a separate η term; fills are still struck at the maker’s quote around the pre-move mid; and the quoting side is one follower model, the same fixed Avellaneda–Stoikov-plus-fixed-spread roster as above, which never widens in response to realized toxicity, so this is still not the equilibrium quote response of Glosten and Milgrom [1985]. What the arm establishes is narrower: on this sampled calibration grid, with the same quotes and information, informed-flow maker profitability declines as flow impact increases, while the paired estimated gap changes modestly. The exogenous and endogenous arms share one code path and one artifact, pinned by a golden hash asserting that the exogenous reports do not move when the endogenous arm runs; the endogenous evidence, per-episode vectors and paired arm changes included, lives under the endogenous key of the same artifact.

7.4 Adversarial realism: the in-band inversion attack

Section 10 names a threat the leak-free interface property does not cover: a scenario is a pure function of a public 64-bit seed under a public deterministic transform, so future bars are in principle computable from past bars, and no forbidden operation is needed. This probe converts that named threat into a measurement. Three adversaries predict the next bar’s return from a causal prefix only, on the three volatility tiers, 16 seeds per tier (4 symbols, 120 days, 30 warmup bars): an *unconditional baseline* that predicts the expanding prefix-mean return per symbol; an *honest statistical adversary*, a ridge AR(5) refit on the prefix at every bar, with no knowledge of the generator; and a *known-seed oracle*, which replays the generator from the true seed and therefore predicts every future bar exactly, the upper bound on any inversion attack. We also run each predictor as a frictionless unit-gross long/short sign-following policy scored through the public `score_run` kernel, and report the kernel’s deflated-Sharpe probability.




Figure 8: The endogenous arm of F6. Left: maker markout per filled unit by horizon, informed versus paired-uninformed flow, exogenous versus endogenous path ($\lambda = 0.1$, $V = 240$), with t -based 95% CIs over 24 paired episodes; paired arm changes are serialized in the evidence. Right: endogenous markout per filled unit at $h = 20$ across an exploratory six-value λ sweep, with pointwise 95% CI bands. Dashed line: the default calibration.

The honest adversary extracts almost nothing. All “ $\pm$ ” quantities here are across-seed standard deviations over the 16 seeds, computed with the population convention (ddof 0), not confidence intervals. On Calm the ridge AR reaches $55.4\% \pm 2.8\%$ directional accuracy against the baseline’s $54.1\% \pm 2.4\%$: it detects the tier’s designed momentum autocorrelation, worth about one descriptive accuracy point over drift-following. On Hard and Extreme the AR accuracies are $51.2\% \pm 3.3\%$ and $50.6\% \pm 2.1\%$, with MSE no better than the unconditional baseline (Hard: 1.24×10^{-4} vs. 1.22×10^{-4} ; Extreme: 1.27×10^{-3} vs. 1.17×10^{-3}). No hypothesis test or equivalence margin was preregistered, so these summaries support only the descriptive statement that the adversary lands near the unconditional baseline. One post-hoc check is worth reporting rather than hiding: a paired t -test of the ridge AR against the baseline over the same 16 Calm seeds gives $t(15) = 3.36$, $p = 0.004$, so the one-point Calm edge over drift-following is larger than seed noise, exactly as the tier’s disclosed momentum structure predicts; the same test is inconclusive on Hard ($p = 0.08$) and Extreme ($p = 0.40$).

The oracle bound is total, and a bounded seed search reaches it. The known-seed oracle scores 100% accuracy, zero MSE, and DSR 1.00 on every tier: determinism means the gap between an honest adversary and a successful inverter is the entire distance from coin flipping to omniscience, so the question is only the cost of inversion. We implement the strongest cheap version, a brute scan over a bounded seed band of width 2^{16} centered on the published `start_level`, matching each candidate’s first-bar closes against the single observed opening bar. The scan recovers the true seed in 16 of 16 trials with zero collisions, from a prefix of one bar, at $14.7 \mu\text{s}$ per candidate (≈ 0.96 s per scan); every recovered seed then reproduces the full 30-bar observed prefix on its tier. When a deployment uses a bounded public (`start_level`, `num_levels`) specification, this in-band inversion is a table scan rather than cryptanalysis, and its scenarios should be treated as known to a motivated adversary (Figure 9). Extrapolating this Python harness’s measured per-candidate cost to the unbounded 2^{64} space gives ~ 8.6 million CPU-years. The figure is harness-specific: a native scanner would be orders of magnitude cheaper per candidate, which moves the number but not the conclusion that brute-force inversion of an unpublished 64-bit seed is infeasible by enumeration.

What remains open. The probe brackets the attack but does not close it. Between the honest regression (no edge beyond disclosed structure) and the band scan (total recovery given a bounded public band) lies the case that matters for unbounded deployments: analytic inversion of SplitMix64 [Jr. et al., 2014] through the mul/add/div/max price transform from observed reals. We implement and test the narrow algebraic step: SplitMix64’s 64-bit finalizer is exactly invertible when its complete




Figure 9: Generator predictability by tier for the unconditional baseline (prefix mean), the causal ridge AR(5) adversary and the known-seed oracle, distinguished by colour and hatching. Left: next-bar directional accuracy; the dotted line is 0.5 accuracy. Right: the kernel’s deflated Sharpe of the frictionless sign-following policy driven by each predictor. Error bars are ± 1 population standard deviation over the 16 seeds, not confidence intervals. The causal AR adversary remains near the unconditional baseline; knowing or recovering the seed gives perfect prediction.

output word is known. The generator’s `next_unit`, however, exposes only the upper 53 bits as a float, leaving 2^{11} candidate post-increment states from one exact unit observation before any price arithmetic is considered. The public interface reveals neither that unit nor an unrounded finalizer output: it reveals prices after data-dependent arithmetic, floors and serialization. Deciding among those candidates from observed prices remains the untested cryptanalytic step; the new primitive reduces the question without claiming a tape-level break. This probe therefore supports only the narrower conclusion that resistance to this attack rests on seed secrecy and entropy rather than transform hardness, and that published bounded seed bands forfeit both.

7.4.1 Sealed evaluation seeds: defeating the bounded-band scanner

The public held-out set is disjoint from training but enumerable, so the band scanner above recovers its scenario identities without violating the observation interface. The opt-in mitigation derives each named slot from a secret salt k :

$$\text{seed}_i = \text{EVAL_SEED_BASE} + (F(k, i) \bmod (2^{64} - 1 - \text{EVAL_SEED_BASE})).$$

The shipped F combines FNV-1a/64 with public SplitMix64 finalizers. It is deterministic and pinned across Rust and Python, and the tested salts and slots produce held-out-band seeds distinct from the public set. No universal injectivity claim is possible for a finite 64-bit mapping, and this construction is PRF-style rather than a certified MAC.

The intended workflow is commit–reveal, as Section 6.3 prescribes it. Before evaluation, the host publishes $\text{SHA-256}(k)$ and withholds k ; after scoring, it reveals k so anyone can recompute the named seeds and replay the run. SharpeBench’s forward-window format can bind that hash before entries arrive. This paper’s evidence is a simulated workflow containing the revealed salt, not proof that any live evaluation used secret custody, a non-reuse policy or a reveal deadline.

On 16 tested slots, the scanner over the public 2^{16} band recovers 16/16 public seeds and 0/16 sealed seeds while the salt is withheld; revealing the salt replays 16/16. Under a uniform-seed approximation,

that scanner covers about 3.6×10^{-15} of the derived band, so its expected number of recoveries over 16 slots is 5.7×10^{-14} : the 0/16 outcome is what the coverage arithmetic already implies, and the run verifies the implementation rather than testing an uncertain hypothesis. The result is therefore narrow: high-entropy salt secrecy defeats this particular bounded-band enumeration budget. It is not a cryptographic security proof, and it does not answer tape-level inversion after observations begin.

8 Calibration findings

The previous section audits the environment. This one runs policies through it and reports what the resulting instruments read. Five findings follow: the baseline board under the corrected scoring kernel and the eligibility witness that answers the question the empty board raises (Sections 8.1 and 8.1.1), regret against the closed-form market-making reference (Section 8.2), the generalization gap and cross-regime transfer matrix (Section 8.3), the failure-mode distribution (Section 8.4), and the replicated ecology probe (Section 8.5). Every follow-up unit sits directly after the finding it extends.

F1, F3, F7, F8 and the witness run on the canonical position-trading generator, whose tape the realism diagnostic of Section 7.1 declines to certify on 23 of 24 seeded panels; F2 runs on the Avellaneda–Stoikov market-making model instead. Every number below is therefore a calibration of an instrument on tape the paper does not claim is market-realistic, and Section 10 carries that constraint forward into every claim built on it. No golden or parity test pins the F1–F3 outputs themselves, which come from Python rollouts through the native binding on one platform; what is tested is narrower. The native and Python suites pin the scenario generator at seven named specifications, and the canonical and clustered ones also reproduce on wasm32 and in the shipped npm bundle (`scenario_kernel_matches_native_and_every_committed_golden`, `exported_generate_scenario_reproduces_every_committed_golden`).

`wasm_kernel_matches_native_engine` requires the WebAssembly JSON kernel, compiled for the host, to reproduce the native engine’s returns and trace for one buy-and-hold fixture. The backtest and kernel goldens pin `run_baseline`, `replay_run`, `walk_forward`, `stress_suite` and `tag_regime` outputs for named inputs natively, on wasm32 and in the npm bundle (`backtest_goldens_reproduce_natively`, `kernel_goldens_reproduce_natively` and their `exported_..._under_wasm32` counterparts). The market-clearing engine’s opt-in entry points are held to its linear path by native equality tests (`a_unit_exponent_is_byte_identical_to_the_linear_path`, `absent_uncertainty_set_is_byte_identical_to_the_point_estimate`). The Python evidence reductions are seed-deterministic on one platform but sit outside all of these. F1 and F3 report 2,000-resample seed bootstrap intervals, F2 reports t-based 95% intervals over committed per-episode vectors, and F7 and F8 are count-based, reporting episode counts and seed-level winner distributions rather than interval estimates. Every entry point, seed set and quantity is named where it is reported, and the commands are in Section A.

8.1 F1: Pipeline validation under the corrected kernel

We roll the reference-policy field (flat, equal-weight long, momentum, minimum-variance, maximum-Sharpe, fractional-Kelly volatility targeting) over sixteen scenario seeds in each of the Calm, Hard and Extreme tiers via `run_baselines`. The `SharpeBench score_run` kernel scores each policy’s pooled return series with the declared trial count equal to the field size, and each row carries a seed-paired bootstrap 95% confidence interval on the deflated Sharpe (2,000 resamples).

Sharpe-ratio convention. Every Sharpe ratio in this paper is the ex-post ratio of Sharpe [1966, 1994] as the `SharpeBench` kernel computes it: the mean per-period return divided by the standard deviation of per-period returns, with no risk-free rate subtracted (cash earns zero, so the differential return is the portfolio return) and no annualization. For n returns with estimated Sharpe ratio $\widehat{SR}$, sample skewness $\hat{\gamma}_3$ and non-excess kurtosis $\hat{\gamma}_4$, the probabilistic Sharpe ratio against a reference

Table 10: Baseline leaderboard under the corrected kernel, 16 seeds per tier. DSR is the pooled deflated Sharpe with its seed-paired percentile bootstrap 95% interval (2,000 resamples); pass^k is the share of seeds whose own run passes the per-run reliability gate; mean return is the pooled mean per-bar return in units of 10^{-4} . DSR and its interval are rounded to two decimals for display, and the prose quotes four decimals. Within each tier, rows are sorted by unrounded DSR, descending, and rows with exactly equal DSR keep the producer’s policy order. Rows that print the same DSR are not separated by this table. [†]The two Calm long books tie only at the displayed precision: their unrounded DSRs differ past the fourth decimal and their intervals overlap, so their order is not a ranking. [‡]`flat` holds no position, so its return series is identically zero and its Sharpe ratio is undefined; the frozen artifact nonetheless reports a DSR of 0.0007 on every tier, the scorer’s value for that undefined statistic rather than an estimate (see text). No policy is rank-eligible on any tier: eligibility requires the gate to pass on every seed (a pass^k rate of exactly 1.00), and the best rate is 0.75.

Tier	Policy	DSR	95% CI	pass^k	Mean return
Calm	<code>kelly_vol_target</code>	1.00 [†]	[0.98, 1.00]	0.75	3.37
	<code>equal_weight_long</code>	1.00 [†]	[0.83, 1.00]	0.62	3.56
	<code>min_variance</code>	0.98	[0.03, 1.00]	0.44	2.40
	<code>flat</code>	0.00 [‡]	[0.00, 0.00]	0.00	0.00
	<code>momentum</code>	0.00	[0.00, 0.00]	0.00	−8.69
	<code>max_sharpe</code>	0.00	[0.00, 0.00]	0.06	−10.02
Hard	<code>equal_weight_long</code>	0.11	[0.00, 0.95]	0.31	2.43
	<code>kelly_vol_target</code>	0.03	[0.00, 0.82]	0.31	1.18
	<code>min_variance</code>	0.00	[0.00, 0.46]	0.19	0.51
	<code>flat</code>	0.00 [‡]	[0.00, 0.00]	0.00	0.00
	<code>max_sharpe</code>	0.00	[0.00, 0.00]	0.00	−9.28
	<code>momentum</code>	0.00	[0.00, 0.00]	0.00	−8.44
Extreme	<code>min_variance</code>	0.02	[0.00, 0.33]	0.06	5.85
	<code>equal_weight_long</code>	0.02	[0.00, 0.31]	0.06	4.73
	<code>kelly_vol_target</code>	0.00	[0.00, 0.11]	0.12	1.39
	<code>flat</code>	0.00 [‡]	[0.00, 0.00]	0.00	0.00
	<code>momentum</code>	0.00	[0.00, 0.00]	0.06	−3.06
	<code>max_sharpe</code>	0.00	[0.00, 0.00]	0.00	−15.78

SR^* is

$$\text{PSR}(SR^*) = \Phi \left(\frac{(\widehat{SR} - SR^*)\sqrt{n-1}}{\sqrt{1 - \hat{\gamma}_3 \widehat{SR} + \frac{\hat{\gamma}_4 - 1}{4} \widehat{SR}^2}} \right)$$

[Bailey and López de Prado, 2012], one minus the p-value, under its normal approximation, of the one-sided test that the Sharpe ratio does not exceed SR^* ; it is not the probability that the true Sharpe ratio exceeds SR^* . The deflated Sharpe (DSR) is the PSR at a reference raised for the declared search breadth [Bailey and de Prado, 2014]. The per-run reliability gate applies the PSR to each seed’s own run, and pass^k is the share of seeds whose run passes it.

This experiment validates the producer-scorer pipeline end to end; the scorer-side discovery it reproduces belongs to the companion paper. Under the SharpeBench 0.0.8 kernel that SharpeArena pinned before its 0.9.0 release, this board was all zeros: every baseline scored a deflated Sharpe of 0.0000 on every tier, which reflected a unit error rather than strictness. The kernel applied its annualized 0.5 deflation prior per period without conversion, which set the bar near an annualized Sharpe of 18 on daily bars, a level nothing clears. Finding and fixing that unit bug is the companion SharpeBench paper’s Finding 1 [Toca, 2026]. The fix, converting once at 252 periods per year, shipped in SharpeBench 0.3.0, and SharpeArena 0.9.0 moved its pin to SharpeBench 0.5.0, which carries it. What this paper claims is only that the producer-scorer separation surfaced the bug, because trajectories produced here could be re-scored under both kernels and the discrepancy attributed to the scorer alone.

Table 10 and Figure 10 show the corrected board. On Calm, drift is free deflated Sharpe: `kelly_vol_target` and `equal_weight_long` saturate the DSR at 1.0000, with bootstrap CIs

[0.9798, 1.0000] and [0.8343, 1.0000]. The equal-weight long book holds the tape and acts on no signal, so its Calm score is the tier’s drift alone: on Calm, drift by itself saturates the deflated Sharpe. The mean-return column is the realized drift each book earned, and it does not order the tiers’ scores, because the DSR reads mean return relative to the volatility of returns, which the committed evidence does not serialize per tier; the equal-weight book earns more per bar on Extreme than on Calm and scores far lower. A DSR on this board therefore mixes a policy’s behavior with its tier’s drift, and the table does not separate the two. Scoring each policy in excess of the equal-weight tape would separate them, but it requires rerunning the frozen producer and is planned rather than reported. `min_variance` reaches 0.9849 on a near-vacuous interval of [0.0264, 1.0000]. The deflated Sharpe is a probability bounded in [0, 1], so 1.0000 is its ceiling rather than an unbounded score, and the percentile bootstrap undercovers at that boundary: the saturated intervals are indicative, and no bias-corrected variant is applied. Yet nothing on the board is rank-eligible, because no policy passes the per-run reliability gate on every seed; the best pass^k rate is 0.75. On Hard the best DSR collapses to 0.1114 and on Extreme to 0.0243, and the long baseline’s pass^k rate degrades monotonically with stress (0.62, 0.31, 0.06). Each rate is 16 Bernoulli seeds, so near a rate of one half a single rate carries a nominal 95% half-width of about 0.24, and a difference between two of them about 0.35: only the Calm-to-Extreme drop of 0.56 exceeds that, so the monotone ordering describes these seeds rather than a resolved ranking. `momentum` and `max_sharpe` lose money outright on every tier (mean per-bar returns down to -1.6×10^{-3} on Extreme).

The `flat` rows are not scores. The flat book holds zero target weight on every bar, so its return series is identically zero on every seed of every tier, and its Sharpe ratio, a zero mean over a zero standard deviation, is undefined. The frozen artifact nonetheless reports a DSR of 0.0007 on each tier, with a zero-width interval: the SharpeBench 0.5.0 scorer that produced it replaces a zero standard deviation with a Sharpe ratio of zero, takes the skewness as zero and the kurtosis as three, and scores the result. The printed value is that convention’s output for an undefined statistic, not evidence about the policy. A repair made after release 0.29.0 withholds these values at SharpeArena’s own boundary, which it was written to do while the pinned scorer still substituted a Sharpe ratio of zero. The pinned scorer is now SharpeBench 0.27.0, which refuses such a track itself, with the same reason and in the same fields, so the refusal no longer rests on SharpeArena’s boundary alone: on every track that boundary refuses, the kernel has already recorded the refusal, and SharpeArena’s check is a second line that changes no reported value. A track whose observations all equal the first, zero or not, is treated as having no Sharpe ratio; the test is exact equality rather than a volatility threshold, so a dispersed track of any volatility scores exactly as before. `score_run` reports such a track with the typed unavailability `deflation_error`, whose reason reads “returns must not be constant: a constant series has no Sharpe ratio”, the reason the corresponding SharpeBench repair uses and the pinned kernel now writes itself. Its deflated Sharpe, PSR and deflation bar print the floor value 0 rather than an estimate, its bootstrap interval and rolling Sharpe summary are absent, and it fails the per-run reliability gate and rank eligibility. The seed-paired bootstrap interval and the paired-difference test refuse a constant pooled band with the same reason, and a track that is not constant but whose computed standard deviation underflows to zero is refused as “Sharpe ratio is not finite”. The repair also closes a worse case of the same shape: a constant nonzero track has a tiny nonzero computed standard deviation, so before it a track of 0.001 on every bar scored a deflated Sharpe and PSR of 1.0 and was rank-eligible. Under the repaired code the `flat` rows report the reason in place of a deflated Sharpe, a pass^k rate and an interval. SharpeBench 0.27.0 carries a second repair of the same shape, which stops a track with no Sharpe ratio voting on other agents’ deflation bar where that bar is measured from the field. No panel in this paper is affected, and no value in Table 10 or in any other table here moves with it. This board’s bar is not measured from a field: `run_baselines` scores each of the six policies as its own one-agent submission through `score_run`, and a one-agent submission deflates against the configured cross-trial dispersion prior in every SharpeBench release this paper has pinned. Every producer here that reaches the kernel at all reaches it that way; none calls the field-ranking entry point where a measured bar is possible. The `flat` track therefore never voted on another policy’s bar. That holds by the producer path rather than by readback, because the committed artifacts do not serialize which dispersion source the scorer took.




Figure 10: Deflated Sharpe per baseline and tier. Error bars are seed-paired percentile bootstrap 95% intervals ($B = 2,000$). Tiers differ by colour and by hatching (Calm plain, Hard diagonal, Extreme dotted). Calm saturates the long baselines; Hard and Extreme collapse them toward zero. The `flat` bars are the frozen artifact’s value for an undefined statistic (Section 8.1).

Eligibility is a conjunction, and each leg catches what the other misses. Deflation alone would crown drift on Calm; pass^k alone would miss the overfit breadth the deflation discounts. The number a real agent has to beat is a positive deflated Sharpe that survives every evaluation seed, and no reference policy produces one. The board alone does not show that the eligibility set is non-empty for this task family. Section 8.1.1 answers that by construction, injecting an out-of-band oracle signal of controlled strength and locating the strength at which the gates first open; the eligibility set is non-empty on every tier, and the per-run reliability gate that fails every baseline here is the leg that binds there too.

8.1.1 Rank-eligibility witness: the eligibility set is non-empty

The baseline board above contains no rank-eligible policy, leaving open whether the conjunction can be satisfied at all. We answer by construction with an out-of-band oracle rather than an entrant. For each seed, the script replays the deterministic, policy-invariant tape one bar ahead and mixes standardized next-bar truth z_t with a seeded noise path ε_t :

$$\text{signal}_t(s) = sz_t + \sqrt{1 - s^2} \varepsilon_t, \quad 0 \leq s \leq 1.$$

The mixing weight s is a nominal signal-strength parameter, not the realized correlation between signal and truth on a path: z_t is each return divided by the episode’s own sample standard deviation without centering, and ε_t is one finite draw, so the per-path correlation differs from s . Within one noise path the same ε_t is reused at every s (common random numbers), so changing signal strength does not silently change the experiment. A sign-following policy trades every bar; a deadband policy trades only when $|\text{signal}| > 1$ and otherwise holds. Both are scored through the unchanged SharpeBench conjunction over 16 seeds. The per-run reliability gate requires every seed’s PSR (Section 8.1), here against a reference Sharpe of zero, to reach 0.90. The script records the full coarse eligibility set, the sampled strengths at which the kernel conjunction accepts. It refines an *observed crossing bracket* by bisection only when that coarse set is monotone; otherwise it reports the observed set without calling any crossing a threshold. A crossing located under one noise path is a functional of that draw, so we repeat the whole procedure (coarse sweep, monotonicity check, bisection) under five independent noise paths per cell: the first is the path serialized as the primary witness, the other four redraw ε_t . Each bisection bracket is numerical resolution on the stated grid: bisection stops once a bracket is no wider than 0.005 in s , and because every bracket starts on a grid interval of 0.05, 0.1 or 0.2, every recorded bracket ends at width 0.0031, a half-width of 0.0016 (Section A); the spread

Table 11: Observed eligibility crossing in the nominal signal strength s : mean [min, max] of the bisection-bracket midpoint over five independent noise paths per cell. The min to max range is variation across noise realizations; the bisection bracket around each individual crossing has a realized half-width of 0.0016, inside the stopping rule’s bound of 0.0025, and is not shown. Crossings are found by bisection after a monotone coarse-grid check on every path, and are not global optima or theorem-level thresholds. “none (0/5)” means that no sampled $s \leq 1$ is eligible on any of the five paths. At every one of the 50 attained crossings, the last failing gate is the per-run reliability gate, the requirement that every seed’s marginal PSR reach 0.90.

Variant	Band	Calm	Hard	Extreme
sign_follow	gap-band witness	none (0/5)	0.730 [0.702, 0.752]	0.343 [0.302, 0.389]
sign_follow	F1 table	none (0/5)	0.783 [0.764, 0.820]	0.353 [0.317, 0.445]
deadband_hold	gap-band witness	0.848 [0.777, 0.930]	0.500 [0.408, 0.623]	0.337 [0.283, 0.383]
deadband_hold	F1 table	0.690 [0.620, 0.758]	0.621 [0.483, 0.755]	0.342 [0.283, 0.398]

of the crossing across noise paths is a separate, statistical quantity, and it is what Table 11 reports. All 50 attained coarse series (10 cells, five paths each) are monotone.

The result is an existence proof and a calibration. The deadband construction is eligible on every tier on every noise path, so the observed eligibility set is non-empty. The per-run reliability gate is last to open at all 50 observed crossings; pooled DSR and bootstrap gates open earlier. The crossing itself is not a constant of the tier: across noise paths it moves by 0.05 to 0.27 in s (a factor of 15 to 85 above the bisection resolution), widest for the deadband policy on Hard, so any single-path threshold is one draw from that range and the table should be read at the precision of the range rather than of the bracket. Differences between the two bands are differences between two 16-seed scenario bands as well as between noise draws, and with five paths per cell we do not test them. The sign follower never becomes eligible on Calm, on any of the five paths, because its every-bar turnover costs more than the low-volatility signal is worth; at $s = 1$ the noise term vanishes, so that endpoint is noise-invariant by construction, and the failure at every $s < 1$ is what the replication adds. The environment therefore imposes no activity requirement, but its cost model does require a Calm policy to ration turnover (Figure 11).

8.2 F2: Regret against the closed-form reference policy

On the Avellaneda–Stoikov market-making task, `mm_regret` runs each candidate quoting policy and the closed-form reference policy (Section 3.7) on the same sixteen seeded episodes ($\sigma = 2.0$, $\gamma = 0.1$, $\kappa = 1.5$, 200 steps) and reports the mean reward gap to the reference. Candidates are the reference itself and fixed-spread quoters over a grid of half-spreads. The evidence commits the per-episode regret vector per grid point; intervals are t-based 95% CIs over the sixteen paired episodes.

The reference’s regret is 0.0 by construction, which confirms the plumbing. The fixed-spread curve is U-shaped (Figure 12): regret is 84.6 (95% CI [68.2, 101.1]) at a half-spread of 0.05 (quoting too tight, filled constantly, run over by inventory), falls to a minimum of 0.037 at 0.5, and rises again to 36.6 ([33.0, 40.2]) at 2.0 and 58.8 ([55.3, 62.3]) at 4.0 (quoting too wide, earning nothing). The minimum is the point the intervals do not resolve: the 0.5 point’s CI is $[-8.7, 8.8]$ and the 1.0 point’s is $[-5.4, 12.4]$. Both intervals include zero, so sixteen episodes do not resolve either point from the analytical reference; this failure to reject zero is not evidence of equivalence. At the observed per-episode dispersion, sixteen paired episodes detect a mean regret of about 12 reward units at 80% power and a two-sided 5% level, roughly a third of the 36.6 measured at a half-spread of 2.0, so differences smaller than that are outside this design’s reach. The tight and wide extremes are separated from zero by wide margins. The instrument resolves the shape of the curve, not the sign of its minimum.

The point of shipping the reference is that agents on this task are scored on regret rather than raw PnL. A market-making leaderboard ranked on PnL rewards whichever policy drew the friendliest flow; regret against a fixed analytical reference on the same episodes removes that luck axis when




Figure 11: Witness sweep on the gap-band witness subset under the primary noise path, one row per policy. Left panels: pooled deflated Sharpe against the nominal signal strength s ; the band around each curve is a seed-paired percentile bootstrap 95% interval (2,000 resamples), filled markers are rank-eligible sampled points and hollow markers are not, and the dotted line is the DSR bar of 0.95. Right panels: the pass^k rate over the 16 seeds, with a dotted line where all 16 pass. In every panel a dashed vertical line marks the primary-path bisection crossing, drawn only after monotonicity is verified, and a shaded column spans the minimum to maximum crossing over five independent noise paths; in the right panels a horizontal bar repeats that range with its marker at the mean. Tiers differ by marker (Calm circle, Hard square, Extreme triangle) and by column hatching (Calm plain, Hard diagonal, Extreme dotted).

both arms draw the same random stream. Identical seeds do not guarantee that stream (Section 3.7), so `mm_regret` now checks it and refuses an unpaired comparison. A committed test runs every half-spread of this subsection through the checked function at the default parameters and reproduces the frozen regrets bit for bit, so every comparison here shares its stream.

8.3 F3: Generalization gap and cross-regime transfer

`generalization_gap` scores the default reference policy (equal-weight long) on disjoint train and held-out seed bands (16 seeds each, separated by a 10,000-seed gap) within each tier. `cross_regime_transfer` then holds sixteen seeds fixed and varies the regime, producing the full three-by-three matrix of transfer gaps, each the in-distribution DSR minus the zero-shot out-of-distribution DSR, whose diagonal is zero by construction because identical regimes reuse byte-identical environments. The policy is fixed and parameter-free, so nothing is trained: “source regime” below means the regime the policy is nominally selected on, and for such a policy the matrix reduces to pairwise differences of three per-tier deflated Sharpes. Those three are the equal-weight long book’s scores on the same sixteen seeds as Table 10, whose mean returns the two artifacts record identically, so the drift confound of Section 8.1 carries over whole: on Calm drift alone saturates the score, and every transfer gap is a difference between drift-driven DSRs. What F3 measures is the instrument, how much a regime switch moves the score of a signal-free long book, and not transfer learning or the transfer of skill. These are the 16-seed subsets Section 6.1 exempts fixed policies from




Figure 12: Mean regret versus the analytical Avellaneda–Stoikov reference policy over 16 paired episodes, with t-based 95% CIs. Fixed-spread quoting is U-shaped in the half-spread; the reference defines zero.

Table 12: Left: within-tier generalization gap (train minus test deflated Sharpe) for the fixed equal-weight reference. Right: the cross-regime transfer gap, in-distribution DSR minus zero-shot out-of-distribution DSR, with rows the source regime and columns the scored regime; a cell is a difference of two scores, not an out-of-distribution score, and the diagonal is zero by construction. Cells are read only through their seed-paired bootstrap intervals: a cell whose interval excludes zero is resolved in sign, and the Hard↔Extreme cells (± 0.096), whose intervals straddle zero, are not interpreted.

Tier	Train DSR	Test DSR	Gap		Calm	Hard	Extreme
Calm	1.0000	0.9988	+0.0012	Calm	0.000	+0.877	+0.973
Hard	0.1231	0.4618	−0.3388	Hard	−0.877	0.000	+0.096
Extreme	0.0269	0.1280	−0.1012	Extreme	−0.973	−0.096	0.000

replacing with the full canonical band. The evidence commits per-seed return series per cell; CIs are seed-resampled percentile bootstrap ($B = 2000$), seed-paired across regimes for the matrix cells.

Table 12 reports both. The within-tier gaps are the expected control: the reference policy has no fitted parameters, so it cannot overfit its train band, and the instrument correctly reads near zero on Calm (+0.0012) and noise-signed values on the stressed tiers (the negative Hard and Extreme gaps mean the held-out band happened to be kinder; with no learning, the gap is pure seed noise, and how much band luck a 16-seed evaluation carries is read from the interval widths rather than from any single realized gap). The per-band bootstrap intervals say the same thing at interval width: Hard’s train-band DSR of 0.1231 carries a CI of $[0.000, 0.944]$ and its test band $[0.037, 0.948]$ (the bands are disjoint seed sets, so these are unpaired), overlapping almost entirely, which is what a pure-noise gap looks like. Hard’s train-band DSR of 0.1231 here and Table 10’s 0.1114 for the same policy on the same sixteen seeds differ only in the declared trial count the deflation prior consumes, six declared policies in F1 against no declared search in the gap metric; rescoring the committed F3 return series at six trials reproduces 0.1114 exactly.

Table 13: Failure-mode counts per tier, 128 episodes each (8 policies $\times$ 16 seeds). No episode reaches bankruptcy or a cascade wipe; failure is dominated by structural mandate breaches, which are flat across tiers, while the stress-sensitive modes are the drawdown family.

Tier	Clean	Stopped out	Mandate struct.	Mandate DD	Mandate inv.	Bankrupt	Clean rate
Calm	51	0	50	19	8	0	0.40
Hard	49	0	50	21	8	0	0.38
Extreme	43	10	43	25	7	0	0.34

The transfer-gap matrix is where regime structure appears, and every cell is classified by one rule, fixed here for all cells: a cell whose seed-paired bootstrap interval excludes zero is resolved in sign, and the width of that interval, not the rule, says how precisely the magnitude is known. The Calm-involved point estimates are the largest: source regime Calm scored zero-shot on Hard loses 0.877 of deflated Sharpe, and Calm to Extreme loses 0.973 of a score bounded in $[0, 1]$. Both intervals exclude zero, so both gaps are resolved in sign, but they differ in precision: Calm to Extreme is $[0.635, 0.999]$, while Calm to Hard is $[0.054, 0.999]$, which leaves its magnitude anywhere from near zero to the whole score. The Hard-to-Extreme cell (0.096, CI $[-0.218, 0.911]$) has an interval straddling zero, so it is not interpreted: at 16 seeds the instrument cannot distinguish it from band luck. A bounded, heavily deflated score bootstrapped over 16 seeds is a coarse instrument, which is itself a calibration result, and Section 6.2 turns it into a prescription. The matrix is antisymmetric because the policy is fixed, and its diagonal is zero by construction. Its off-diagonal magnitude is the quantity a within-tier gap is structurally blind to: a policy scored only on calm seeds can carry a saturated DSR into a leaderboard while being worth 0.03 under stress. Both metrics are cheap, and the protocol requires reporting both.

8.4 F7: Failure-mode distributions

`classify_episode_failure` classifies every rollout of the eight-policy field (the six reference policies plus the two behavioral counterparties), across all three tiers and sixteen seeds (384 episodes), into the worst-case-first taxonomy: cascade-wiped, bankrupt, stopped-out, mandate breach (structural, drawdown or inventory), or clean, with a per-scenario sampled mandate and a 50% drawdown stop in force. The two evidence dispositions of Section 3.8, invalid evidence and execution failure, are counted in the same rollout and are zero on every tier, so they are omitted from Table 13. This artifact was regenerated for SharpeArena 0.29.0 by rerunning its producer; the other F1–F8 artifacts were not.

Mandates are sampled per scenario and assigned to fixed policies that cannot observe or adapt to them, so a structural breach is guaranteed whenever the draw pairs, say, a shorting policy with a long-only mandate. The flat structural-breach counts across tiers (50, 50, 43) are the signature of that assignment mechanism, not a behavioral finding about agents; a mandate-aware field could invert the distribution entirely, and measuring one is the interesting follow-up. The draw itself changed in SharpeArena 0.29.0: the momentum style is no longer sampled, because `mandate_breach` reads per-bar portfolio weights and pooled returns and has no per-symbol returns to grade “lean into recent winners” against (Section 3.8). Episodes that drew it in the earlier artifact are redrawn as styles that do carry a rule, so 264 of the 384 episodes are now graded against a structural constraint where 168 were before, and the structural counts rise accordingly. The remaining 120 draw the declared permissive control, which is scored by its optional caps alone and says so in its prompt text.

Table 13 and Figure 13 show the shape of failure as stress rises. The clean rate falls from 0.40 on Calm to 0.34 on Extreme. Each rate is 128 episodes, and the independent-binomial 95% half-width of a difference between two such rates is about 0.12. That width is not a floor or a ceiling for this design: episodes that share a seed or a policy are not independent, dependence does not fix the sign of their covariance, and a paired cross-tier difference over the same policies and seed indices has a covariance of its own. We therefore read the counts descriptively. Every tier-to-tier step, Calm to Extreme included at 0.40 versus 0.34, is smaller than the nominal width, and the monotone ordering of clean rates describes these 16 seeds rather than a resolved difference. What the tiers




Figure 13: Failure-mode counts per tier, each summed over the 8 policies and 16 seeds (128 episodes per tier, 384 in total), for every disposition in the rollout schema; the bankruptcy, cascade-wipe, invalid-evidence and execution-failure bars are zero. Tiers differ by colour and by hatching (Calm plain, Hard diagonal, Extreme dotted).

do separate is the composition of failure rather than its rate. Structural mandate breaches are flat, as the mechanism above dictates: a shorting policy under a sampled long-only mandate breaks the rule on every tier. The stress-sensitive modes are the drawdown family: mandate drawdown breaches rise from 19 on Calm to 25 on Extreme, and hard stop-outs at the 50% line appear only on Extreme (10 episodes: 5 overconfident, 4 max_sharpe, 1 disposition_effect). No policy in this field ever reaches bankruptcy or a cascade wipe; the taxonomy’s worst states are reachable but not by these baselines at this stop level. Per policy, flat is clean 48/48 by construction, while overconfident and the momentum policy are the worst on Calm (3 clean of 16 each, 8 of each one’s failures a drawdown breach) and overconfident is the worst on Extreme (2 clean of 16). On this mandate-blind reference field, most episodes end in a process failure that a return-ranked board does not record: 231 of 384. Counted against failures rather than episodes the share is 231 of 241, because only ten episodes fail on PnL at all, every one of them a stop-out.

8.5 F8: Ecology outcomes under shocks, replicated over seeds

run_ecology seeds the replicator with the eight baseline species over a shared-book market via market_payoffs (12 generations, field size 8, innovation every 4 generations breeding parameter-perturbed variants of the leader) and runs trajectories under a steady all-Calm control schedule and a regime_shocks schedule rotating Calm (generations 0 to 3), Hard (4 to 7) and Extreme (8 to 11). We replicate the probe over eight replicator seeds per schedule; Table 14 reports the cross-seed winner distribution, and the seed-0 trajectories (Figure 14) remain the illustrative detail. The two schedules chain episode seeds differently (the shock schedule strides seeds across generations), so control and shocked runs at the same replicator seed are matched in configuration, not in price paths.

The replicator flow rule sits in the market-ecology tradition: strategies as species whose profitability decays as their share of the flow grows [Farmer, 2002], and populations whose composition, rather than any single agent’s optimization, drives aggregate dynamics [Lux and Marchesi, 1999]. The implementation abstracts away much of what those models study: capital flows react instantly (no fund-flow lags), there are no leverage or margin constraints beyond the strategies’ own sizing, and there is no entry beyond the single leader-perturbation innovator. Its outputs are trajectories of this replicator map, not claims about market ecology in general.

Table 14: Replicated ecology outcomes over 8 replicator seeds per schedule: distribution of the final dominant lineage (bred variants credited to their founding species). The root winner differs between the schedules on 5 of 8 seeds, and on 3 of the 6 seeds where both schedules resolve a dominant lineage; the seed-0 replacement of the volatility targeter by the equal-weight long book occurs on 1 of 8 seeds.

Final dominant lineage	Control (all Calm)	Shocked (Calm/Hard/Extreme)
kelly_vol_target	4	7
equal_weight_long	2	1
flat (no resolved dominant)	2	0


The replication does not support the single-seed narrative. At seed 0 the story was clean: `kelly_vol_target` sweeps the control board (peak share 1.00, final 0.81), collapses under the shock schedule once the regime turns Hard, and `equal_weight_long` inherits dominance (final 0.67). Across eight seeds that story is the exception. Under the shock schedule a `kelly_vol_target` lineage ends dominant on 7 of 8 seeds (final shares near 0.90); the seed-0 handoff to the long book happens once. The control side is itself seed-dependent: `kelly_vol_target` wins 4 seeds, `equal_weight_long` wins 2, and on 2 seeds no species resolves dominance at all (all eleven end persistent at small shares, with a `flat` variant holding the largest share at 0.10). The root-level winner differs between the two schedules on 5 of 8 seeds, but on 2 of those the control side never resolves a dominant lineage at all; among the 6 seeds where both schedules resolve one, the winner differs on 3. At eight seeds the reordering is neither established nor excluded, and it has no consistent direction. Extinction pressure is more stable than identity: shocked runs produce a resolved dominant on 8/8 seeds with 4 to 7 extinctions per run (at seed 0, both behavioral counterparties are gone or collapsed on both schedules, `overconfident` extinct at generation 0).

One replicator trajectory is one draw from a heavy-tailed outcome distribution, and single-run ecology narratives do not survive replication. What the probe reliably produces at this scale is the outcome distribution itself: which lineages can win, how often shocks reorder the board, and how much of the population goes extinct, per committed seed set. Sharper claims need more seeds, varied shock orderings and path-matched schedule pairs, all of which are mechanical extensions of the committed script. A comparable ecological simulator already works at a larger replication scale, with 128 independent runs per configuration and 95% intervals across them [Zou et al., 2026, pp. 4, 9].


9 Related work

Environment interfaces. Gym [Brockman et al., 2016] demonstrated that a small, stable environment interface is worth more than any single environment, and Gymnasium [Towers et al., 2025] carried the interface forward with versioned environment IDs and a conformance checker; SharpeArena’s Python surface is a first-class Gymnasium citizen with versioned, namespaced IDs and adopts its `check_env`. PettingZoo [Terry et al., 2021] standardized the multi-agent case, which SharpeArena’s shared-book and batched-competition environments implement. Minari [Farama Foundation, 2024] standardizes offline-RL datasets; SharpeArena exports leak-safe rollouts to it. The wire contract of Section 4 differs from all of these in scope: it is a language-agnostic JSON boundary for external agent processes, governed like a protocol rather than a library API.

Evaluation rigor in RL. The environment’s evaluation stance imports three lessons from the RL evaluation literature. From the ALE revisitation [Machado et al., 2018]: determinism in the environment invites trajectory memorization, so evaluation must be designed against it; SharpeArena keeps determinism for verifiability and defeats memorization with procedural scenario distributions and held-out seed bands instead of injected stochasticity, and Section 7.4 measures how far that gets. From Procgen [Cobbe et al., 2020]: procedural generation over integer seed intervals makes generalization measurable, which Section 3.6 ports to markets and Section 6.1 states as a protocol. Prioritized Level Replay [Jiang et al., 2021] shows seed-level curricula shape what agents learn, which



(a) Control: steady Calm.



(b) Shocked: Calm, then Hard, then Extreme.

Figure 14: Seed-0 population-share trajectories over 12 replicator generations, the illustrative single run behind Table 14's distribution; seed 0 only. Both schedules start from the same eight founders (flat, equal_weight_long, momentum, min_variance, max_sharpe, kelly_vol_target, disposition_effect and overconfident). The innovation operator adds three bred variants to each: kelly_vol_target variants labelled ~lookback4, ~target_vol8 and ~kelly_fraction12 under control, and equal_weight_long variants labelled ~v4, ~v8 and ~v12 under shocks. Every band is labelled directly on the plot with a leader line and carries its own hatch; there is no legend. At this seed the regime rotation replaces the calm-regime winner with the equal-weight long book; across the eight-seed replication that replacement occurs once.

Table 15: Financial RL environments on the five property axes of Section 1. Entries describe the mechanism each system provides for the axis, not the quality of the system on its own terms. Where an entry rests on our reading of a system’s paper and repository rather than on a property its authors claim, it is written descriptively and marked “not targeted” rather than “absent”.

System	Leak-freedom mechanism	Determinism guarantee	guarantee	Trajectory verifiability	External-agent contract	con-	Generalization protocol
FinRL [Liu et al., 2020]	pipeline discipline over replayed history	not targeted		not targeted	in-process Gym API		train/test date split
FinRL-Meta [Liu et al., 2022]	DataOps pipeline discipline	not targeted		not targeted	in-process Gym API		train/test date split
ABIDES-Gym [Amrouni et al., 2021]	discrete-event simulation kernel	not targeted		not targeted	in-process Gym API		not targeted
TradeMaster [Sun et al., 2023]	workflow convention	not targeted		evaluation toolkit inside the agent’s trust domain	in-process API	Python	train/valid/test split
mbt_gym [Jerome et al., 2023]	model-generated state, no history to leak	not targeted		not targeted	in-process Gym API		not targeted
ABIDES [Byrd et al., 2020]	discrete-event simulation kernel	not targeted		not targeted	in-process classes	agent	not targeted
ABIDES-MARL [Cheridito et al., 2025]	synchronized agent actions; global observation lags one step, so opponents’ same-step actions stay hidden	not targeted		not targeted	in-process API compatible with standard RL libraries		holdout evaluation episodes after training
SharpeArena	interface shape: no operation returns a post-cursor bar (Section 3.1)	committed goldens across native, WASM and Python (Section 3.3)	FNV-1a	recompute-from-decisions replay with an adversarial tamper test (Section 3.4)	governed JSON wire contract v1.0 with schemas and fixtures (Section 4)		disjoint integer seed bands with a 10,000-seed gap (Section 6.1)

motivates keeping the train band operator-owned and the held-out band frozen. τ -bench [Yao et al., 2025] argues that agent reliability must be measured across repeated runs rather than on average; that requirement lives in the SharpeBench kernel this environment feeds [Toca, 2026].

Financial RL environments. A Gym-for-finance lineage predates this work. Table 15 places SharpeArena against it on the five property axes the environment competes on, which are the axes Section 1 names and the rest of this paper enforces.

The table compresses five substantive differences. FinRL and FinRL-Meta wrap historical market data into gym-style training environments at scale, and FinRL-Meta’s DataOps pipeline explicitly targets the data-quality half of the problem, survivorship bias and backtest overfitting included; because their environments replay recorded data inside a general-purpose Python process, leak-freedom is a property of pipeline discipline rather than of interface shape, and a caller who reaches around the wrapper is not stopped by anything. ABIDES-Gym wraps the ABIDES discrete-event market simulator in Gym environments, bringing multi-agent market microstructure to RL training; its focus is simulation fidelity, and to our knowledge it does not target cross-runtime byte-determinism or tamper-evident trajectories. TradeMaster is a holistic platform spanning data, environments, algorithms and evaluation for RL trading research; it standardizes the workflow rather than hardening the producer, and its evaluation toolkit runs inside the same trust domain as the agent, which is the arrangement Section 6.5 declines. mbt_gym provides model-based Gym environments for limit-order-book problems, including the same Avellaneda–Stoikov family as Section 3.7, with closed-form solutions available as baselines; it is the closest relative of the market-making task specifically, and Section 3.7 notes the relation. Where these descriptions rest on our reading of the systems’ papers and repositories rather than on properties their authors claim, we have kept them descriptive; none of these systems, to our knowledge, sets out to make look-ahead structurally unavailable at the interface, pin cross-runtime byte-identity with committed hashes, or make submitted trajectories third-party recomputable, which is the specific combination this paper contributes.

Agents across a process boundary. The external-agent column of Table 15 reads “in-process Gym API” or “in-process agent classes” for every prior financial environment, and that entry decides more than it appears to. An in-process agent is a Python object; calling it either returns a decision or raises, and there is no third outcome to type. A wire contract admits one, because a transport can fail without either party being wrong, and the failure then has to be represented somewhere or it will be absorbed by whatever the transport returns instead. The failure mode of Section 4.5 therefore does not exist for these systems, and its absence from their designs is not an oversight: they do not have the capability that creates it. The agent-evaluation literature outside finance does have the capability, and τ -bench [Yao et al., 2025] is the closest neighbour on the measurement question, arguing that agent reliability is a property of repeated runs rather than of an average and building the pass^k statistic to say so. The per-run reliability gate built on that idea lives in the SharpeBench kernel this environment feeds [Toca, 2026]. What SharpeArena adds on this axis is upstream of reliability rather than a competitor to it: a run that recorded a wire fault is not a low-reliability run, it is not a run, and the distinction has to be made before any statistic is computed over the k repetitions.

Market simulation and its validation. ABIDES [Byrd et al., 2020] is the closest simulator in spirit: a high-fidelity multi-agent market with a matching engine and background agent populations. SharpeArena trades some of that fidelity for properties ABIDES does not target, listed in Table 15. The microstructure models are the classical ones in stylized form: a Kyle-style linear permanent-impact term [Kyle, 1985], Almgren–Chriss temporary impact and execution cost [Almgren and Chriss, 2001], and the Avellaneda–Stoikov analytical quoting reference [Avellaneda and Stoikov, 2008]. The realism diagnostic compares the generator with selected stylized facts from Cont [2001], and Section 7 argues from the validation literature [Platt and Gebbie, 2018, Fagiolo et al., 2019, Lamperti, 2018, Vyetrenko et al., 2020, Nagy et al., 2025] that a calibrated gate a public generator can fail is worth more than an uncalibrated one it passes. That literature is also where SharpeArena’s gate stops: it is a directional three-check conjunction, not a distributional divergence in the sense of Lamperti [2018] or a full L1-style benchmark suite in the sense of Nagy et al. [2025], and adopting one of those is the natural next revision of the diagnostic. ABIDES-MARL, added to Table 15, extends ABIDES-Gym to several learning agents that act in one synchronized step, recovers gradual price discovery in an extended Kyle model, and states that its minimal agent mix does not reproduce volatility clustering [Cheridito et al., 2025, pp. 1, 18, 20].

Shared markets with model agents. Benchmarks for language-model agents in shared economies choose their counterparties differently. CoffeeBench’s background firms run Claude Sonnet 4.6, which is also one of the evaluated models [Sugiura et al., 2026, pp. 5–6]. In the EconEvals collusion litmus test, two model agents set prices for competing products, and the test scores GPT-4o at 0.71 and GPT-5 at 0.04 on a scale from competitive (0) to maximally collusive (1) pricing [Fish et al., 2026, pp. 60–62]. SharpeArena’s counterparties are a deterministic roster of fixed policies, which gives up that behavioral range in exchange for replayable tape and a fixed baseline for every shock schedule.

Scoring. SharpeArena exposes the published SharpeBench scoring and confidence helpers used by its diagnostics and adapters; it does not define a separate ranking contract. The board-level deflated Sharpe [Bailey and de Prado, 2014], reliability and process gates that consume this environment’s trajectories are specified in the SharpeBench paper [Toca, 2026]; this paper remains the producer half of that funnel.

10 Limitations

One synthetic generator family. All three tiers come from a single procedural vol-and-jump generator; no real price series enters the environment, and the tiers differ in that generator’s parameters, not in kind. F4 quantifies the cost: the tape has fat tails on the stressed tiers but weak volatility clustering at episode length, and the finite-panel-calibrated three-check conjunction passes on only 1 of 24 seeded panels. The Fano ratio is conditional-IID-calibrated but exploratory and does not gate.

The existing Calm calibration grid produces no preset that meets its diagnostic, confirmation and volatility rule. An agent that excels here has demonstrated skill against this generator’s structure, not against a market, and the realism gate says so rather than hiding it. A 2025 survey states that evaluation sets should never contain synthetic items [Fu, 2025, p. 15]; every table in this paper evaluates on synthetic tape, which is one reason the paper reads its findings as calibrations of instruments rather than as evidence about agents. Section 7 reads the failure as showing that this gate can refuse this generator, not as general discriminating power, and neither reading defends the tape: on the four-level empirical-validity scheme of Axtell and Epstein [1994], as surveyed by Fagiolo et al. [2019], the canonical generator sits at Level 0 on the gated axis, and every finding built on it is instrument calibration on uncertified tape. The paragraph below states which experiments run on other processes instead.

The interface guarantee is not an information guarantee. The leak-freedom of Section 3.1 is a statement about the interface’s shape: no operation returns a post-cursor bar. It is not a claim that public deterministic scenarios are unpredictable. Section 7.4 makes the distinction empirical. A causal AR(5) adversary stays near its unconditional baseline, but an exhaustive scan of a public 2^{16} seed band recovers all 16 tested seeds from one observed bar in about one second. Full-space brute force remains infeasible. A complete 64-bit SplitMix64 output inverts to one state, whereas its published 53-bit unit representation leaves 2^{11} candidates; tape-level inversion through the later price arithmetic remains open. The opt-in sealed derivation of Section 7.4.1 defeats this specific bounded-band scanner 0/16 when its salt is withheld, and a forward protocol can commit the salt hash before entry. The evidence here is a simulated workflow, not evidence that any live evaluation used it. Any public bounded-band evaluation should be treated as compromised by this scanner until it documents sealed seed custody, a pre-entry commitment, non-reuse and a reveal deadline. The derivation is PRF-style, not a certified MAC, so its security rests on salt entropy and secrecy rather than a cryptographic proof.

The baselines are trivial policies, not trained agents. Every number in Sections 7 and 8 is produced by fixed reference policies (buy-and-hold analogs, one-step tilts, closed-form allocators) and two deliberately biased behavioral counterparties. They fix the zero point of each instrument and calibrate its resolution; they do not establish what a trained agent scores, and this paper’s claims are scoped to environment and protocol for that reason. The constraint is the learner’s, not the simulator’s: Section 6.7 measures the environment at 14,240 steps per second on the slowest of the three shipped surfaces, so a PPO run is bounded by the policy’s own training cost rather than by environment time. In particular, the near-zero within-tier generalization gaps in F3 are the expected control for policies with nothing to overfit, not evidence that the environment resists overfitting by learners. Relatedly, no reference policy is rank-eligible under the kernel’s full conjunction (Section 8.1). Section 8.1.1 establishes that the eligibility set is non-empty on every tier, but it does so with an out-of-band oracle that replays the seed one bar ahead, which is not an entrant and is not scored on any board; no policy that reads only the observation interface has yet been shown eligible.

Named future experiments. Two follow-ups remain specific. (1) A learning baseline: PPO trained through the shipped Gymnasium adapter on the train band, evaluated on a named disjoint evaluation band and cross-regime, so F3 contains a policy that can actually overfit. (2) Tape-level analytic state recovery from quantized SplitMix64-derived prices. The exact-finalizer inverse and the resulting 2^{11} candidates per 53-bit unit are tested, but no observed price has been inverted; sealed seeds (Section 7.4.1) defeat this specific bounded-band scanner but do not answer this one. The manipulation positive control in Section 7.2.2 finds a profitable sampled round trip at $\beta = 0.5$ and none among its sampled linear-impact cells, though the search did not cover all interactions or transient impact. The rank-eligibility existence question of Section 8.1 is answered by the witness of Section 8.1.1, by an oracle rather than by an entrant.

The Python probe layer is outside the golden-hash guarantee. The byte-identical claim covers the Rust core, the WebAssembly build and the generated scenario bytes returned through the Python

binding. The native and Python test suites assert all seven committed scenario fingerprints. Three committed golden files are read alike by the native suite, a `wasm32` build of the exported entry points and the shipped npm bundle: the canonical and clustered scenarios; four entries on the execution and replay path (`run_baseline` and `replay_run`); and nine entries for `walk_forward`, `stress_suite` and `tag_regime`. Direct native–WebAssembly parity tests compare stepping and scenario generation beyond those entries. The backtest and kernel entries pin output bytes of functions that come from the pinned `sharpebench-sim`; they are not inputs to `SPEC_HASH`. The Python binding delegates environment dynamics to that same native Rust core and has deterministic replay tests, but the golden coverage on the Python side is the scenario serialization, not the package’s other serialized artifacts. The Python reductions that produce the realism, manipulation, adverse-selection, failure, ecology and predictability evidence are deterministic in their seeds, but are off the golden-hash path; their floating-point reproducibility is the platform’s, not the paper’s. The concave-impact engine itself is also excluded from the canonical byte-identity claim because a general exponent uses `powf`; the linear default remains on the untouched golden path.

Current engineering checks are not historical treatment variables. Apart from F7, the committed F1–F8 artifacts were not rerun after effective-configuration readback, `SPEC_HASH`, canonical pre-hash fixtures, typed FFI errors, strict input decoding, path-containment checks, the prospective forecast contract or its formal model were added. F7 is the exception: it was regenerated for SharpeArena 0.29.0 after the mandate draw changed, and its artifact carries the effective-configuration block, though not a runtime version. F1 identifies SharpeArena 0.9.0; the other historical evidence files do not carry a runtime version, and the commits that added them lock SharpeBench 0.5.0. Those additions strengthen present and future executions, but they do not change the observations of the artifacts that were not rerun, or prove that each guard was active when an older artifact was created. The manifest, rather than a version number copied into prose, identifies the present engineering tree (Section A).

Some post-freeze repairs do change what the code computes. The paragraph above concerns additions that leave the observations intact. A second group does not. The September 2026 audit of the 0.24.1 tree found defects in code that produced the committed F1–F8 artifacts, and the repairs merged after those artifacts were frozen. The standardized-moment estimators mixed finite-sample normalizations, which moves deflated Sharpes and their intervals. SharpeBench corrected its estimators in 0.19.0 and SharpeArena corrected its own confidence estimator in 0.25.0. The current exact pin, SharpeBench 0.27.0, carries the corrected moments; later SharpeBench releases reach SharpeArena only through a further pin move, which would rebind `SPEC_HASH`. A current baseline row attaches the kernel’s interval only when the pinned kernel reproduces SharpeArena’s point estimate bit for bit on that row’s pooled returns. The frozen F1, F3, witness and predictability artifacts predate both corrections (the commits that added them lock SharpeBench 0.5.0), so their deflated Sharpes and intervals carry the superseded moments and have not been rescored. Episode metrics changed in ways that move reported values rather than merely guarding them: the initial net asset value now contributes to full-episode metrics, the risk-radar drawdown direction and its anchors were reversed, the Sortino target uses full-sample downside root mean square, resource penalties are monotonic for signed base scores, and an episode that failed, went missing or returned a nonfinite value can no longer default to a clean disposition, which changes the failure taxonomy’s denominator. The reward schemes now gate on full-horizon and process eligibility, reading each engine process event’s severity from a contract generated out of the pinned SharpeBench scorer and refusing an event whose severity cannot be established, and floor a failed or incomplete episode at a composite -1 . A market-scenario constraint that previously accepted an infeasible parameter now solves the active-axis cross-section, so a configuration this paper could have expressed is refused as outside the feasible set. The sealed-seed reveal is a further case: its committed record reports sixteen slots replayed, but that count was measured by a check that compared only each symbol’s opening close in a two-day calm environment, while the declared evaluation is 120 days at the hard tier. The check now regenerates the complete tape and requires equality on every bar and symbol, so the frozen number is not evidence for the stronger property the reveal now tests. Changes merged in

mid-September 2026 form a third group. The market-making regret now refuses two arms that drew different random streams, and the committed F2 regrets pass that check unchanged (Section 8.2). The order-book environment’s default inventory mark changed twice, so rewards from runs before the first change reproduce only under `mark="book_mid"` and rewards from runs under the intervening fallback do not reproduce at all; no committed artifact uses the environment. The manipulation probe gained follower externality fields that the committed F5 artifact does not carry, and the calibration pairing of Section 3.2 changed, although no committed artifact carries a calibration field. Of the committed artifacts only F7 has been regenerated, and its rerun reflects the failure-taxonomy repair above and the mandate draw change of Section 8.4; rerunning any other would be a new experiment rather than a correction to what it measured. The repository’s audit checklist records the disposition of every finding.

The realism diagnostic does not certify every experiment’s process. F4 applies to the canonical position-trading generator. The Avellaneda–Stoikov market-making process in F2 and the market-clearing models in F5 and F6 are separate simulators with their own structural and positive-control checks; they are not passed through F4. Their results are therefore calibrations of those specified models, not claims conditioned on the canonical tape’s realism verdict. The distinction matters in both directions: F4’s failure does not mechanically invalidate a different process, and the separate process is not market-valid merely because its internal controls pass.

The impact models are stylized. The endogenous market composes a Kyle-style linear permanent-impact term with an Almgren–Chriss temporary term; the adverse-selection module does not separate a meta-order’s information from its impact; the follower population is a one-parameter caricature. The concave ablation is not a universal test: its dimensionless flow $|Q/V|$ sits two to three orders of magnitude below the unit crossover, so exponents below one amplify impact relative to the linear arm at the tested calibration, and the symmetric schedule does not search the asymmetric round trips nonlinear theory permits. Section 7.2.2 does search them and finds cells that pay at $\beta = 0.5$ with no followers, so the symmetric null says this pump does not pay here, and the positive control says the instrument can see schedules that do; the extended sweep enlarges the sampled axes but does not cover interactions, other flow scales or transient impact.

Permanent impact in the endogenous market accumulates in a multiplier that compounds every bar and never decays, so an exponent below one makes each bar’s increment concave in that bar’s flow but does not produce the square-root law of meta-order impact. For small per-bar increments, a meta-order’s impact grows about linearly in executed quantity at every exponent, a fixed total spread over D bars grows as $D^{1-\beta}$, and nothing relaxes after execution ends. Jafree et al. [2025, pp. 3–4] reproduce the empirical shape in a Hawkes order-book simulator: impact grows as the square root of the quantity executed so far and decays as a power law after execution, with a decay exponent of 0.168 against 0.19 in the empirical study they cite. The rank-neutral probe `meta_order_impact_shape` reports the execution exponent, the post-execution relaxation ratio and the duration exponent for any β . At $\beta = 1$ and 0.5 a test pins the first within 0.05 of one, the second within 10^{-12} of one, and the third within 0.05 of $1 - \beta$. The exponent is also hard to identify from finite data [Tsaknaki et al., 2026, p. 12], so the concave arm here is a stress setting, not an estimate. A transient-impact kernel is future work.

The guard is a deny-list. The structural guarantee covers the environment’s own surface. LookaheadGuard extends it to harness tools by pattern, and a deny-list is exactly as good as its patterns; a novel leak path through a side channel the environment does not own (wall-clock time, filesystem state, a network call) is out of scope for both layers. The replay check verifies decisions against frozen data and therefore catches falsified results, not information an agent obtained elsewhere.

The agent paths are built, and no current-model result is admitted. The contract, conformance kit, replay verifier, strict shared parser, transport-fault gate, edge manifest and its ledger, counterfactual ledger, batched model scheduler, resumable evidence journal, prospective forecast ledger, trace-

promotion path, paper-execution lifecycle and SharpeBench artifact compiler exist and are tested with deterministic model doubles. They do not constitute agent evidence. One obsolete convenience-sample forecast pilot completed the lifecycle and remains committed so its history cannot be rewritten; it is explicitly superseded and its scores are excluded from this paper, model comparison and trading rank. No current-model field has completed, no third party has submitted an agent, and no model-performance number appears here. Model weights may contain public historical prices; controlled inference does not remove pretraining contamination. A hosted intake, live leaderboard and multi-tenant service do not exist. The significance claim remains the verifiability infrastructure that is checkable today; adoption and model performance are future evidence rather than assumptions.

An evaluation sandbox is not hostile-code containment. In this paper, *evaluation sandbox* names the controlled point-in-time market interface: the environment cursor makes future bars unavailable, and LookaheadGuard adds a harness-side refusal layer. SharpeArena itself does not isolate processes or containers. Executable-entry containment belongs to the companion SharpeBench -image path. By the companion paper's own account, its live Docker boundary has been exercised in one hosted-runner and daemon configuration with one benign fixture and a set of named probes, no hostile entrant has been evaluated, and those observations do not rule out a Docker or kernel escape [Toca, 2026]. That is companion engineering evidence, not a SharpeArena empirical result. The inference service remains outside that boundary and inside the trusted computing base. The generated-strategy path of Section 5.4 avoids executable input by accepting only a closed candidate language. A hosted multi-tenant service and a hostile-entrant campaign have not been run.

Partial fills are last-write-wins. The order lifecycle's partial-fill transition sets the cumulative filled quantity from the value the caller supplies rather than summing deltas, so a sequence of partial reports leaves the last one standing. That is what the paper-broker adapters shipped here require, because each report carries a cumulative reading, and the self-transition exists precisely so successive readings are legal. A broker that instead sends incremental deltas would accumulate wrongly against this code and needs different handling; the interface does not currently distinguish the two conventions, and adopting such a broker is a change to the lifecycle rather than a configuration choice.

Reconciliation does not retry, so a flaky broker halts a session. The sweep over unresolved orders is a single pass with no retry and no backoff. An unanswerable query raises, and the raise propagates out of the sweep, so orders it had not yet reached are left for a later call and the session stops. This is the safe direction and it is deliberate: the alternative, swallowing the failure and continuing, is exactly how an unanswerable query becomes a recorded absence and a recorded absence authorizes a duplicate order. The cost is a real one. A broker that answers intermittently will halt a forward session that a retry with backoff would have completed, and the operator has to restart the sweep by hand.

Forward evidence is not historical replay. The paper-trading arm is deliberately paper-only, places no real-capital order and gives the model no credential. Its trusted supervisor applies a fixed risk gate and records every decision, refusal, submission and broker response. Even so, provider-time-dependent bars and simulated fills cannot be reproduced byte for byte after the fact. A paper result must therefore be labeled as forward, non-replayable evidence and bound by commitment, raw-response hashes and broker acknowledgments; it cannot be pooled silently with the deterministic historical tables in this paper, and the admissibility rule of Section 6.6 is what keeps the two apart. A forward record can also be incomplete and still honest: an order left in `submission_unknown` by an unanswered query (Section 5.7) is a fact about what the run observed, and resolving it to accepted or absent without asking the broker would be the only way to make the record look complete. The dual-use surfaces of this release, and the norms we state for them, are in Section 13.

11 Conclusion

What the environment establishes. The environment’s guarantees are properties of committed code, each checkable by a listed test (Section A). The data layer has no operation that returns a bar past the cursor. The determinism-critical core reproduces committed goldens on the native, WebAssembly and packaged surfaces. Replay recomputes performance from recorded decisions and frozen engine inputs, so a doctored decision stream does not reproduce the honest run. The checked and repository scoring paths turn a recorded wire fault into a typed failed cell rather than a hold. Each guarantee has a stated boundary (Section 10): leak-freedom at the interface is not an information guarantee, the Python probe layer sits outside the golden-hash path, and replay does not authenticate every trace field.

What the probes establish. The probes establish less, and several of their results go against the environment. The canonical generator fails its own finite-panel-calibrated realism conjunction on 23 of 24 seeded panels, and a predeclared Calm sweep finds no setting that survives its rule (Sections 7.1 and 7.1.1). No reference policy is rank-eligible, while an out-of-band oracle shows that the eligibility set is non-empty on every tier (Sections 8.1 and 8.1.1). On the market-clearing models the symmetric pump loses money at every sampled coefficient, an asymmetric schedule under concave impact pays on sampled cells without follower flow, and makers keep a positive markout against informed flow at the default calibration (Sections 7.2 and 7.3). A public bounded seed band is inverted from one observed bar, and the sealed derivation defeats that scanner (Sections 7.4 and 7.4.1). The single-seed ecology narrative does not survive replication (Section 8.5). Every one of these numbers calibrates an instrument with fixed reference policies, on tape or on specified models the paper does not claim are market-valid.

What stays open. Two questions remain open, and neither is answered by rereading the committed evidence. The first is learning. No agent is trained here, the near-zero within-tier generalization gaps are the expected control for policies with nothing to overfit, and no current-model field has been admitted, so what a trained or model agent scores, and whether the seed-band protocol separates its skill from band luck, is future evidence. The second is market validity. The canonical tape sits at the lowest level of the empirical-validity scheme on the axis the realism gate grades, the impact models are stylized, and the market-making and market-clearing simulators are calibrations of specified models rather than validated markets. SharpeArena offers a producer whose present guarantees can be checked without trusting its authors, and diagnostics that report where its tape is not yet market-like; closing either open question requires a new experiment rather than a new reading of the old ones.

12 Reproducibility statement

The Cargo workspace has two members, `sharpearena` and `sharpearena-wasm`; the separately built and explicitly excluded `sharpearena-py` crate supplies the PyPI package. The surfaces are published under the MIT and Apache-2.0 licenses, with a pinned toolchain and a lockfile, and depend on published SharpeBench engine crates rather than a vendored copy. The core forbids unsafe code. The determinism claims are enforced rather than asserted: exact canonical JSON fixtures precede the committed FNV-1a hashes of generated scenarios and the matching engine’s fill tape; the WebAssembly crate asserts byte-identity with the native engine for stepping and scenario generation; the scenario, backtest and kernel golden fixtures execute natively, on `wasm32` under Node and against the committed `.wasm`; and the npm suite performs the replay-and-decision-tamper check of Section 3.4. The cross-wrapper `SPEC_HASH` adds a refusal boundary for a mismatched engine and wrapper. The contract artifacts, schemas, conformance fixtures and governance document are committed beside the code they govern, and a conformance test exercises the fixtures.

Continuous integration covers the Rust surface with formatting, lints as errors, tests, a WASM target build and a `wasm32` execution of the committed goldens, plus dependency auditing, npm, a check that the committed WebAssembly bundle answers a fixed call battery exactly as a fresh build of the same

source does, the Python wheel and the evidence manifest. Two jobs test consumer usability rather than source-tree identity: one packs the npm tarball and installs it offline into a fresh project before calling a real export; the other builds the wheel, installs it into a fresh virtual environment outside the checkout, imports it and steps one bar. The latter also asserts that `py.typed` and the native-extension stub ship. The evidence validator runs on every push to main and every pull request and binds the commit it runs on rather than the commits behind it (Section A).

Every number in Sections 7 and 8 is produced by a committed script under `paper/src/` from fixed seeds, writing JSON evidence to `paper/evidence/` and figures to `paper/figures/`; the commands are listed in Section A. The scripts call the package’s public Python API. The current manifest binds the checked artifact set to the present source snapshot, but it cannot reconstruct a runtime version that an older evidence file did not record. Only F7 was regenerated from the current package, for release 0.29.0; the other artifacts come from earlier source revisions. Exact clean-room regeneration therefore requires the original source revision and locked dependencies; where those are absent, the claim is integrity of the committed artifact and producer, not demonstrated regeneration from the current package.

The model-agent, generated-strategy, prospective-forecast and paper-only forward paths of Section 5 are product surfaces, not sources for any admitted number in this manuscript. Nothing has been promoted through the path of Section 6.6 into the F1–F8 evidence set. The archived prospective pilot is a separately labeled engineering trace, not a finding; retaining its raw bytes does not admit its obsolete checkpoint scores. Deterministic tests use injected model transports and fixed decisions; continuous-batched inference, model-server scheduling and live provider data are outside the golden-hash claim. A future current-model field must bind the exact model artifact, inference runtime, prompt and scaffold, decoding configuration, sampling seed, raw completion, semantic-validation outcome and fault class. Environment bytes and replay can remain fixed while agent outputs vary, and the evidence schema preserves that distinction.

The products compose through files rather than a mutual dependency. SharpeArena writes and validates complete append-only cell journals and prospective forecast ledgers; SharpeBench independently reads those artifacts for field-level evaluation. The executable cross-product tutorial fixes a synthetic two-producer artifact and checks the resulting report; the scoring it exercises is specified in the companion paper [Toca, 2026]. The two Lean projects model complementary invariants: Arena’s deadline, revision, binding and fixed-point Brier rules, and Bench’s exact pair support as the intersection of two agents’ resolved contracts, Holm monotonicity and bound, finite-bootstrap correction and trading-rank isolation. These proofs do not extract or verify the Python and Rust implementations. Executable tests cover the same rules in code independently of the models, and nothing mechanically links a model to the code it describes.

No large language model is a component of the environment, the contract or any admitted result. The environment has no judge. The superseded engineering pilot is retained outside the paper’s empirical conclusions. Section 14 states how language models were used in preparing the manuscript.

13 Ethics and broader impact

Dual-use surfaces. Three parts of this release can be turned against the purpose they were built for. The manipulation probe of Section 7.2 is a payoff-search harness over trade-based manipulation schedules [Allen and Gale, 1992], and is dual-use by construction. The in-band inversion probe of Section 7.4 publishes a scan that recovers a scenario’s seed, and with it every future bar, from one observed bar of a public bounded seed band in about one second. And a byte-verifiable leaderboard score is more persuasive than an unverifiable one without being more externally valid.

The profitable manipulation cells. The paper reports manipulation schedules that pay, and states where. Every sampled coefficient of the symmetric linear sweep loses money, and so does every one of the 45 sampled linear cells of the asymmetric search, which is consistent with the no-manipulation conditions of the linear model under its assumptions. Under normalized concave permanent impact

at $\beta = 0.5$, slow accumulation followed by block liquidation pays on the two arms without follower flow: four of the 135 cells keep familywise-positive intervals, the selected cell replicates on 32 fresh seeds, and the extended sweep’s largest sampled cell is a longer trip of the same shape; the canonical follower arm has no pointwise-positive cell (Sections 7.2.2 and 7.2.3). These are statements about a stylized simulator at one calibration, gross of fees, in discrete time with no latency and with a one-parameter follower population. They are not estimates of what any schedule earns in a market.

What is and is not released. Released: the environment, the probe modules, the producer scripts and the committed per-cell evidence behind every reported number, the manipulation grid included, in the public repository under the MIT and Apache-2.0 licenses, with the environment and probe modules also in the published packages. Not released, because none exists: a real-market price series inside the environment, a trained policy, model weights, an admitted model-performance result, an execution path from the manipulation module to any order adapter, and a real-capital endpoint. Inside the Python package the manipulation module is imported only by the package namespace, not by the forward paper arm, and that arm’s order adapter refuses every origin other than a paper-trading endpoint (Section 5.7). The module’s own documentation states that it is a simulator diagnostic and not a strategy, that its schedules are written to be crude and obvious, and that no output of it should be handed to an execution surface. That is a stated norm, not an enforced one.

Misuse. We name three misuses. The first is presenting a SharpeArena score, verified or not, to investors or retail audiences as evidence of real-market skill: a verified number is a claim about a synthetic environment only, and SharpeArena scores are not performance marketing. The second is reading the profitable concave cells as a trading recipe, or fitting or deploying schedules taken from the manipulation module. The third is running the seed-band scan against an evaluation one has entered, which the protocol forbids (Section 6.3). No license term enforces any of the three; this section states them as norms.

Why we publish the probes. We judge publishing them net positive. The linear arm is a diagnostic against a stated no-manipulation assumption rather than a search for deployable trades, and the paper makes no guarantee outside its finite grid. The concave positive control shows that the instrument can see a paying schedule, which a probe that has never returned a positive cannot claim. Publishing the seed scan beside the sealed-seed derivation of Section 7.4.1 turns an attack a motivated entrant could run silently into a documented threat with a shipped mitigation.

Data and people. No human subjects and no personal data are involved. The environment’s tape is synthetic, and the forward paper arm reads public market data read-only. No model inference contributes to any admitted result; the use of language models in preparing the manuscript is stated in Section 14.

14 Declarations

Use of language models. Language models helped draft and edit the prose, read the cited papers and propose candidate facts, and write parts of the code the paper describes. Every fact taken from a cited paper was checked against that paper’s text, and the page is given where the fact is used; every bibliography entry was checked against its source. No language model is a component of the environment, the contract or any admitted result (Section 12).

Competing interests. The author is affiliated with General Liquidity, which publishes both SharpeArena and the companion SharpeBench. Every reference policy, probe and producer evaluated here was built for this repository, and the companion scorer is the same author’s work [Toca, 2026], so the producer–scorer separation of Section 6.5 separates code and artifacts, not authors.

Data and code availability. The source code, the contract schemas and conformance fixtures, the producer scripts and the committed evidence behind every reported number are in the public

repository <https://github.com/general-liquidity/sharpearena>, dual-licensed under MIT or Apache-2.0; Section A lists the command behind each artifact. No market data enters any reported result: the tape is synthetic and every scenario is regenerated from its seed.

References

- Franklin Allen and Douglas Gale. Stock-price manipulation. *The Review of Financial Studies*, 5(3): 503–529, 1992. doi: 10.1093/rfs/5.3.503.
- Robert Almgren and Neil Chriss. Optimal execution of portfolio transactions. *Journal of Risk*, 3(2): 5–39, 2001. doi: 10.21314/JOR.2001.041.
- Robert Almgren, Chee Thum, Emmanuel Hauptmann, and Hong Li. Direct estimation of equity market impact. *Risk*, 18(7):58–62, 2005.
- Selim Amrouni, Aymeric Moulin, Jared Vann, Svitlana Vyetrenko, Tucker Balch, and Manuela Veloso. ABIDES-Gym: Gym environments for multi-agent discrete event simulation and application to financial markets. In *Proceedings of the Second ACM International Conference on AI in Finance (ICAIF)*, pages 1–9, 2021. doi: 10.1145/3490354.3494433.
- Marco Avellaneda and Sasha Stoikov. High-frequency trading in a limit order book. *Quantitative Finance*, 8(3):217–224, 2008. doi: 10.1080/14697680701381228.
- Robert L. Axtell and Joshua M. Epstein. Agent-based modeling: Understanding our creations. *The Bulletin of the Santa Fe Institute*, 9(2):28–32, 1994.
- David H. Bailey and Marcos López de Prado. The deflated sharpe ratio: Correcting for selection bias, backtest overfitting, and non-normality. *The Journal of Portfolio Management*, 40(5):94–107, 2014. doi: 10.3905/jpm.2014.40.5.094.
- David H. Bailey and Marcos López de Prado. The Sharpe ratio efficient frontier. *Journal of Risk*, 15(2):3–44, 2012. doi: 10.21314/JOR.2012.255.
- Sidnei Barbieri, Wellington Vargas, and Ágney Lopes Roth Ferraz. Auditing AI investment recommendations as executable actions. arXiv preprint arXiv:2606.27570v1, 2026.
- Sylvain Barde and Sander van der Hoog. An empirical validation protocol for large-scale agent-based models. Studies in Economics 1712, School of Economics, University of Kent, July 2017.
- T. J. Barton, Chris Constantakis, Patti Hauseman, Annie Mous, Alaska Hoffman, Brian Bergeron, and Hunter Goodreau. What LLM trading agents actually do in production: A six-month, population-scale record from two fleets. arXiv preprint arXiv:2609.05663v1, 2026.
- Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. OpenAI Gym. arXiv preprint arXiv:1606.01540, 2016.
- Eric Budish, Peter Cramton, and John Shim. The high-frequency trading arms race: Frequent batch auctions as a market design response. *The Quarterly Journal of Economics*, 130(4):1547–1621, 2015. doi: 10.1093/qje/qjv027.
- David Byrd, Maria Hybinette, and Tucker Hybinette Balch. ABIDES: Towards high-fidelity multi-agent market simulation. In *Proceedings of the 2020 ACM SIGSIM Conference on Principles of Advanced Discrete Simulation (SIGSIM-PADS)*, pages 11–22, 2020. doi: 10.1145/3384441.3395986.
- Patrick Cheridito, Jean-Loup Dupret, and Zhexin Wu. ABIDES-MARL: A multi-agent reinforcement learning environment for endogenous price formation and execution in a limit order book. arXiv preprint arXiv:2511.02016v1, 2025.

- Karl Cobbe, Christopher Hesse, Jacob Hilton, and John Schulman. Leveraging procedural generation to benchmark reinforcement learning. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 2048–2056, 2020.
- Rama Cont. Empirical properties of asset returns: Stylized facts and statistical issues. *Quantitative Finance*, 1(2):223–236, 2001. doi: 10.1080/713665670.
- Yimin Du. Machine learning enhanced multi-factor quantitative trading: A cross-sectional portfolio optimization approach with bias correction. arXiv preprint arXiv:2507.07107v2, 2026.
- Giorgio Fagiolo, Mattia Guerini, Francesco Lamperti, Alessio Moneta, and Andrea Roventini. Validation of agent-based models in economics and finance. In Claus Beisbart and Nicole J. Saam, editors, *Computer Simulation Validation: Fundamental Concepts, Methodological Frameworks, and Philosophical Perspectives*, Simulation Foundations, Methods and Applications, pages 763–787. Springer International Publishing, Cham, 2019. doi: 10.1007/978-3-319-70766-2_31. Circulated as LEM Working Paper 2017/23, Sant’Anna School of Advanced Studies.
- Farama Foundation. Minari: A dataset API for offline reinforcement learning. <https://minari.farama.org>, 2024. Software documentation, accessed 15 September 2026; release 0.5.0 of September 2024 archived at doi:10.5281/zenodo.13767625.
- J. Doyne Farmer. Market force, ecology and evolution. *Industrial and Corporate Change*, 11(5): 895–953, 2002. doi: 10.1093/icc/11.5.895.
- Óscar Fernández Vicente. *Market Making Strategies with Reinforcement Learning*. PhD thesis, Universidad Carlos III de Madrid, January 2025.
- Sara Fish, Julia Shephard, Minkai Li, Ran Shorrer, and Yannai A. Gonczarowski. EconEvals: Benchmarks and litmus tests for economic decision-making by LLM agents. arXiv preprint arXiv:2503.18825v4, 2026.
- Weilong Fu. The new quant: A survey of large language models in financial prediction and trading. arXiv preprint arXiv:2510.05533v1, 2025.
- Jim Gatheral. No-dynamic-arbitrage and market impact. *Quantitative Finance*, 10(7):749–759, 2010. doi: 10.1080/14697680903373692.
- Eray Gençay. What survives honest evaluation? Leakage-safe, search-aware assessment of LLM-driven trading strategy discovery. arXiv preprint arXiv:2608.27734v1, 2026.
- Lawrence R. Glosten and Paul R. Milgrom. Bid, ask and transaction prices in a specialist market with heterogeneously informed traders. *Journal of Financial Economics*, 14(1):71–100, 1985. doi: 10.1016/0304-405X(85)90044-3.
- Tilmann Gneiting and Adrian E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007. doi: 10.1198/016214506000001437.
- Olivier Guéant, Charles-Albert Lehalle, and Joaquin Fernandez-Tapia. Dealing with the inventory risk: A solution to the market making problem. *Mathematics and Financial Economics*, 7(4): 477–507, 2013. doi: 10.1007/s11579-012-0087-0.
- Jiacheng Guo, Suozhi Huang, Yunlong Gao, Zihao Li, Jian Ge, Xu Kuang, and Mengdi Wang. AQuA: Recursively self-improving quantitative trading research agents. arXiv preprint arXiv:2608.12841v1, 2026.
- Ryuji Hashimoto, Ryosuke Takata, Masahiro Suzuki, Yuki Tanaka, and Kiyoshi Izumi. Financial market as a self-organized ecosystem: Simulation via learning with heterogeneous preferences. arXiv preprint arXiv:2604.23975v1, 2026.

- Gur Huberman and Werner Stanzl. Price manipulation and quasi-arbitrage. *Econometrica*, 72(4): 1247–1275, 2004. doi: 10.1111/j.1468-0262.2004.00531.x.
- Ali Raza Jafree, Konark Jain, and Nick Firoozye. When AI trading agents compete: Adverse selection of meta-orders by reinforcement learning-based market making. arXiv preprint arXiv:2510.27334v1, 2025.
- Shakya Jayakody and Prarthinie Jayakody. KineticSim: A lightweight, high-performance execution engine for real-time market simulators. arXiv preprint arXiv:2606.21784v2, 2026.
- Joseph Jerome, Leandro Sánchez-Betancourt, Rahul Savani, and Martin Herdegen. mbt-gym: Reinforcement learning for model-based limit order book trading. In *Proceedings of the Fourth ACM International Conference on AI in Finance (ICAIF)*, pages 619–627, 2023. doi: 10.1145/3604237.3626873.
- Minqi Jiang, Edward Grefenstette, and Tim Rocktäschel. Prioritized level replay. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 4940–4950, 2021.
- Guy L. Steele Jr., Doug Lea, and Christine H. Flood. Fast splittable pseudorandom number generators. In *Proceedings of the 2014 ACM International Conference on Object Oriented Programming Systems Languages and Applications (OOPSLA)*, pages 453–472, 2014. doi: 10.1145/2660193.2660195.
- Albert S. Kyle. Continuous auctions and insider trading. *Econometrica*, 53(6):1315–1335, 1985. doi: 10.2307/1913210.
- Francesco Lamperti. An information theoretic criterion for empirical validation of simulation models. *Econometrics and Statistics*, 5:83–106, 2018. doi: 10.1016/j.ecosta.2017.01.006.
- Yuqi Li, Siyuan Liu, and Bingjun Liu. AI trading’s alpha singularity: Emergent market reasoning through agent-to-agent self-evolution. arXiv preprint arXiv:2606.29194v1, 2026.
- Xiao-Yang Liu, Hongyang Yang, Qian Chen, Runjia Zhang, Liuqing Yang, Bowen Xiao, and Christina Dan Wang. FinRL: A deep reinforcement learning library for automated stock trading in quantitative finance. arXiv preprint arXiv:2011.09607, 2020. Deep RL Workshop, NeurIPS 2020.
- Xiao-Yang Liu, Ziyi Xia, Jingyang Rui, Jiechao Gao, Hongyang Yang, Ming Zhu, Christina Dan Wang, Zhaoran Wang, and Jian Guo. FinRL-Meta: Market environments and benchmarks for data-driven financial reinforcement learning. In *Advances in Neural Information Processing Systems 35, Datasets and Benchmarks Track*, 2022.
- Sichun Luo, Yi Huang, Shichang Meng, Fengyuan Liu, Mukai Li, Qinghua Yao, Zefa Hu, Junlan Feng, and Qi Liu. AlphaQT-Bench: Diagnosing the gap between financial code generation and quantitative reasoning in LLMs. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 2871–2886, 2026.
- Thomas Lux and Michele Marchesi. Scaling and criticality in a stochastic multi-agent model of a financial market. *Nature*, 397(6719):498–500, 1999. doi: 10.1038/17290.
- Shaocong Ma and Heng Huang. Robust reinforcement learning in finance: Modeling market impact with elliptic uncertainty sets. In *Advances in Neural Information Processing Systems 38*, 2025. arXiv:2510.19950v3.
- Marlos C. Machado, Marc G. Bellemare, Erik Talvitie, Joel Veness, Matthew Hausknecht, and Michael Bowling. Revisiting the arcade learning environment: Evaluation protocols and open problems for general agents. *Journal of Artificial Intelligence Research*, 61:523–562, 2018. doi: 10.1613/jair.5699.

- Peer Nagy, Sascha Yves Frey, Kang Li, Bidipta Sarkar, Svitlana Vyetrenko, Stefan Zohren, Ani Calinescu, and Jakob Nicolaus Foerster. LOB-Bench: Benchmarking generative AI for finance – an application to limit order book data. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, volume 267 of *Proceedings of Machine Learning Research*, pages 45437–45460, 2025. arXiv:2502.09172.
- Donovan Platt and Tim Gebbie. Can agent-based models probe market microstructure? *Physica A: Statistical Mechanics and its Applications*, 503:1092–1106, 2018. doi: 10.1016/j.physa.2018.08.055. arXiv:1611.08510.
- William F. Sharpe. Mutual fund performance. *The Journal of Business*, 39(1, Part 2):119–138, 1966. doi: 10.1086/294846.
- William F. Sharpe. The Sharpe ratio. *The Journal of Portfolio Management*, 21(1):49–58, 1994. doi: 10.3905/jpm.1994.409501.
- Issa Sugiura, Daichi Hattori, Kazuo Araragi, Keita Ogawa, Shota Onose, Taro Makino, Teppei Usuki, and Takashi Ishida. CoffeeBench: Benchmarking long-horizon LLM agents in heterogeneous multi-agent economies. arXiv preprint arXiv:2606.16613v1, 2026.
- Shuo Sun, Molei Qin, Wentao Zhang, Haochong Xia, Chuqiao Zong, Jie Ying, Yonggang Xie, Lingxuan Zhao, Xinrun Wang, and Bo An. TradeMaster: A holistic quantitative trading platform empowered by reinforcement learning. In *Advances in Neural Information Processing Systems 36, Datasets and Benchmarks Track*, 2023.
- Jordan Terry, Benjamin Black, Nathaniel Grammel, Mario Jayakumar, Ananth Hari, Ryan Sullivan, Luis S. Santos, Clemens Dieffendahl, Caroline Horsch, Rodrigo Perez-Vicente, Niall Williams, Yashas Lokesh, and Praveen Ravi. Pettingzoo: Gym for multi-agent reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 34, 2021.
- Tiberiu Toca. SharpeBench: A Luck-Robust Benchmark for Trading Agents. Preprint, <https://github.com/general-liquidity/sharpebench>, 2026. Version 0.27.0. The SharpeArena evidence cited here was produced against SharpeBench 0.5.0, except the F7 artifact, which was regenerated against SharpeBench 0.21.0.
- Mark Towers, Ariel Kwiatkowski, Jordan Terry, John U. Balis, Gianluca De Cola, Tristan Deleu, Manuel Goulão, Andreas Kallinteris, Markus Krimmel, Arjun KG, Rodrigo Perez-Vicente, Andrea Pierré, Sander Schulhoff, Jun Jet Tai, Hannah Tan, and Omar G. Younis. Gymnasium: A standard interface for reinforcement learning environments. In *Advances in Neural Information Processing Systems 38, Datasets and Benchmarks Track*, 2025. arXiv:2407.17032.
- Ioanna-Yvonne Tsaknaki, Andrea Macrì, and Fabrizio Lillo. Can reinforcement learning efficiently discover price manipulation? arXiv preprint arXiv:2607.06121v1, 2026.
- Svitlana Vyetrenko, David Byrd, Nick Petosa, Mahmoud Mahfouz, Danial Dervovic, Manuela Veloso, and Tucker Hybinette Balch. Get real: Realism metrics for robust limit order book market simulations. In *Proceedings of the First ACM International Conference on AI in Finance (ICAIF)*, pages 1–8, 2020. doi: 10.1145/3383455.3422561. arXiv:1912.04941.
- Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. In *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025. arXiv:2406.12045.
- Zhiyuan Zeng, Jiashuo Liu, et al. FutureX: An advanced live benchmark for LLM agents in future prediction. arXiv preprint arXiv:2508.11987v3, 2025.
- Kaicheng Zhang, Wen Ge, Lei Jiang, Weixin Yang, Jordan Langham-Lopez, Jialin Yu, Lukasz Szpruch, and Hao Ni. OpenFinGym: A verifiable multi-task gym environment for evaluating quant agents. arXiv preprint arXiv:2606.26350v1, 2026.

Mingxi Zou, Jiaxiang Chen, Aotian Luo, Jingyi Dai, Chi Zhang, Dongning Sun, and Zenglin Xu.
FinEvo: From isolated backtests to ecological market games for multi-agent financial strategy
evolution. arXiv preprint arXiv:2602.00948v1, 2026.

Gilles Zumbach. Time reversal invariance in finance. *Quantitative Finance*, 9(5):505–515, 2009. doi:
10.1080/14697680802616712.

A Commands that produce every number

All commands run from the repository root with the `sharpearena` Python package installed. Each script fixes and serializes its seeds, and the evidence manifest records file hashes and producer commands.

Version provenance. This is the paper’s single provenance statement. F1 serializes `sharpearena` 0.9.0; the other F1–F8 evidence files do not serialize a runtime version, so this paper does not reconstruct one from their creation dates. Historical review records identify SharpeBench 0.5.0 as the companion scorer used during that evidence cycle, and the commits that added every artifact except F7 lock SharpeBench 0.5.0, but the empirical JSON does not serialize a SharpeBench version. F7 is the one artifact regenerated since: it was rerun for SharpeArena 0.29.0, whose crates pin SharpeBench at `=0.21.0`, and it carries the effective-configuration readback. The present engineering guarantees are verified against the source and tests bound by the manifest; they do not retroactively establish that every current guard was active when the historical empirical tables were generated. F4 and F5 are generated by the documented finite-panel estimator and normalized-flow producers, and the witness artifact serializes its producer and common-random-number configuration. No current-model artifact or trained-policy result is admitted. A superseded convenience-sample forecast pilot is retained as an engineering trace and excluded from every result in this paper.

Content provenance.

```
python paper/src/make-provenance.py
python paper/src/check-provenance.py
```

The first command writes `paper/evidence/provenance.json`; the second validates it against the tree. The manifest records the commit it was generated at, whether the working tree was dirty at that moment, a content hash over the Rust, Python, manifest and producer-source snapshot, and the SHA-256 digest of every JSON and figure artifact. A separate model-artifact scope accepts closed identity files; the manifest binds each file and exposes its model digest and inference-runtime identity as first-class records. An empty model-artifact list therefore means no model result has entered the paper, rather than that model provenance was omitted. Archived protocol JSON may still be present in the ordinary artifact scope without being admitted as model evidence. The source snapshot is source only: paths under build directories, virtual environments and bytecode caches are excluded, and a text file’s digest is taken with CRLF collapsed to LF, so the same tree yields the same snapshot hash on a different machine and on a Windows checkout. Files carrying a NUL byte are hashed verbatim. The manifest is a release-time snapshot rather than a continuously maintained view of the working tree: regenerating it is part of the release procedure, and `check-provenance.py` recomputes every recorded digest, re-expands both artifact scopes and exits nonzero on any mismatch, so a manifest that has fallen behind its tree is a detectable failure rather than a silent one. Continuous integration runs that check on every push to main and every pull request, so the failure is detectable without anyone remembering to look. What it establishes is a property of the commit the job runs on, which is the tip of the pushed branch and the head of the pull request, and not a property of every commit in the history. The manifest is bound in a commit of its own after the work it describes, so a commit that moved a source file carries a manifest that has not caught up until the binding commit lands; that is how the artifact is produced rather than a defect in it, and the gate is what makes the state a reader fetches or reviews bind its own tree. Three properties of the validator matter for what the check is

worth. The writer and the validator share one rule module, so the scope, exclusions, schema, digest convention, validator path, reproduction entry point and model identity fields are validator policy rather than values a manifest may narrow. Both missing and extra records fail, and every exact model record is opened and its identity recomputed; a model identity file whose artifact digest or runtime identity reads as unresolved is refused by both rather than written by one and accepted by the other. `generated_at_head_dirty` is a claim the manifest makes about itself, so the validator re-derives it and fails a manifest asserting a clean generation against a tree that has uncommitted work. And a manifest that claims a clean generation is held to what that claim means: every recorded digest is compared against the bytes committed at `generated_at_head`, read back out of the repository; a missing commit is a failure, so editing the flag or naming an unavailable commit cannot bypass the check. The writer can record a dirty candidate for inspection, but the validator refuses it because no later check can reconstruct what that uncommitted tree looked like when the manifest was written. A validated manifest therefore always names a clean committed generation. The checked release driver waits until the version bump is committed and clean, regenerates the manifest, commits only that manifest and tags the resulting provenance-only child. Before any registry job starts, the release validator checks that parent relation, every recorded digest and expanded scope, each model identity summary, the snapshot hash and every package version against the exact tagged tree. A published tag must be an ancestor of `origin/main`; only the private pre-push check admits the prospective child that the driver will push atomically. Validation passes the exact checked commit identifier to every publishing job, so none resolves the mutable tag again. Read that way, the content snapshot identifies the current source state the manifest binds, and a package label does not; neither identifies the source state from which the original numbers were produced, which, as the version paragraph above records, was not always serialized.

The design-property checks (Sections 3.1, 3.3, 3.4, 4, 4.5 and 5).

```
cargo test -p sharpearena          # golden hashes, leak-free
                                   # property tests, conformance kit
cargo test -p sharpearena-wasm     # native == wasm parity + golden
wasm-pack test --node crates/sharpearena-wasm
                                   # same goldens, run on wasm32
npm test --prefix npm/sharpearena  # replay + tamper detection, and the
                                   # goldens against the shipped .wasm
python -m pytest crates/sharpearena-py # Python surface, check_env,
                                   # seed-band disjointness, canonical
                                   # parser, model-field scheduler,
                                   # strategy ledger and paper-only arm
cd formal && lake build              # abstract forecast invariants
```

The `cargo test` line compiles the WebAssembly crate's JSON kernels for the host, so it checks the shared source rather than the WebAssembly build; the `wasm-pack` line is the one that executes the `wasm-bindgen` export layer on `wasm32`, and the `npm` line is the one that reads the binary the package publishes. All three assert the same three committed golden files, for scenario generation, the backtest and replay path, and the `walk_forward`, `stress_suite` and `tag_regime` exports. The model-field tests inject deterministic model transports and do not download weights or produce an agent-performance result. The MCP regression cases assert that unknown and duplicate symbols and action/weight contradictions fail without advancing the environment. The paper-arm tests use in-memory transports and fixed endpoints; they do not send a broker order. The transport-gate cases of Section 4.5 include the characterization test that runs a wedged agent through the unguarded engine and asserts a full-length scoreable return series, so the defect the gate closes stays visible in the suite rather than only in the changelog. Companion-container evidence belongs to SharpeBench and is outside the commands above (Section 10).

The current golden tests compare the literal canonical JSON fixtures before computing the FNV-1a fingerprints. The `SPEC_HASH` tests separately bind the tape-defining source set to both wrappers and

the committed npm binary. Python tests exercise effective-configuration readback, typed native errors, sealed-salt redaction, unknown-field refusal, plan-path containment, the shared target-weight event contract and the process-event severity contract. These are current engineering checks: running them does not regenerate or reclassify the historical evidence files.

Prospective forecast protocol and cross-product tutorial (Section 5.3).

```
python examples/forecast-quality/tutorial.py \  
  --output-dir /tmp/sharpe-forecast-tutorial \  
  --sharpebench-dir ../sharpebench
```

The deterministic tutorial writes two fields of strict forecast ledgers with a digest manifest each, asks the sibling SharpeBench checkout to score their exact common support, and compares each result with the committed fixture. The supported field's twelve synthetic questions span six resolution-time blocks and the withheld field's first eight of them span two; both include an effective revision, a late attempt that cannot replace it, a pre-open rejection, and blind and consensus-visible information exposures. It is a conformance example, not an empirical model run. The separate prospective runner has prepare, forecast, seal and resolve phases, but no invocation of it contributes a result to this paper.

F1: baseline leaderboard (Section 8.1).

```
python paper/src/make-f1-baselines.py  
-> paper/evidence/f1-baselines.json, paper/figures/f1-baselines.pdf
```

The script passes the bootstrap parameters explicitly ($B = 2000$ percentile resamples, $\alpha = 0.05$, resample seed 0x5BA7_2026) and serializes them in the evidence config; they equal the library defaults defined in `crates/sharpearena-py/python/sharpearena/confidence.py` (DEFAULT_N_BOOT, DEFAULT_RESAMPLE_SEED), which reuse the scoring kernel's canonical bootstrap seed.

F2: regret vs the analytical reference (Section 8.2).

```
python paper/src/make-f2-regret.py  
-> paper/evidence/f2-regret.json, paper/figures/f2-regret.pdf
```

F3: generalization gap and cross-regime transfer (Section 8.3).

```
python paper/src/make-f3-generalization.py  
-> paper/evidence/f3-generalization.json
```

The transfer-gap matrix is reported as Table 12 only; the paper carries no F3 figure.

F4: stylized-facts realism certification (Section 7.1).

```
python paper/src/make-f4-realism.py  
-> paper/evidence/f4-realism.json, paper/figures/f4-realism.pdf,  
   paper/figures/f4-calm-calibration.pdf
```

The same command also runs the 99-configuration Calm calibration sweep, serializes it under `calm_calibration` and writes `paper/figures/f4-calm-calibration.pdf` (Section 7.1.1).

F5: manipulation boundary and size response (Section 7.2).

```
python paper/src/make-f5-manipulation.py  
-> paper/evidence/f5-manipulation.json,  
   paper/figures/f5-boundaries.pdf, paper/figures/f5-size-response.pdf
```

The same command also runs the concave-impact extension at $\beta = 0.5$ and 0.7 , serializes it under `concave`, and writes `paper/figures/f5-concave.pdf` (Section 7.2.1). It then runs the asymmetric positive control (Section 7.2.2): 135 grid points over five duration ratios, three block fractions and $\beta \in \{1.0, 0.7, 0.5\}$ on the theory, temporary-impact and canonical arms, eight seeds each, plus one fixed selected cell on 32 fresh seeds under `positive_control.confirmation`; the producer serializes all configurations and writes `paper/figures/f5-positive-control.pdf`. It also runs the 60-cell extended sweep using independently normalized uniform increments for unequal legs, recorded under `extended_sweeps` and written to `paper/figures/f5-extended-sweeps.pdf`. Because its anchor is selected from the first grid on the same eight seeds, each extension interval corrects over the global 195-cell family (135 initial plus 60 extension cells); that correction is descriptive for the data-adaptively chosen extension cells rather than simultaneous coverage (Section 7.2.3).

Rank-eligibility existence witness (Section 8.1.1).

```
python paper/src/make-witness.py
-> paper/evidence/witness.json, paper/figures/witness.pdf
```

The command replays each seed one bar ahead to recover the true next-bar return (asserting that the canonical tape is policy-invariant), mixes it with a standardized noise path reused at every strength within that path (common random numbers), scores `sign_follow` and `deadband_hold` through `score_run` on every seed and on the pooled series, and first checks eligibility is monotone over the 13-point coarse grid. It bisects only a monotone first crossing, on the gap-band witness subset and the F1 table band, and stops once the bracket is no wider than 0.005 in s (`bisect_resolution`; the `bracket_half_width` of 0.0025 the evidence records is half that bound). Every bracket starts on a coarse-grid interval of 0.05, 0.1 or 0.2, so halving stops at width 0.003125: every recorded bracket, primary and replicate, has width 0.0031 and half-width 0.0016, which is the resolution Table 11 reports. The whole procedure is repeated under five independent noise paths per cell, with noise-seed uniqueness asserted in the producer; the first path is serialized as the primary witness and all five are serialized under `noise_replicates`, which is what Table 11 summarizes. The JSON records every strength's per-seed PSR, pooled DSR, bootstrap p -value, eligibility flag, monotonicity check and every bracketed crossing.

Sealed evaluation seeds (Section 7.4.1).

```
python paper/src/make-sealed-seeds.py
-> paper/evidence/sealed-seeds.json
```

The command derives 16 named slots publicly (the eight committed offsets plus eight drawn from the same 2^{16} band) and sealed (the same slots under a 32-byte salt from the operating system, committed by SHA-256 before the run), re-runs the predictability probe's band-scan adversary unchanged against both, then reveals the salt and replays the sealed scenarios against the recomputed seeds. The JSON records the commitment, the revealed salt, the table-build time and every per-trial recovery and replay record.

F6: adverse-selection markouts (Section 7.3).

```
python paper/src/make-f6-adverse-selection.py
-> paper/evidence/f6-adverse-selection.json,
    paper/figures/f6-markouts.pdf, paper/figures/f6-endogenous.pdf
```

The same command runs the endogenous arm and its six-value impact sweep, serializes them under the `endogenous` key and writes `paper/figures/f6-endogenous.pdf` (Section 7.3.1).

F7: failure-mode distributions (Section 8.4).

```
python paper/src/make-f7-failures.py
-> paper/evidence/f7-failures.json, paper/figures/f7-failures.pdf
```

F8: ecology under shocks, 8 seeds per schedule (Section 8.5).

```
python paper/src/make-f8-ecology.py
-> paper/evidence/f8-ecology.json,
   paper/figures/f8-ecology-control.pdf,
   paper/figures/f8-ecology-shocked.pdf
```

The evidence carries the full seed-0 reports (the figures' source) plus, under `multi_seed`, per-seed winners, outcome counts and the cross-seed winner distributions behind Table 14.

Generator predictability (Section 7.4).

```
python paper/src/make-predictability.py
-> paper/evidence/predictability.json,
   paper/figures/predictability.pdf
```

The command evaluates baseline, causal AR(5) and true-seed oracle predictors on 16 seeds per tier, then times exhaustive recovery in a 2^{16} candidate band. The JSON contains every per-seed score and the measured search timing.

Throughput and compute cost (Section 6.7).

```
python paper/src/make-throughput.py
-> paper/evidence/throughput.json
```

The command times the Python surface directly (scenario generation, Gymnasium-adapter stepping, and one full six-policy evaluation over the canonical held-out band on all three tiers), then shells out to the native and WebAssembly benchmarks and parses their output, so all three rows of Table 4 come from one invocation. The two sub-benchmarks also run standalone:

```
cargo run --release -p sharpearena --example bench-steps
node npm/sharpearena/bench/throughput.js
```

Every rate is the median of three repeats of 500 episodes on a 4×120 panel, after a warm-up invocation that is discarded; the evidence JSON records every individual repeat and the host's platform string, so a reader can see the spread rather than only the median. Throughput is hardware-dependent and is the one quantity in this paper that a replicator should expect to differ.

Dispersion layers. The F2, F3, F5 and F6 scripts serialize the per-episode or per-seed vectors behind every reported interval into their evidence JSONs (`per_episode_regret`, `per_seed_returns` with `band_dsr_ci` and `transfer_gap_ci`, `dispersion`, and `per_episode_markout_per_unit` with `gap_stats` respectively), each with its CI convention recorded in the file, so every interval in Sections 7 and 8 is recomputable from the committed evidence.

Regenerating every figure from the committed evidence.

```
python paper/src/make-figures.py
```

Each `make-f*` script produces its own figure at run time; `make-figures.py` re-renders all figures from the JSON already in `paper/evidence/` without re-running any experiment, so a reader can check that no figure contains a number the evidence does not.